

Adventures in SIMD Thinking (Part 2 of 2)

Bob Steagall
CppCon 2020

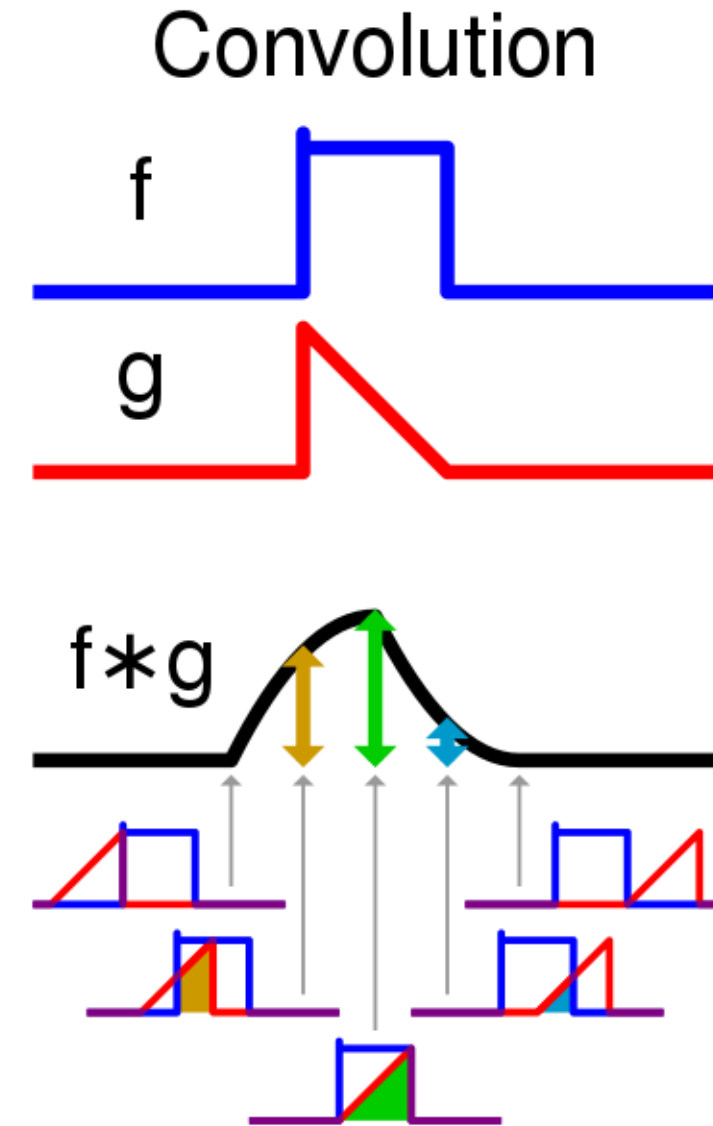


- Learn a little about Intel's SIMD facilities (disclaimer: I don't work for Intel)
- Create some useful functions in terms of AVX-512 intrinsics
- Try some SIMD-style thinking to tackle some interesting problems
 - Intra-register sorting
 - Fast linear median-of-seven filter
 - Fast small-kernel convolution
 - Faster (?) UTF-8 to UTF-32 conversion (with AVX2)
- No heavy code, but lots of pictures
 - Thinking "vertically"

Small-Kernel Convolution

Convolution

- f is a signal
- g is a kernel
- Output $f * g$ is the convolution
 - Every point of result $f * g$ is f weighted by every point of g
- Useful for smoothing and de-noising



Convolution

$$\begin{array}{rcccccccc} S & = & s_0 & & s_1 & & s_2 & & s_3 & & s_4 & & s_5 & & s_6 & \dots \\ K & = & k_0 & & k_1 & & k_2 & & & & & & & & & \end{array}$$

$$r_0 = \dots$$

$$r_1 = s_0k_0 + s_1k_1 + s_2k_2$$

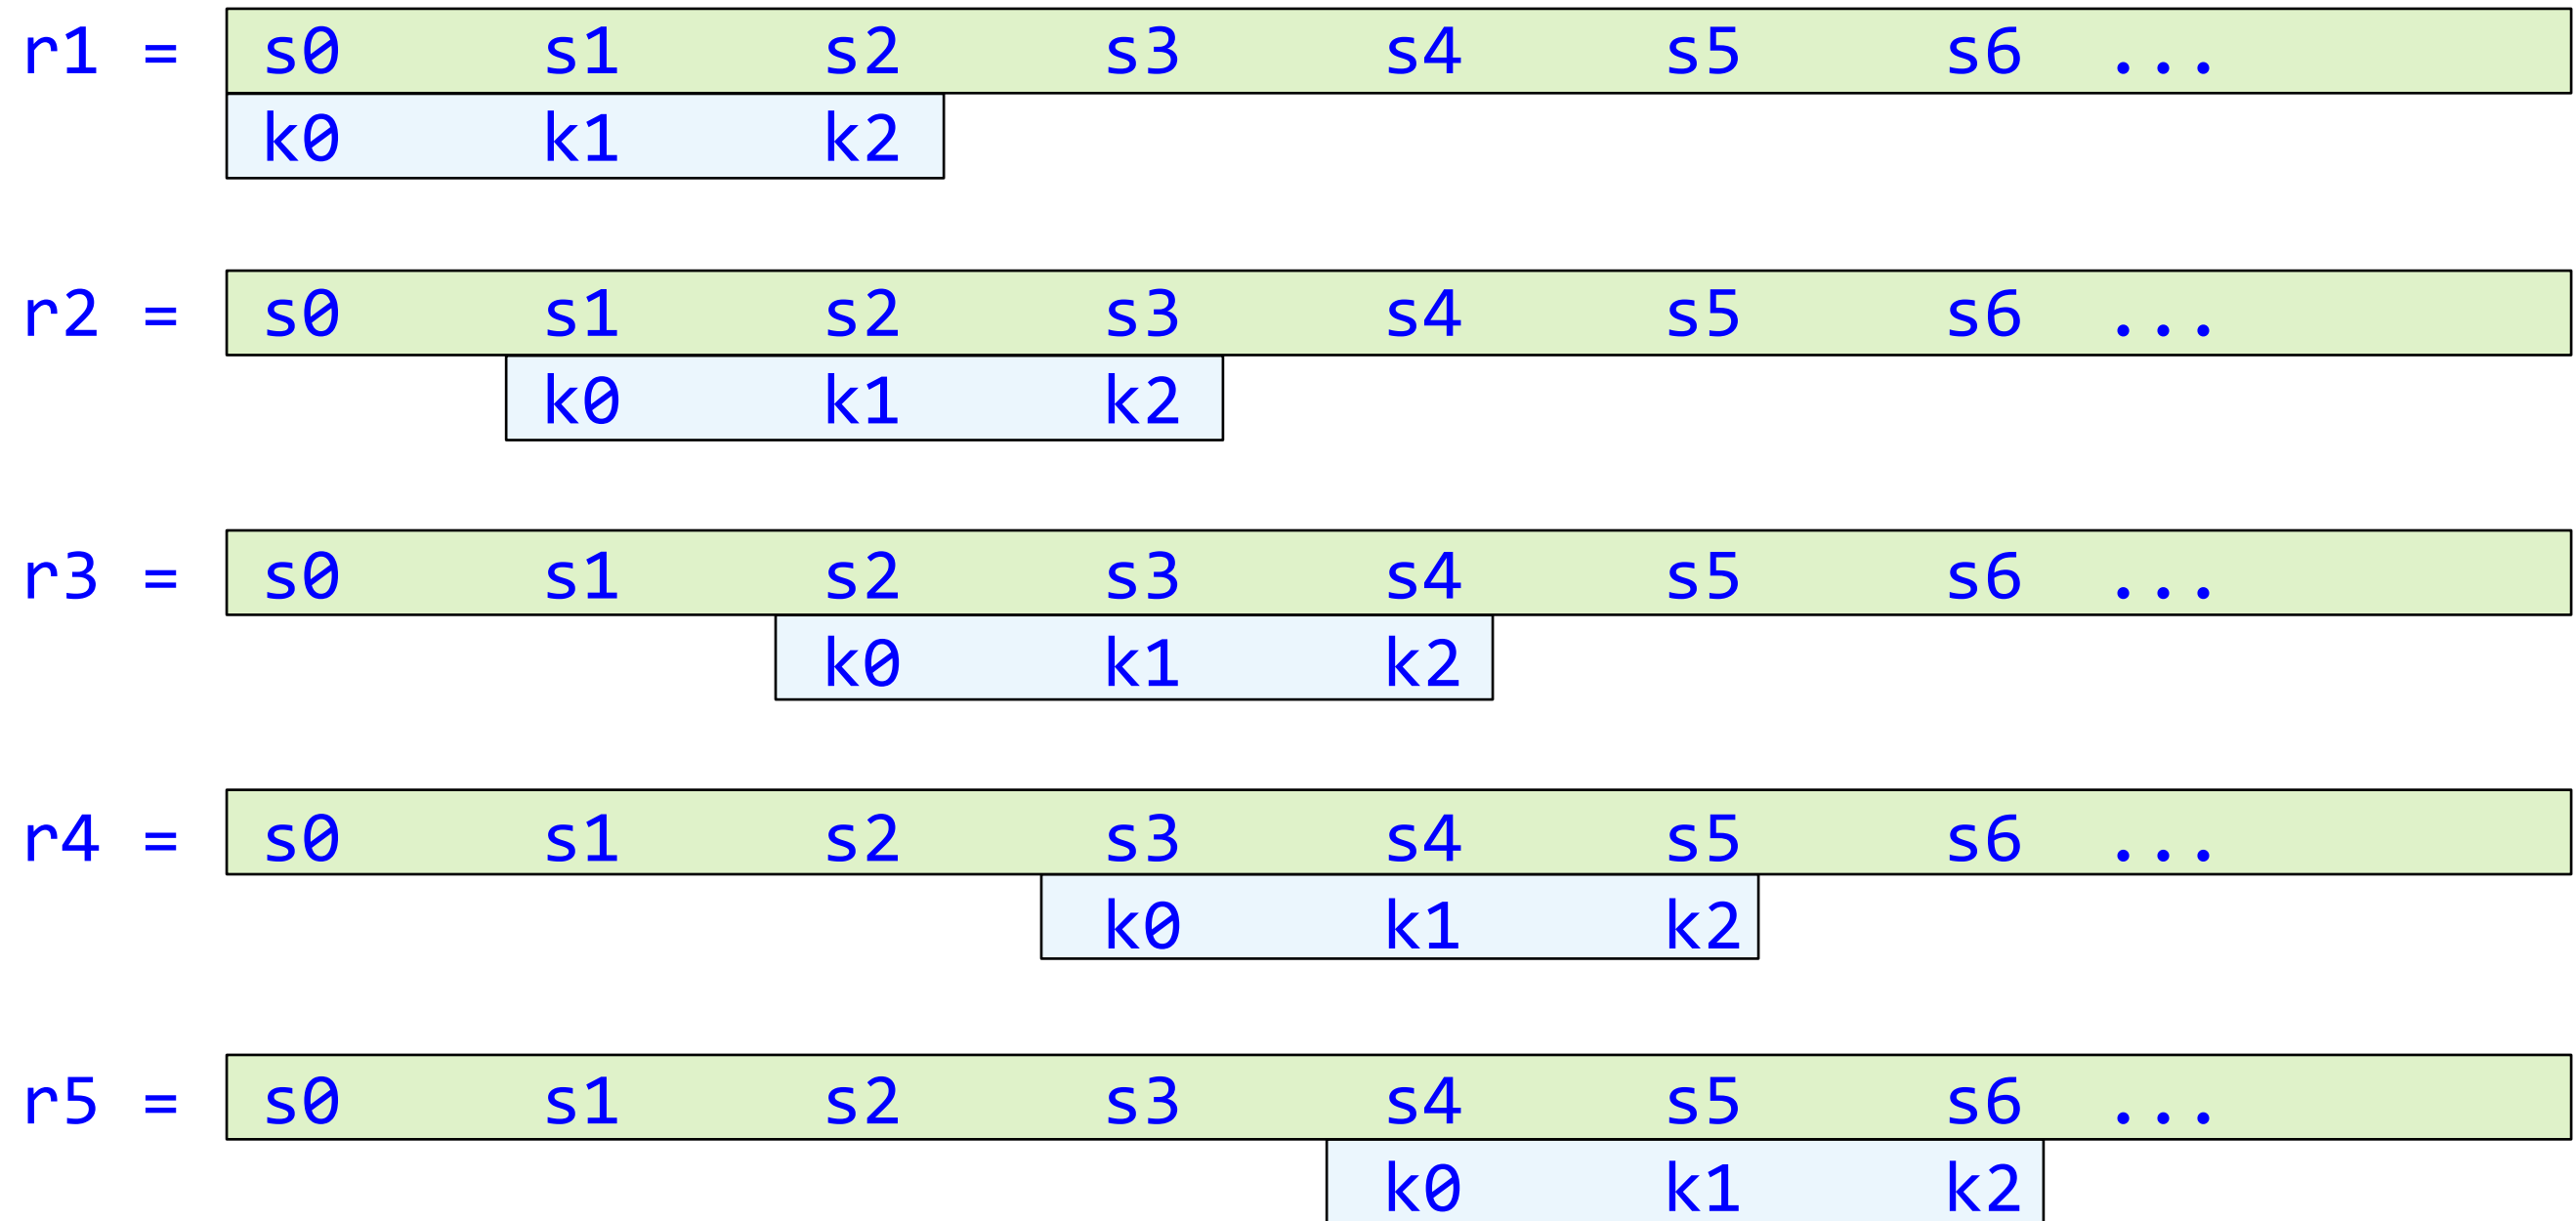
$$r_2 = \quad s_1k_0 + s_2k_1 + s_3k_2$$

$$r_3 = \quad \quad s_2k_0 + s_3k_1 + s_4k_2$$

$$r_4 = \quad \quad \quad s_3k_0 + s_4k_1 + s_5k_2$$

$$r_5 = \quad \quad \quad \quad s_4k_0 + s_5k_1 + s_6k_2$$

Convolution



Convolution

$S = s_0 \quad s_1 \quad s_2 \quad s_3 \quad s_4 \quad s_5 \quad s_6 \dots$
 $K = k_0 \quad k_1 \quad k_2$

$r_0 = \dots$

$r_1 = s_0k_0 + s_1k_1 + s_2k_2$

$r_2 = s_1k_0 + s_2k_1 + s_3k_2$

$r_3 = s_2k_0 + s_3k_1 + s_4k_2$

$r_4 = s_3k_0 + s_4k_1 + s_5k_2$

$r_5 = s_4k_0 + s_5k_1 + s_6k_2$

Convolution

$S = s_0 \quad s_1 \quad s_2 \quad s_3 \quad s_4 \quad s_5 \quad s_6 \dots$
 $K = k_0 \quad k_1 \quad k_2$

$r_0 = \dots$

$r_1 = s_0k_0 + s_1k_1 + s_2k_2$
 $r_2 = s_1k_0 + s_2k_1 + s_3k_2$
 $r_3 = s_2k_0 + s_3k_1 + s_4k_2$
 $r_4 = s_3k_0 + s_4k_1 + s_5k_2$
 $r_5 = s_4k_0 + s_5k_1 + s_6k_2$

Convolution

s0	s1	s2	s3	s4	s5	s6	s7
k0	k0	k0	k0	k0	k0	k0	k0

s1	s2	s3	s4	s5	s6	s7	s8
k1	k1	k1	k1	k1	k1	k1	k1

s2	s3	s4	s5	s6	s7	s8	s9
k2	k2	k2	k2	k2	k2	k2	k2

r1	r2	r3	r4	r5	r6	r7	r8
----	----	----	----	----	----	----	----

Function Template avx_convolve()

```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    //- The convolution kernel must have non-negative size and fit with a single register.
    //
    static_assert(KernelSize > 1 && KernelSize <= 16);

    //- The index of the kernel center must be valid.
    //
    static_assert(KernelCenter >= 0 && KernelCenter < KernelSize);

    //- Convolution flips the kernel, so the kernel center must be adjusted.
    //
    constexpr int WindowCenter = KernelSize - KernelCenter - 1;

    rf_512 prev;    //- Bottom of the input data window
    rf_512 curr;    //- Middle of the input data windows
    rf_512 next;    //- Top of the input data window
    rf_512 lo;      //- Primary work data register, used to multiply kernel coefficients
    rf_512 hi;      //- Upper work data register, supplies values to the top of 'lo'
    rf_512 sum;     //- Accumulated value

    ...
}
```

Function Template avx_convolve()

```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    rf_512 kcoeff[KernelSize];    //- Coefficients of the convolution kernel

    //- Broadcast each kernel coefficient into its own register, to be used later in the FMA call.
    //
    for (int i = 0, j = KernelSize - 1; i < KernelSize; ++i, --j)
    {
        kcoeff[i] = load_value(pkrnl[j]);
    }

    //- Preload the initial input data window; note the zeroes in the register representing data
    // preceding the input array.
    //
    prev = load_value(0.0f);
    curr = load_from(psrc);
    next = load_from(psrc + 16);

    ...
}
```

Function Template avx_convolve()


```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    rf_512 kcoeff[KernelSize];    //- Coefficients of the convolution kernel

    //- Broadcast each kernel coefficient into its own register, to be used later in the FMA call.
    //
    for (int i = 0, j = KernelSize - 1; i < KernelSize; ++i, --j)
    {
        kcoeff[i] = load_value(pkrnl[j]);
    }

    //- Preload the initial input data window; note the zeroes in the register representing data
    // preceding the input array.
    //
    prev = load_value(0.0f);
    curr = load_from(psrc);
    next = load_from(psrc + 16);

    ...
}
```

Function Template `avx_convolve()`



Function Template avx_convolve()

```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi);
        }
        store_to(pdst, sum);
    }
}
```

Function Template avx_convolve()


```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi); // - Slide the input data down by 1
        }
        store_to(pdst, sum);
    }
}
```

Function Template avx_convolve()



Function Template avx_convolve()


```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi);
        }
        store_to(pdst, sum);
    }
}
```

Function Template `avx_convolve()`



Function Template avx_convolve()

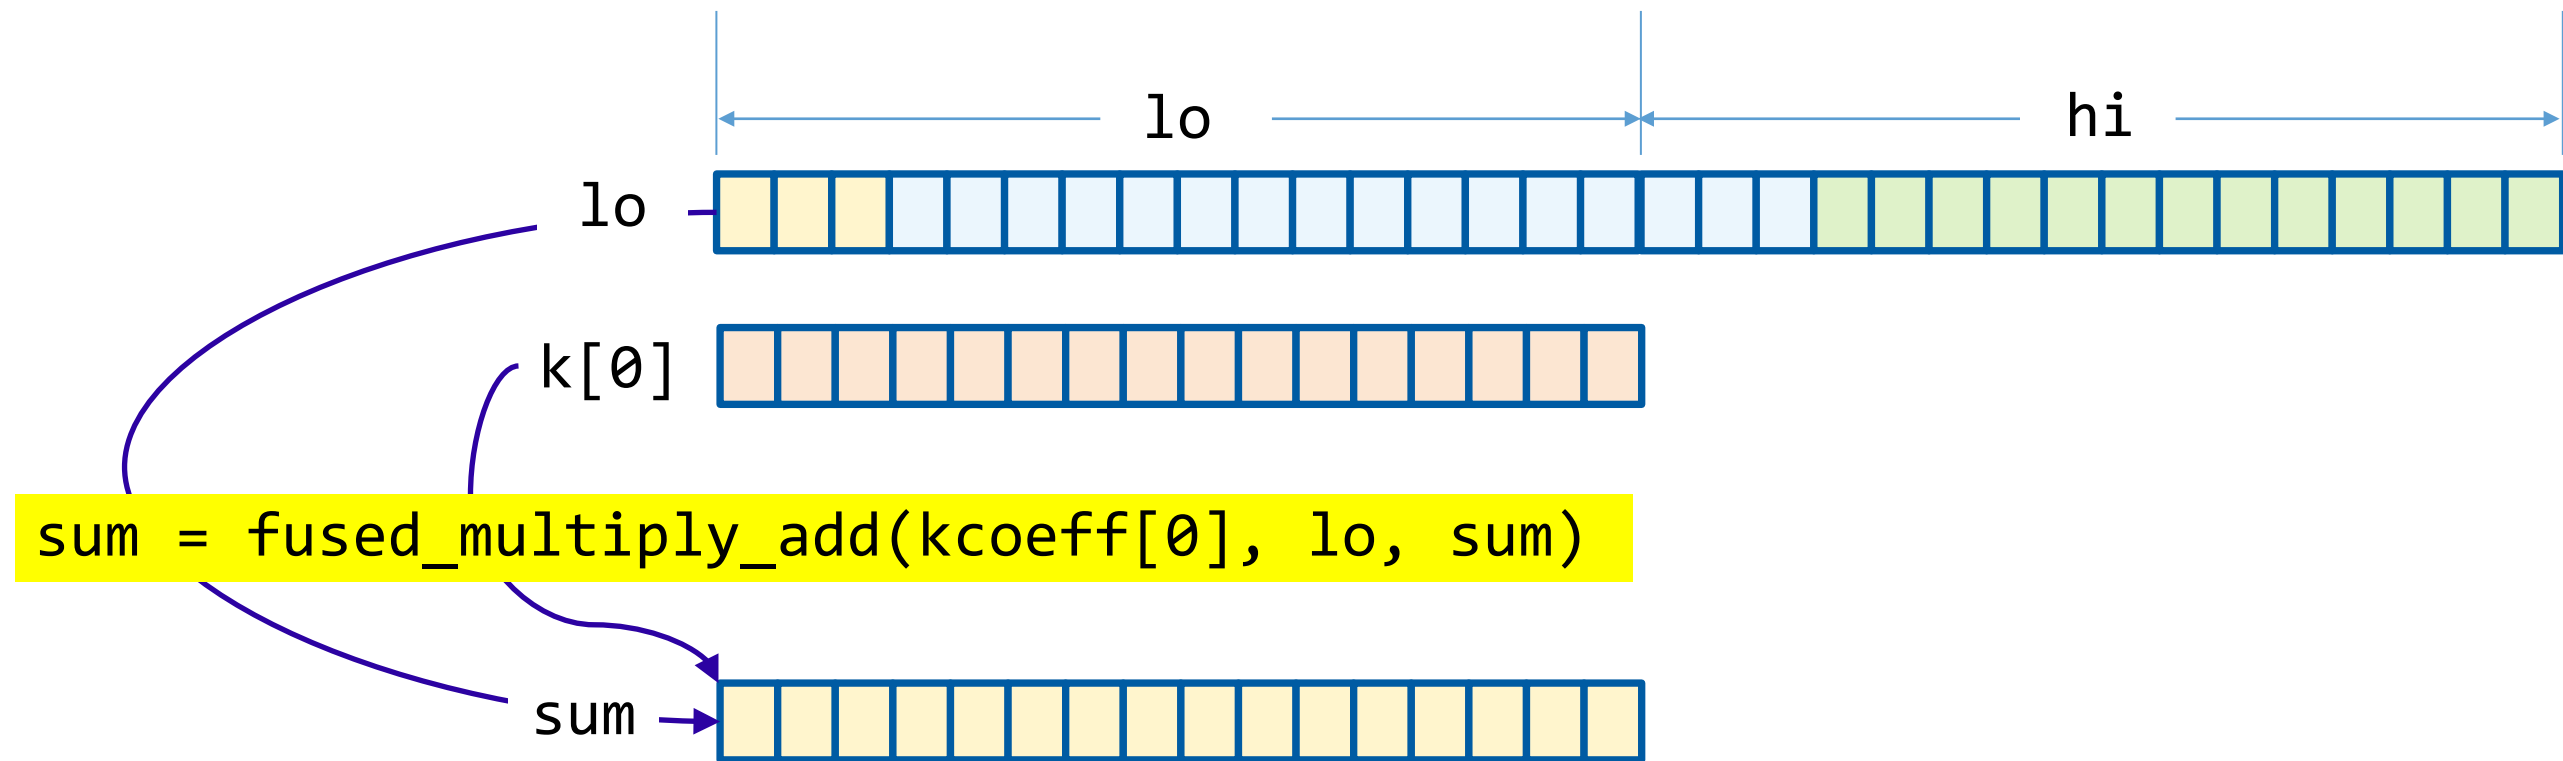
```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi);
        }
        store_to(pdst, sum);
    }
}
```

Function Template avx_convolve()



Function Template avx_convolve()

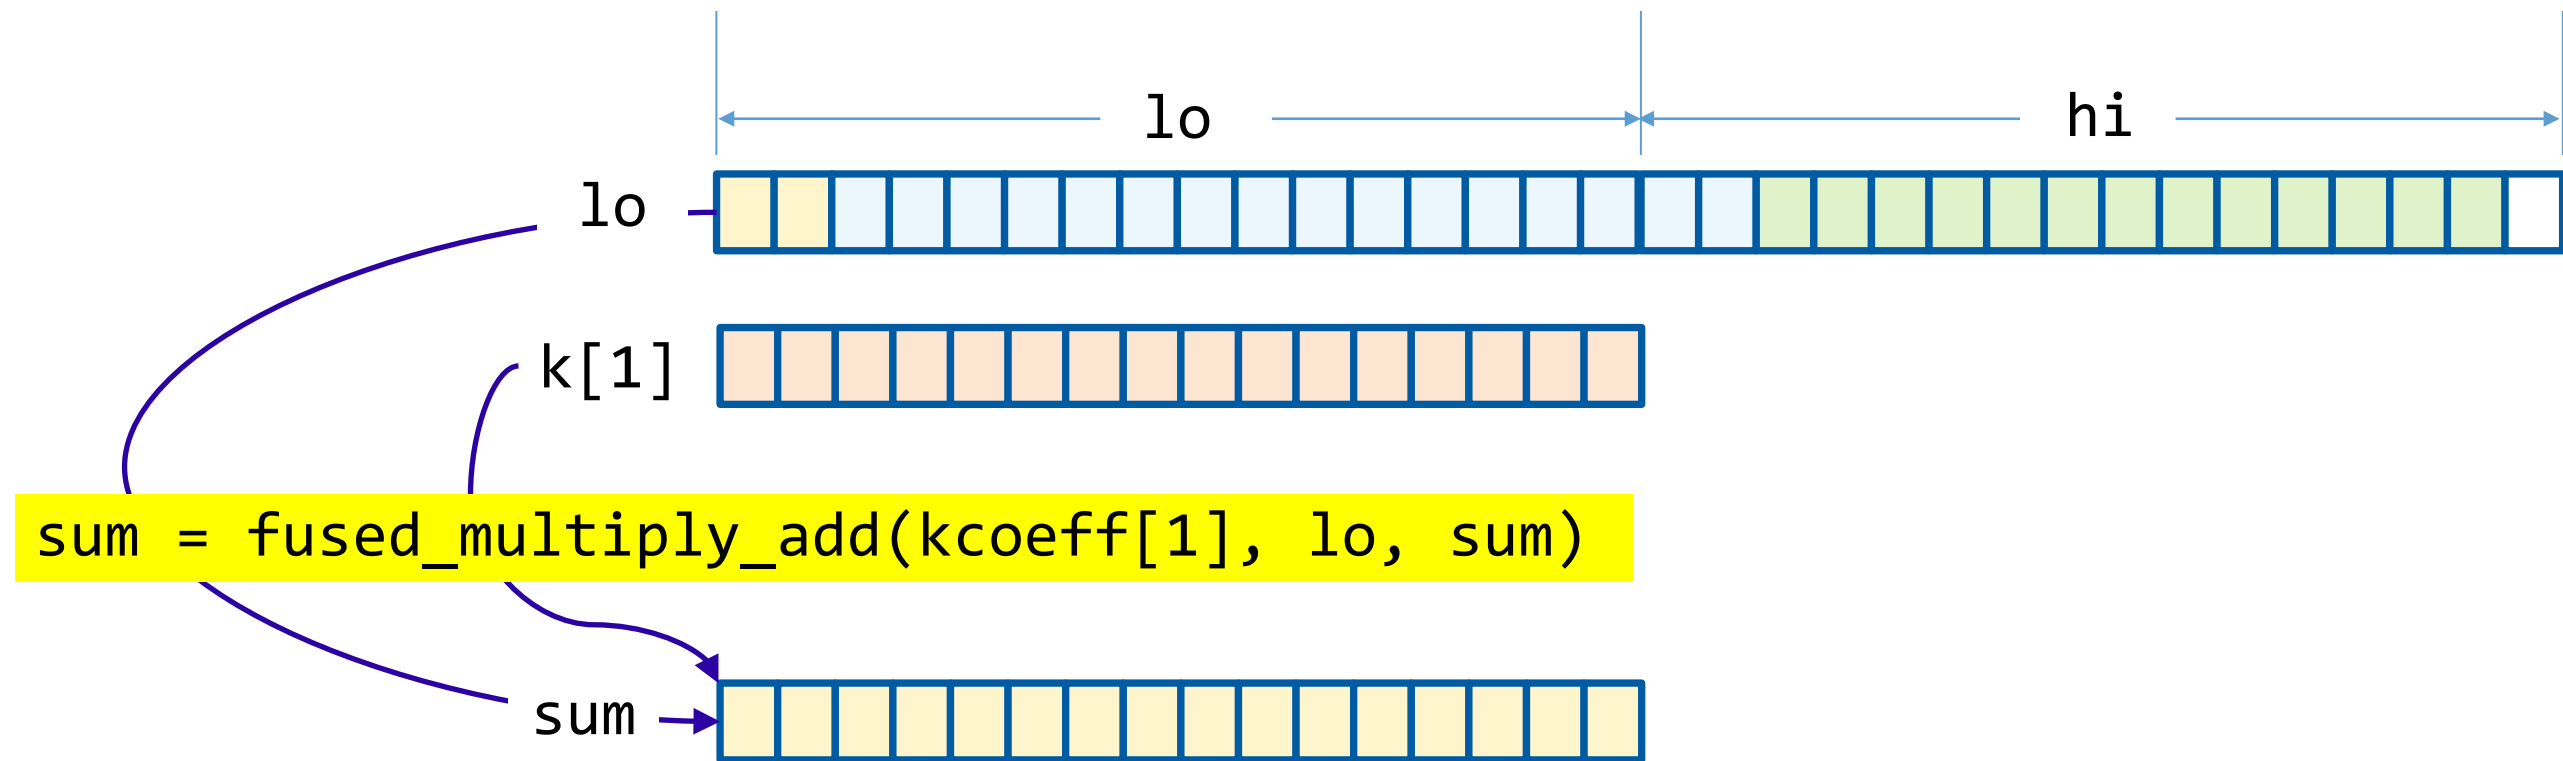
```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

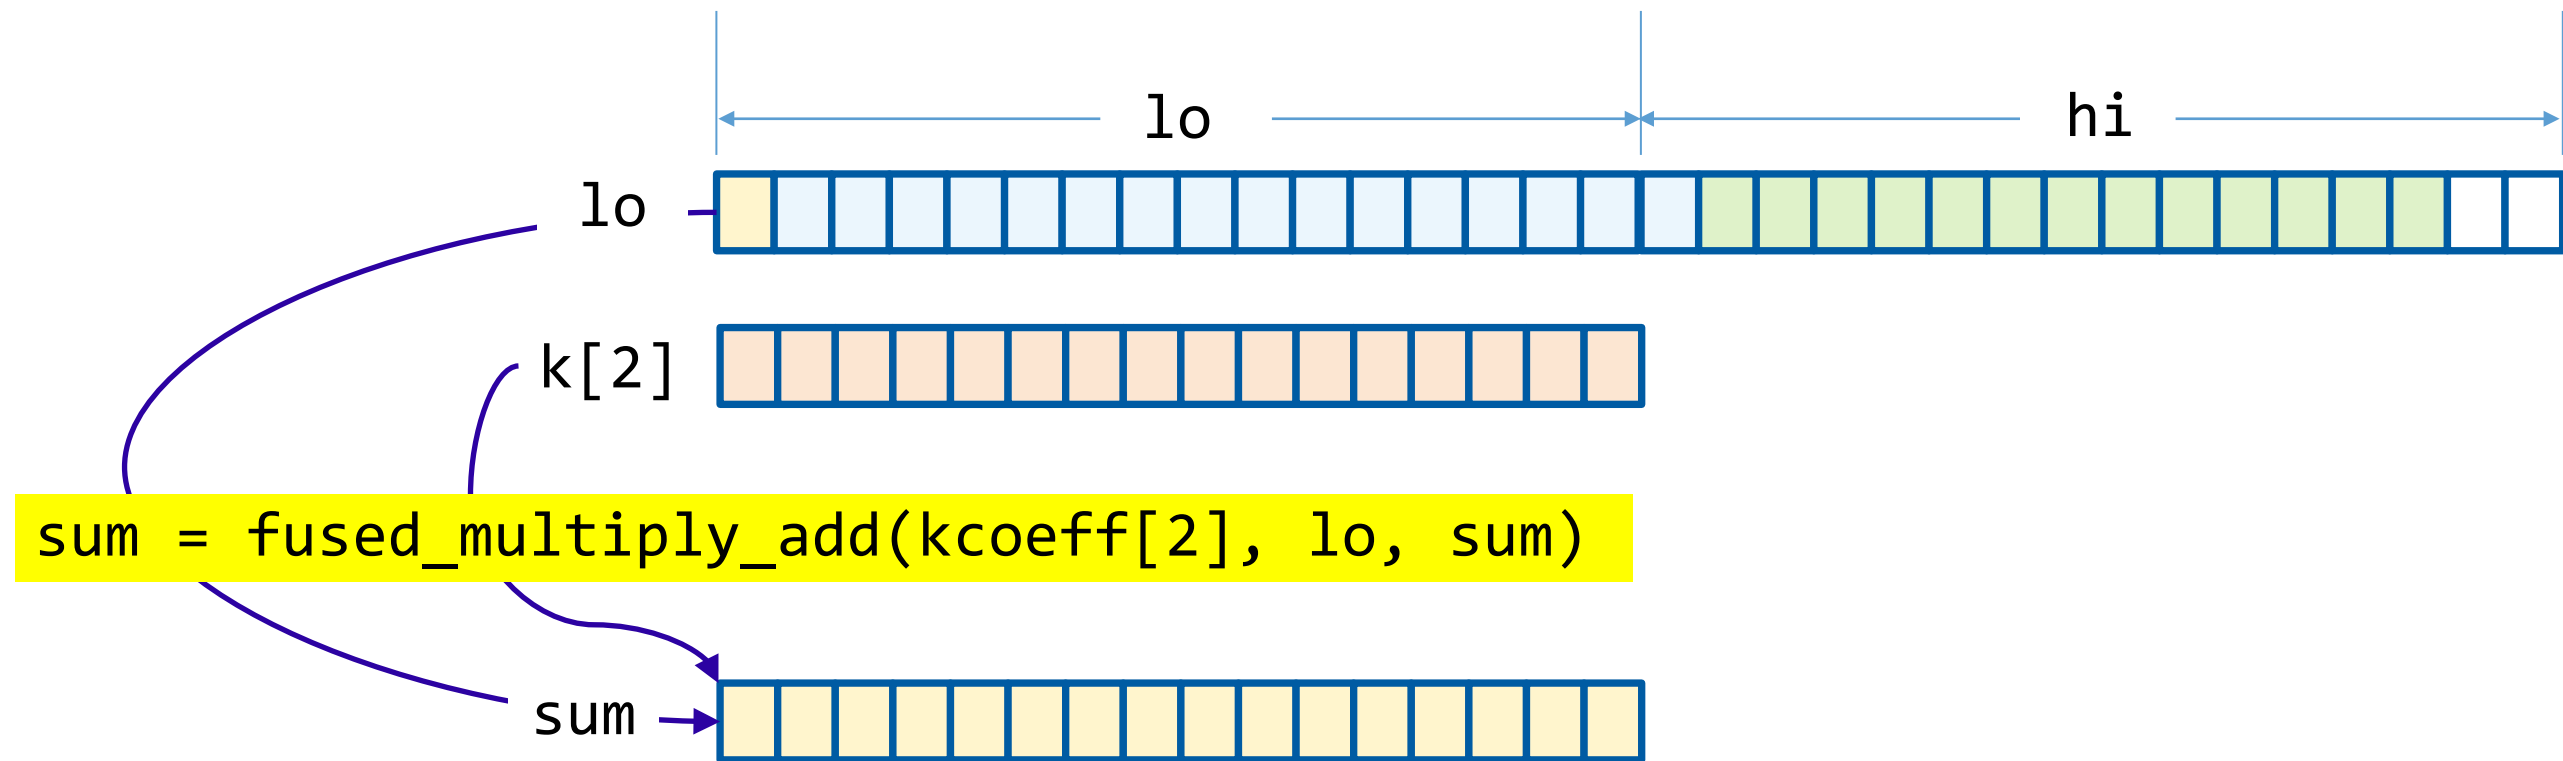
        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi);
        }
        store_to(pdst, sum);
    }
}
```

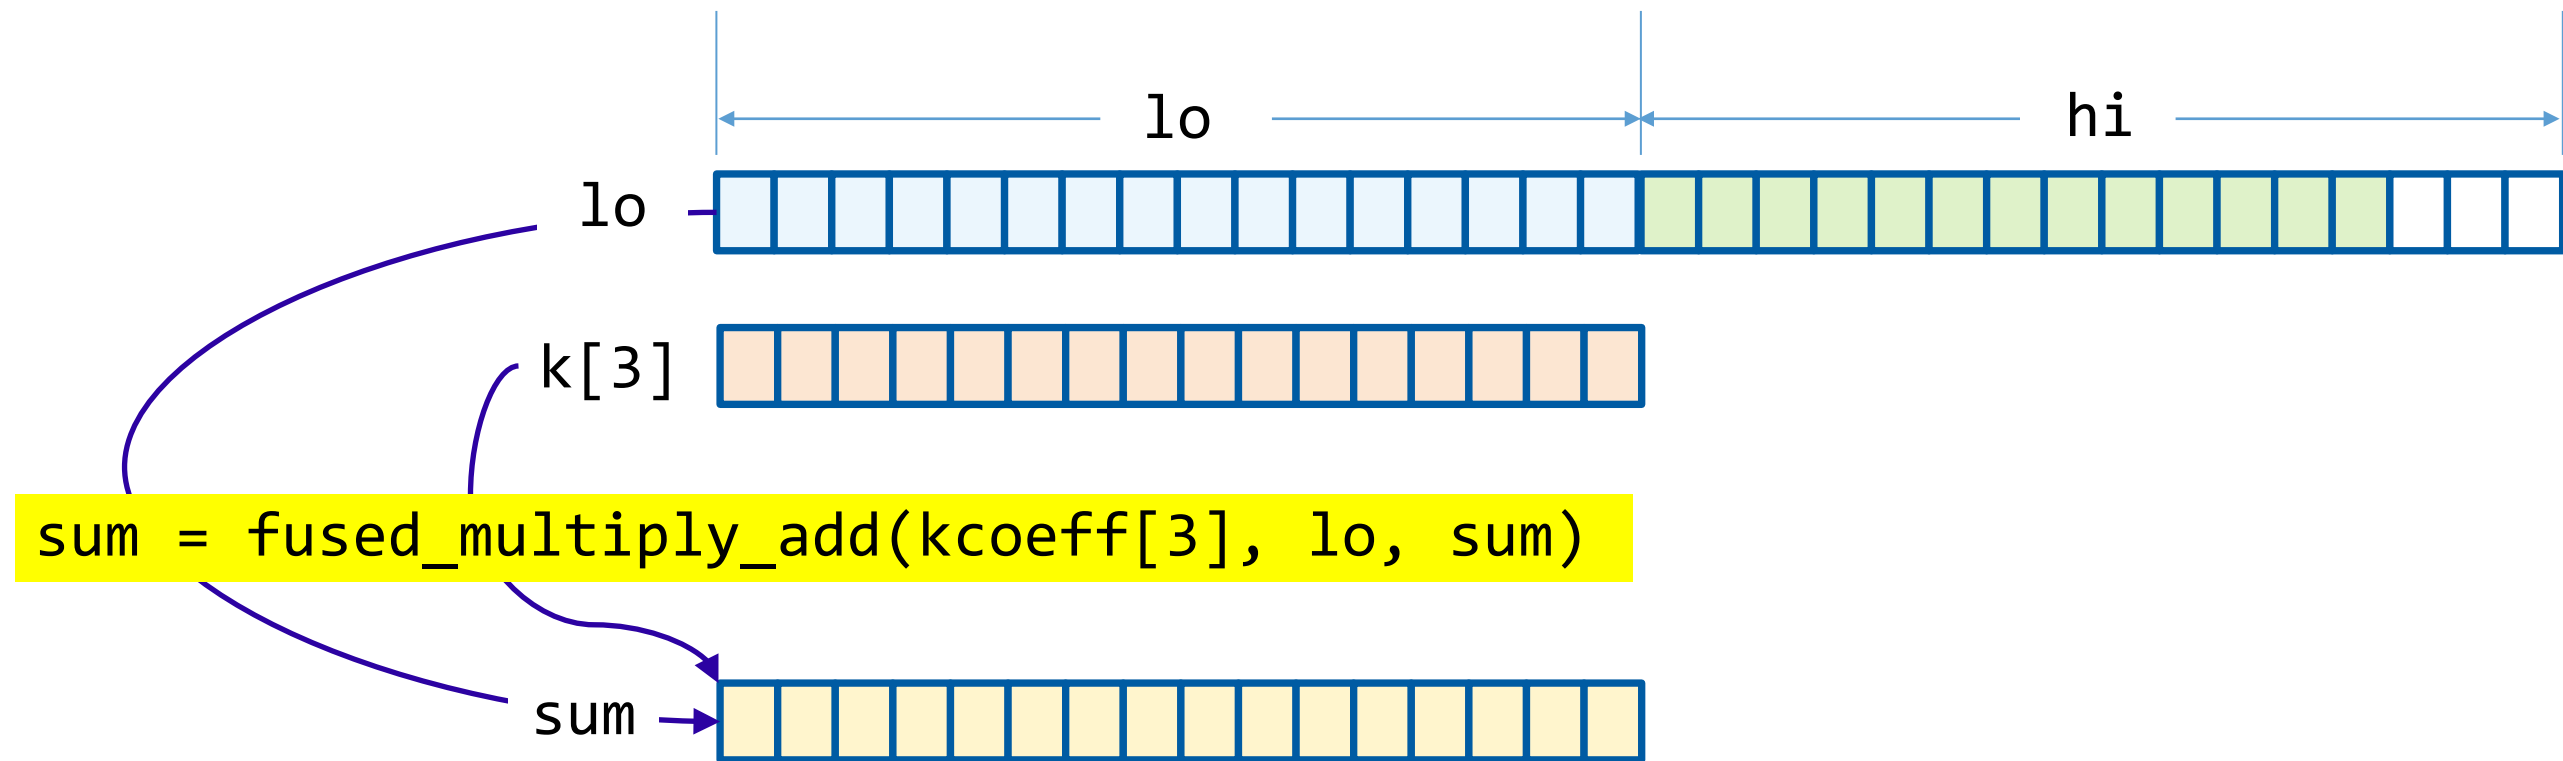
Function Template avx_convolve()



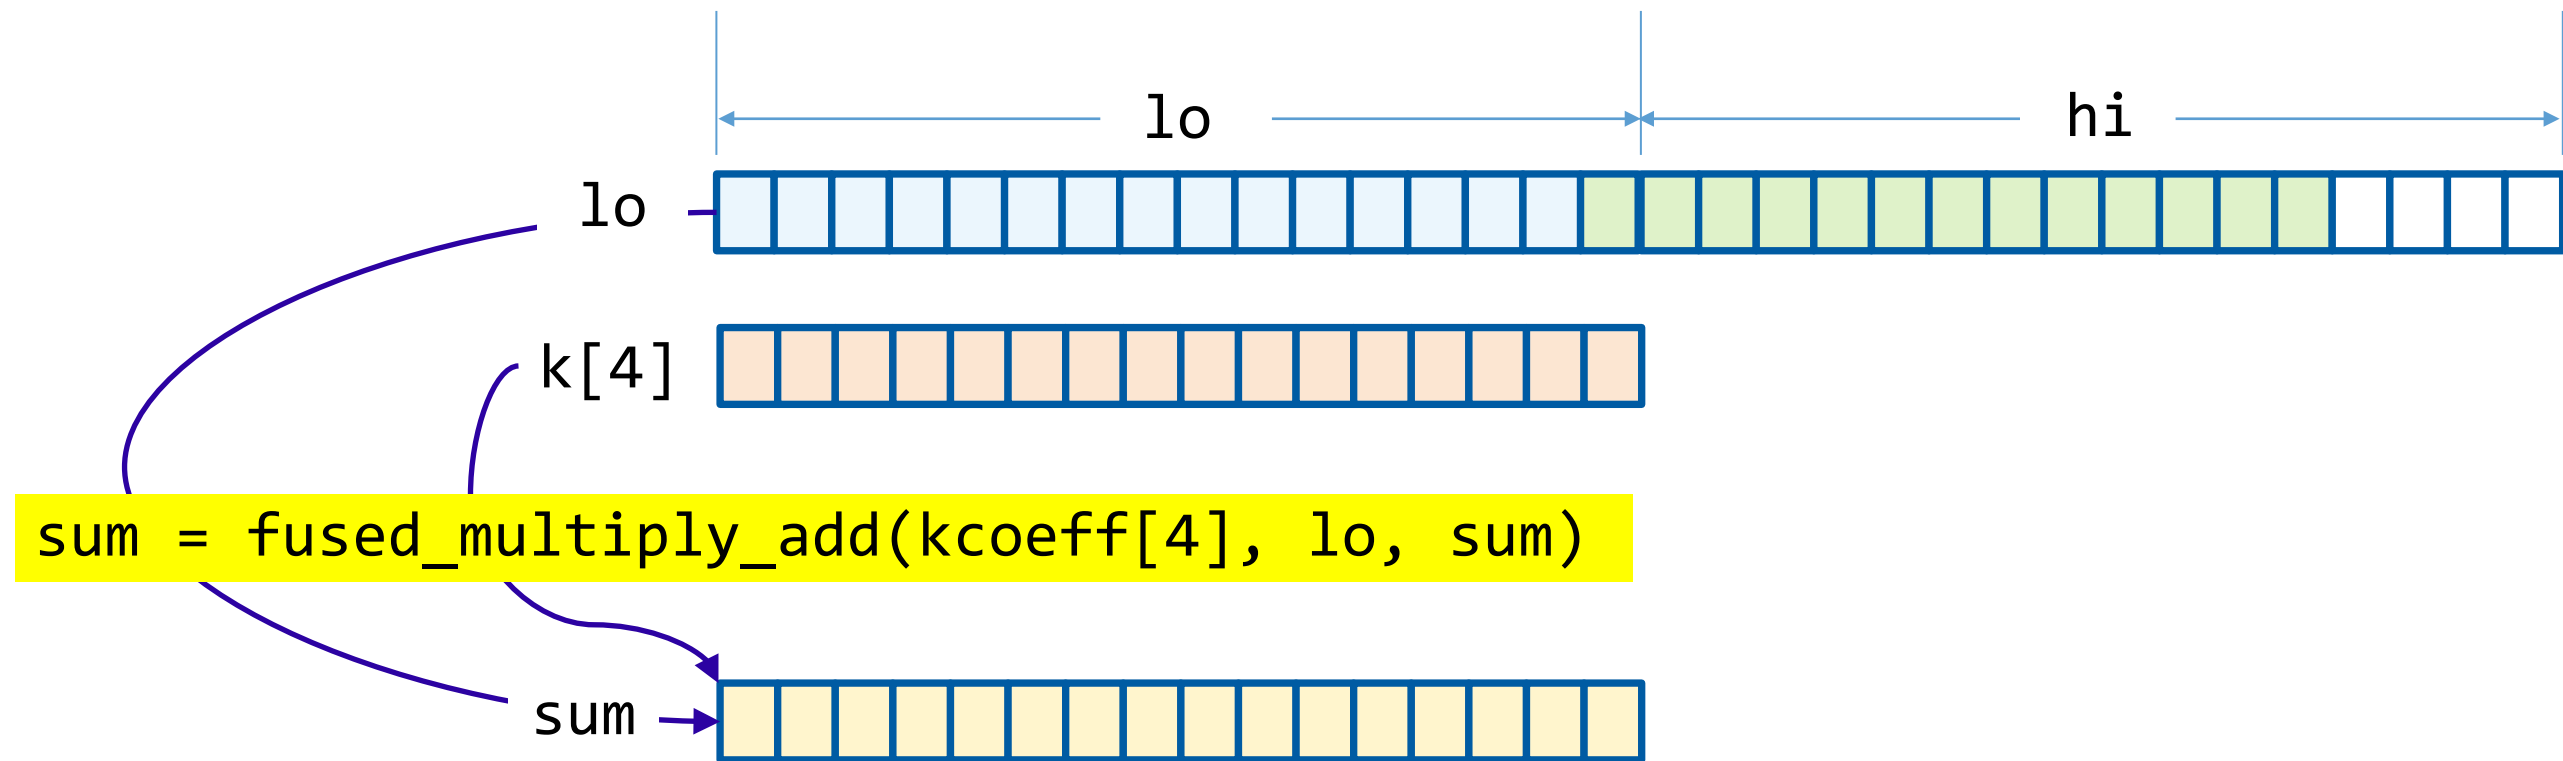
Function Template avx_convolve()



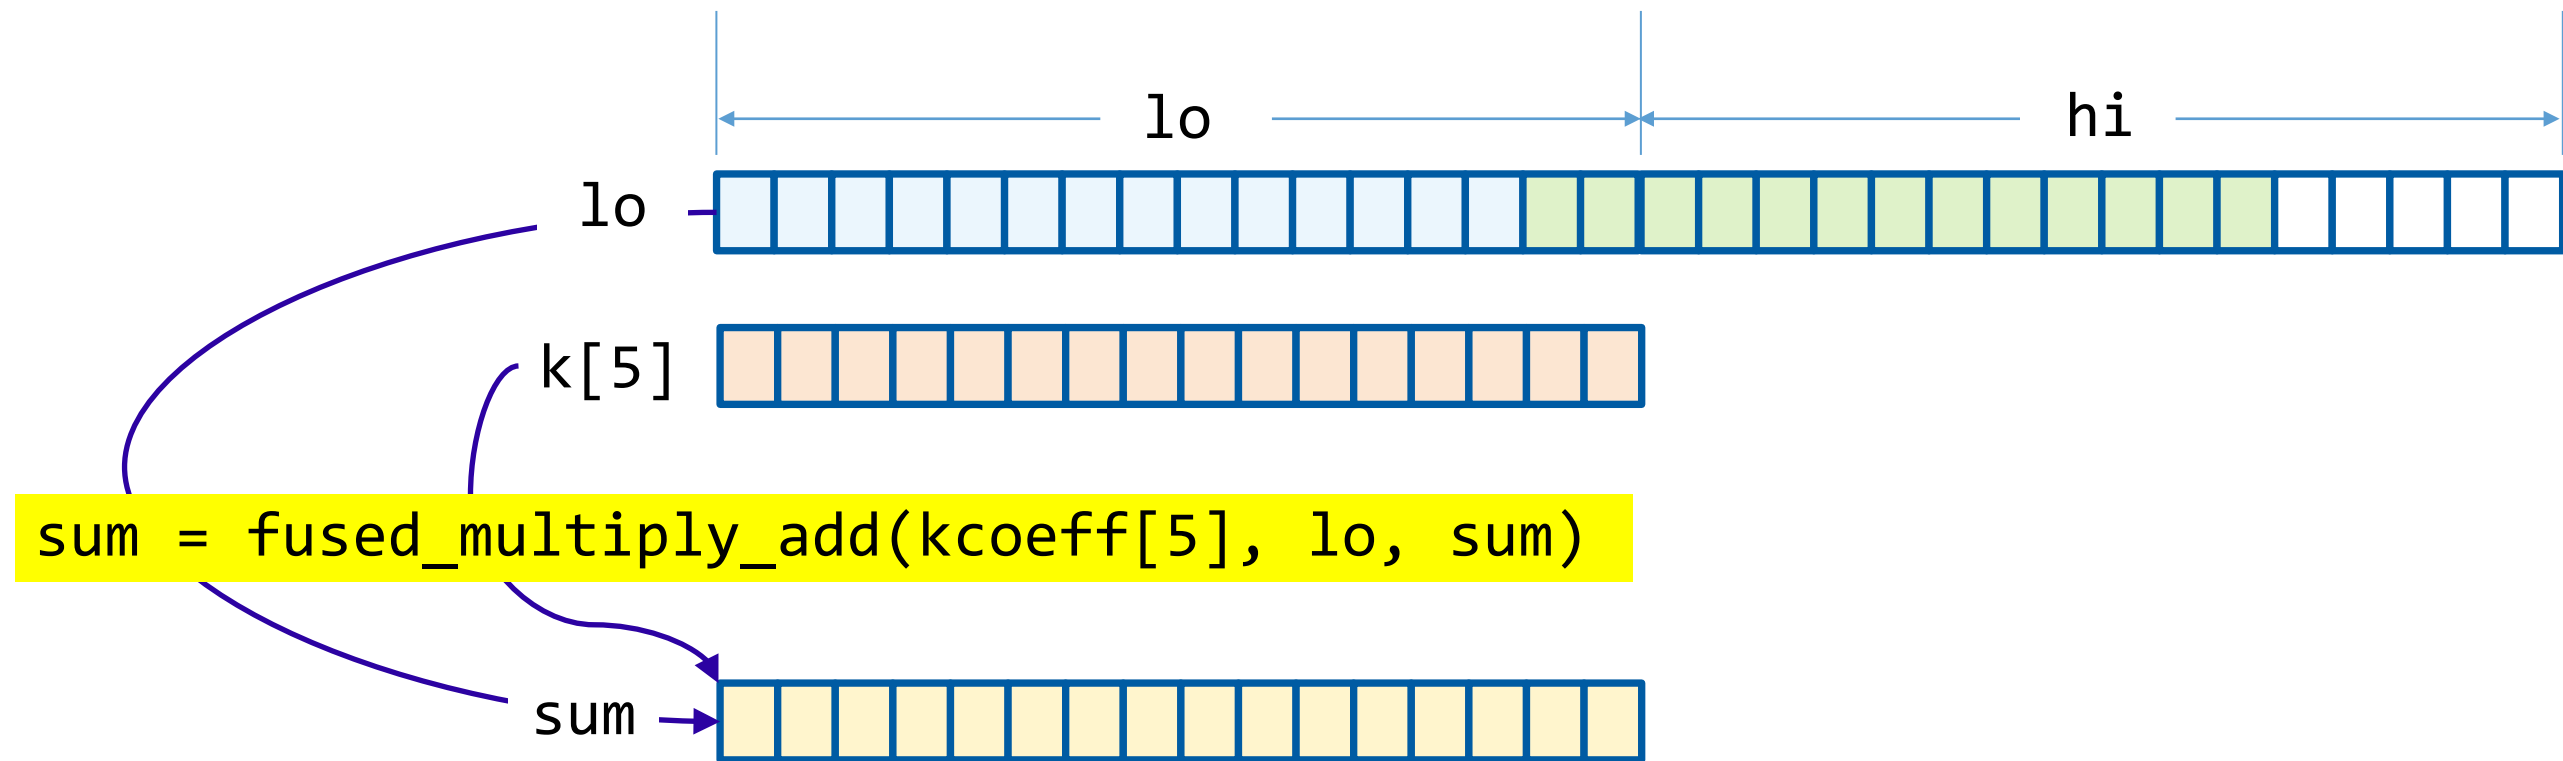
Function Template avx_convolve()



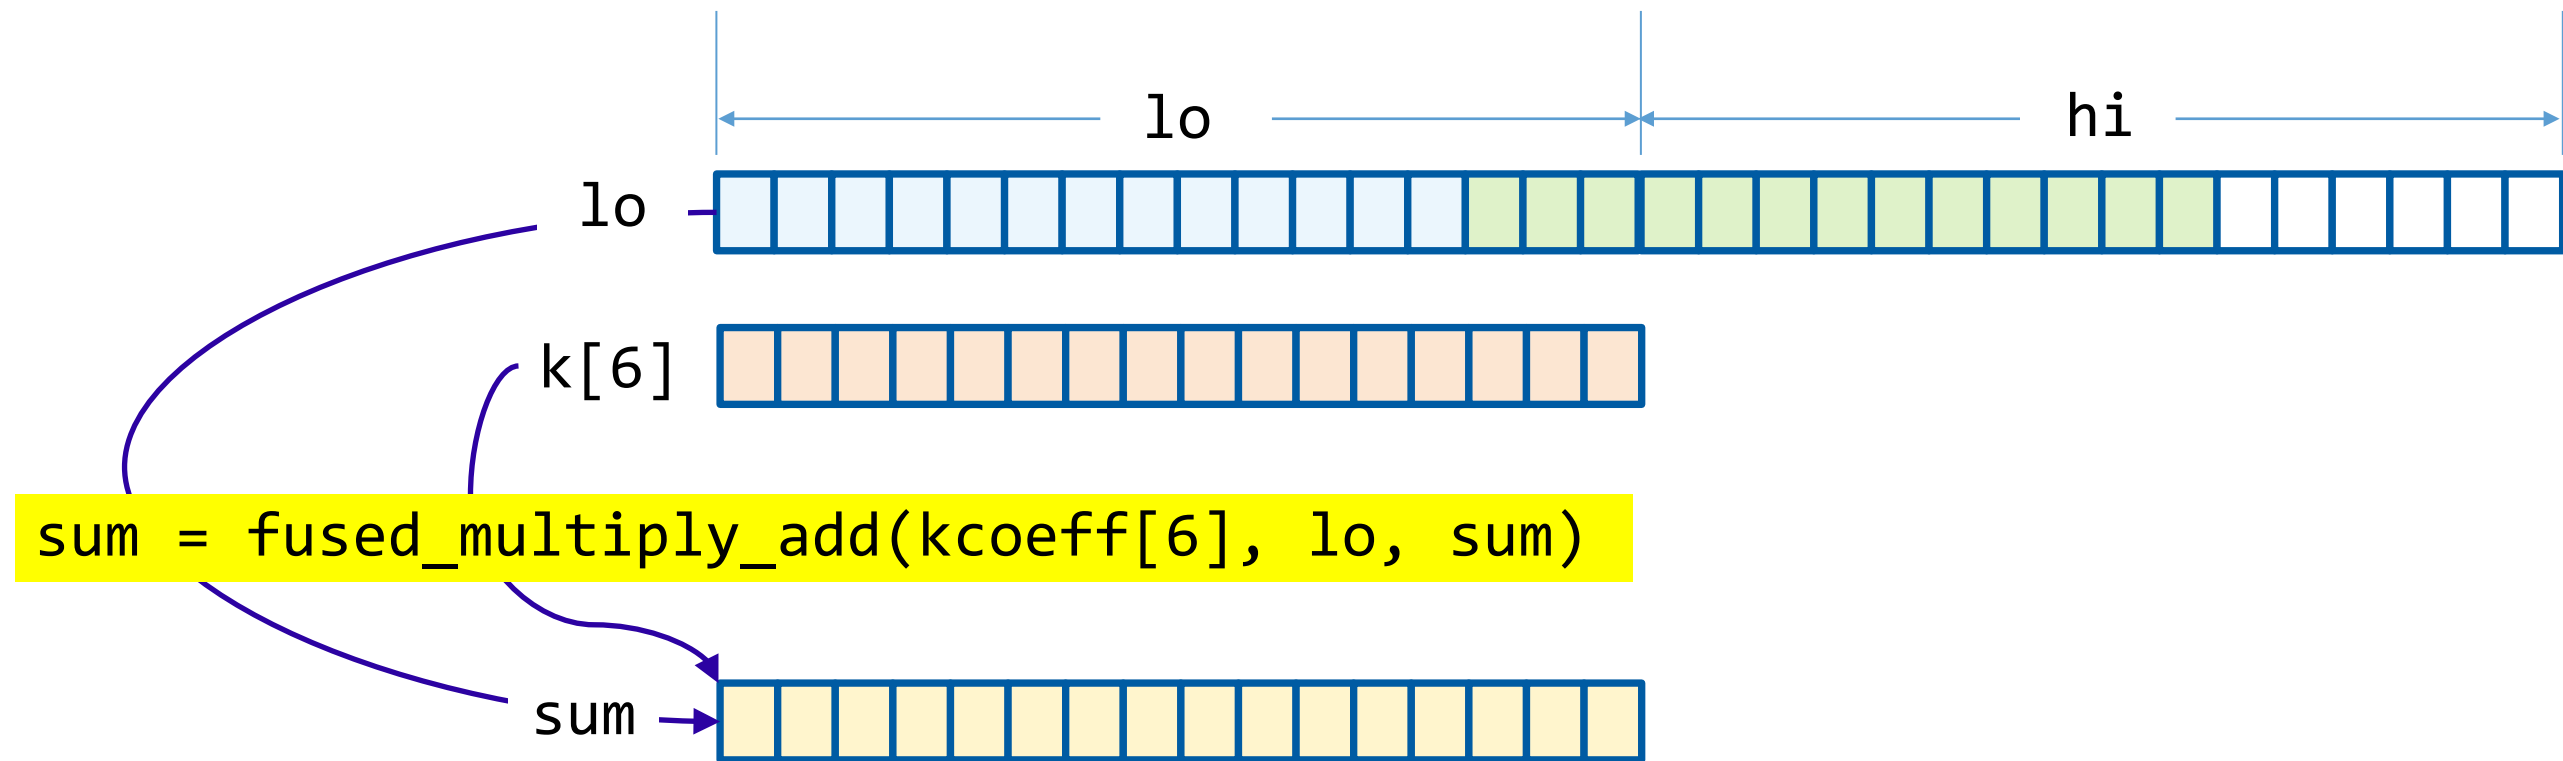
Function Template avx_convolve()



Function Template avx_convolve()



Function Template avx_convolve()



Function Template avx_convolve()

```
template<int KernelSize, int KernelCenter> void
avx_convolve(float* pdst, float const* pkrnl, float const* psrc, size_t len)
{
    ...
    for (auto pEnd = psrc + len - 16; psrc < pEnd; psrc += 16, pdst += 16)
    {
        sum = load_value(0.0f); // - Init the accumulator

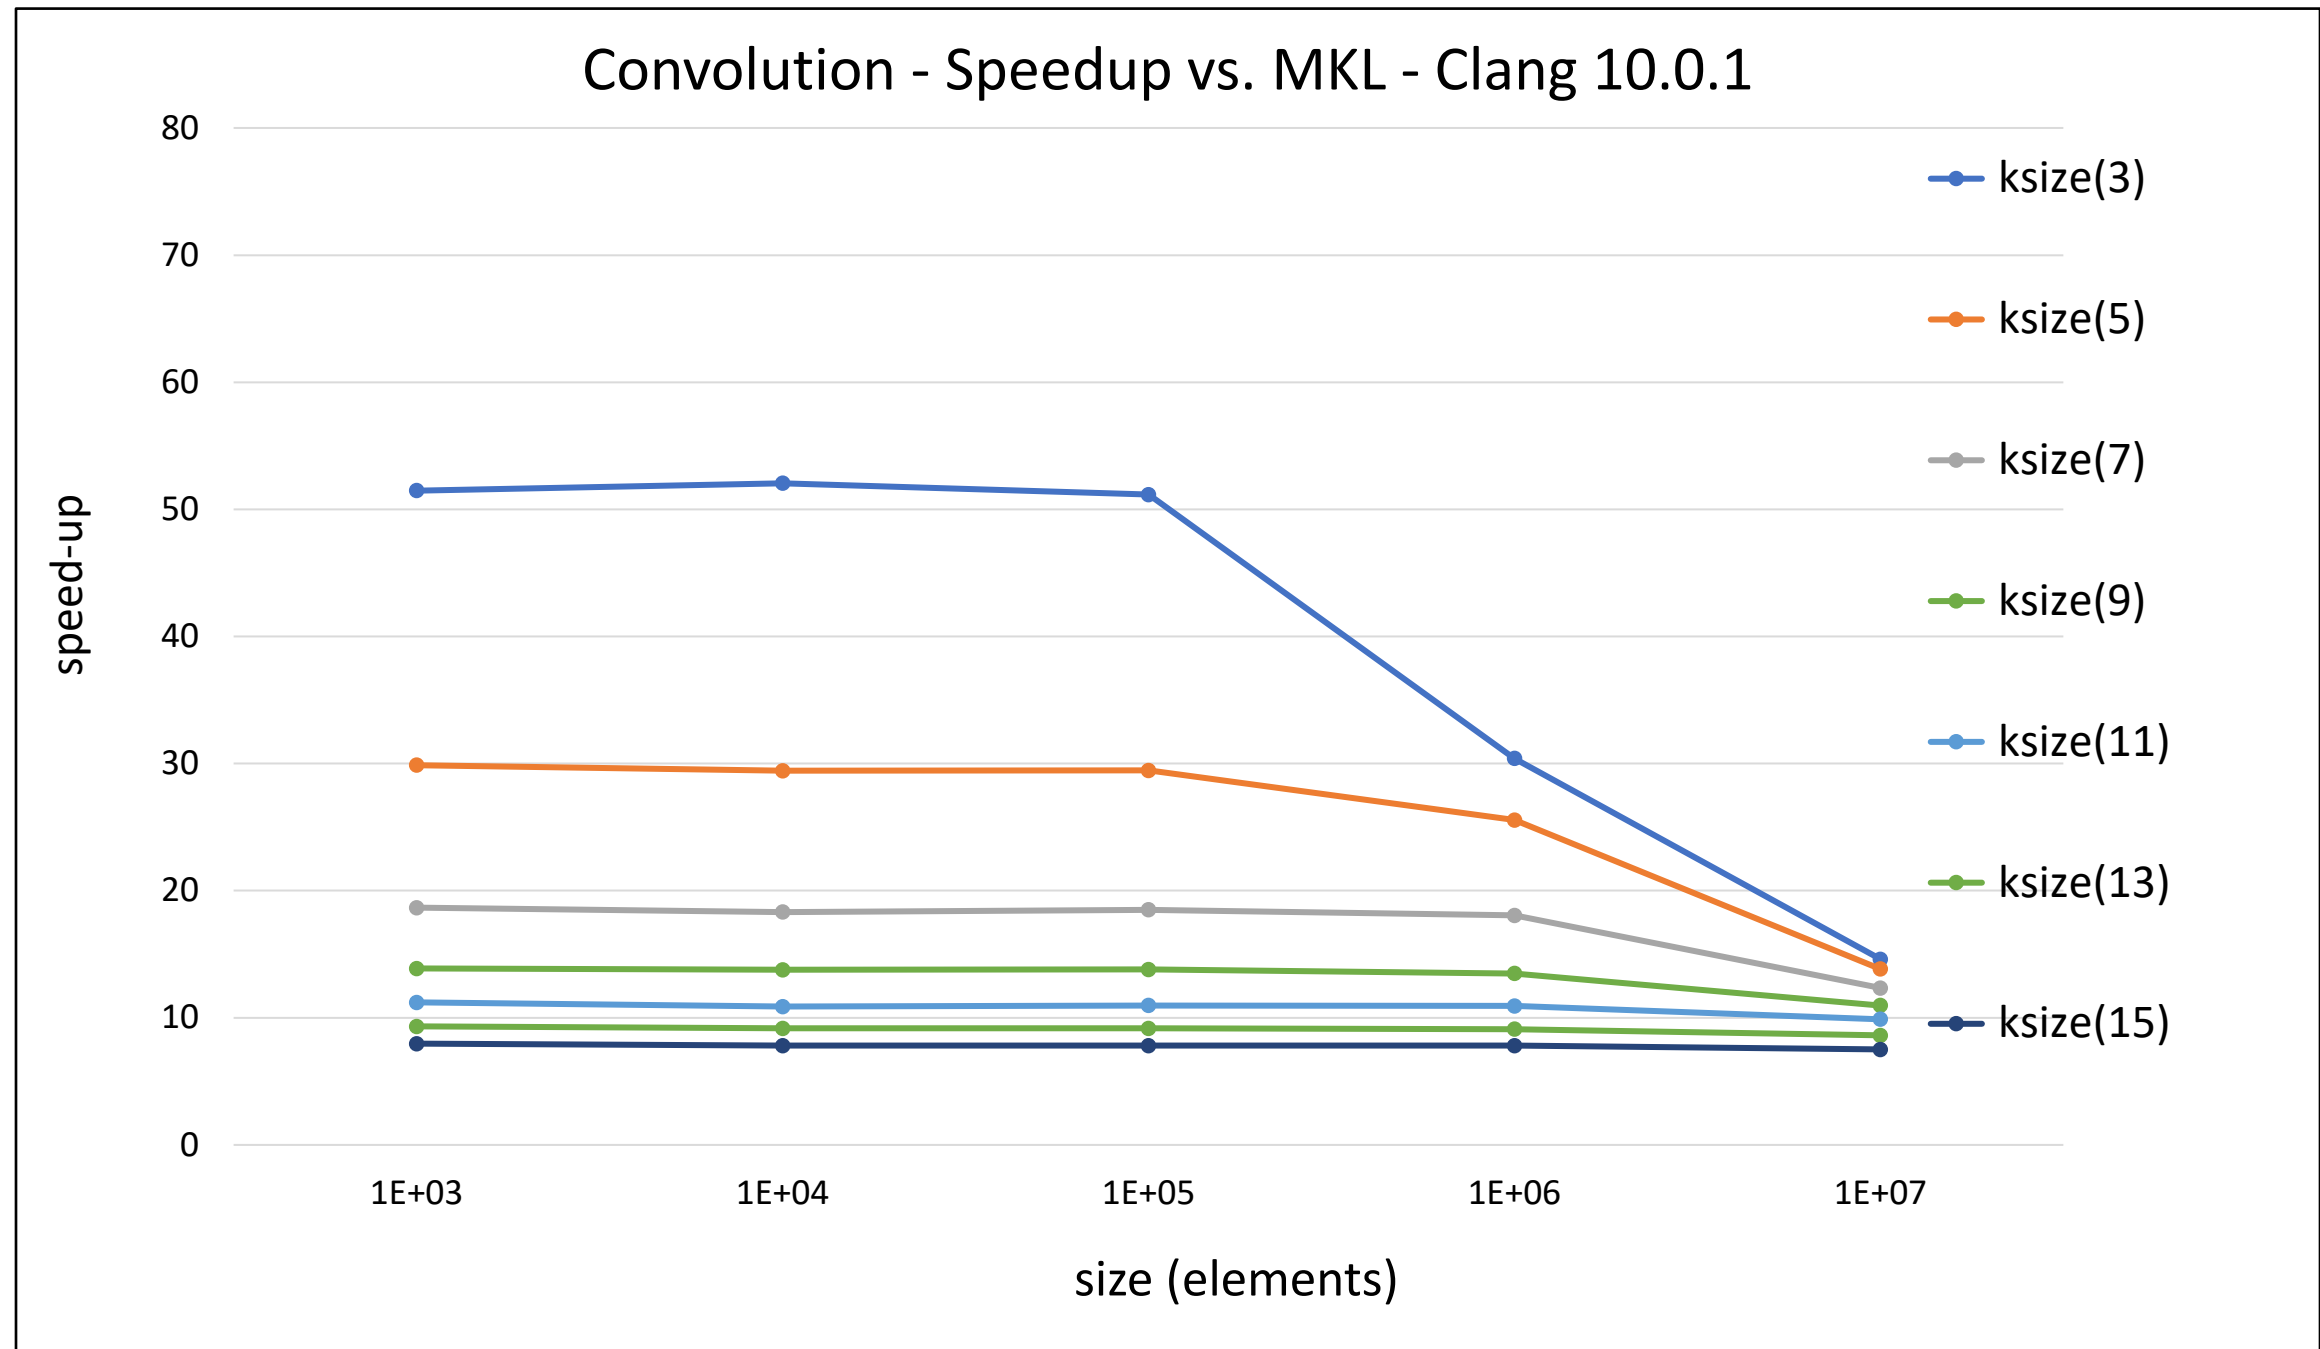
        lo = shift_up_with_carry<WindowCenter>(prev, curr); // - Init the work data registers
        hi = shift_up_with_carry<WindowCenter>(curr, next);

        prev = curr; // - Slide the input data window up
        curr = next; // by a register's work of values
        next = load_from(psrc + 32);

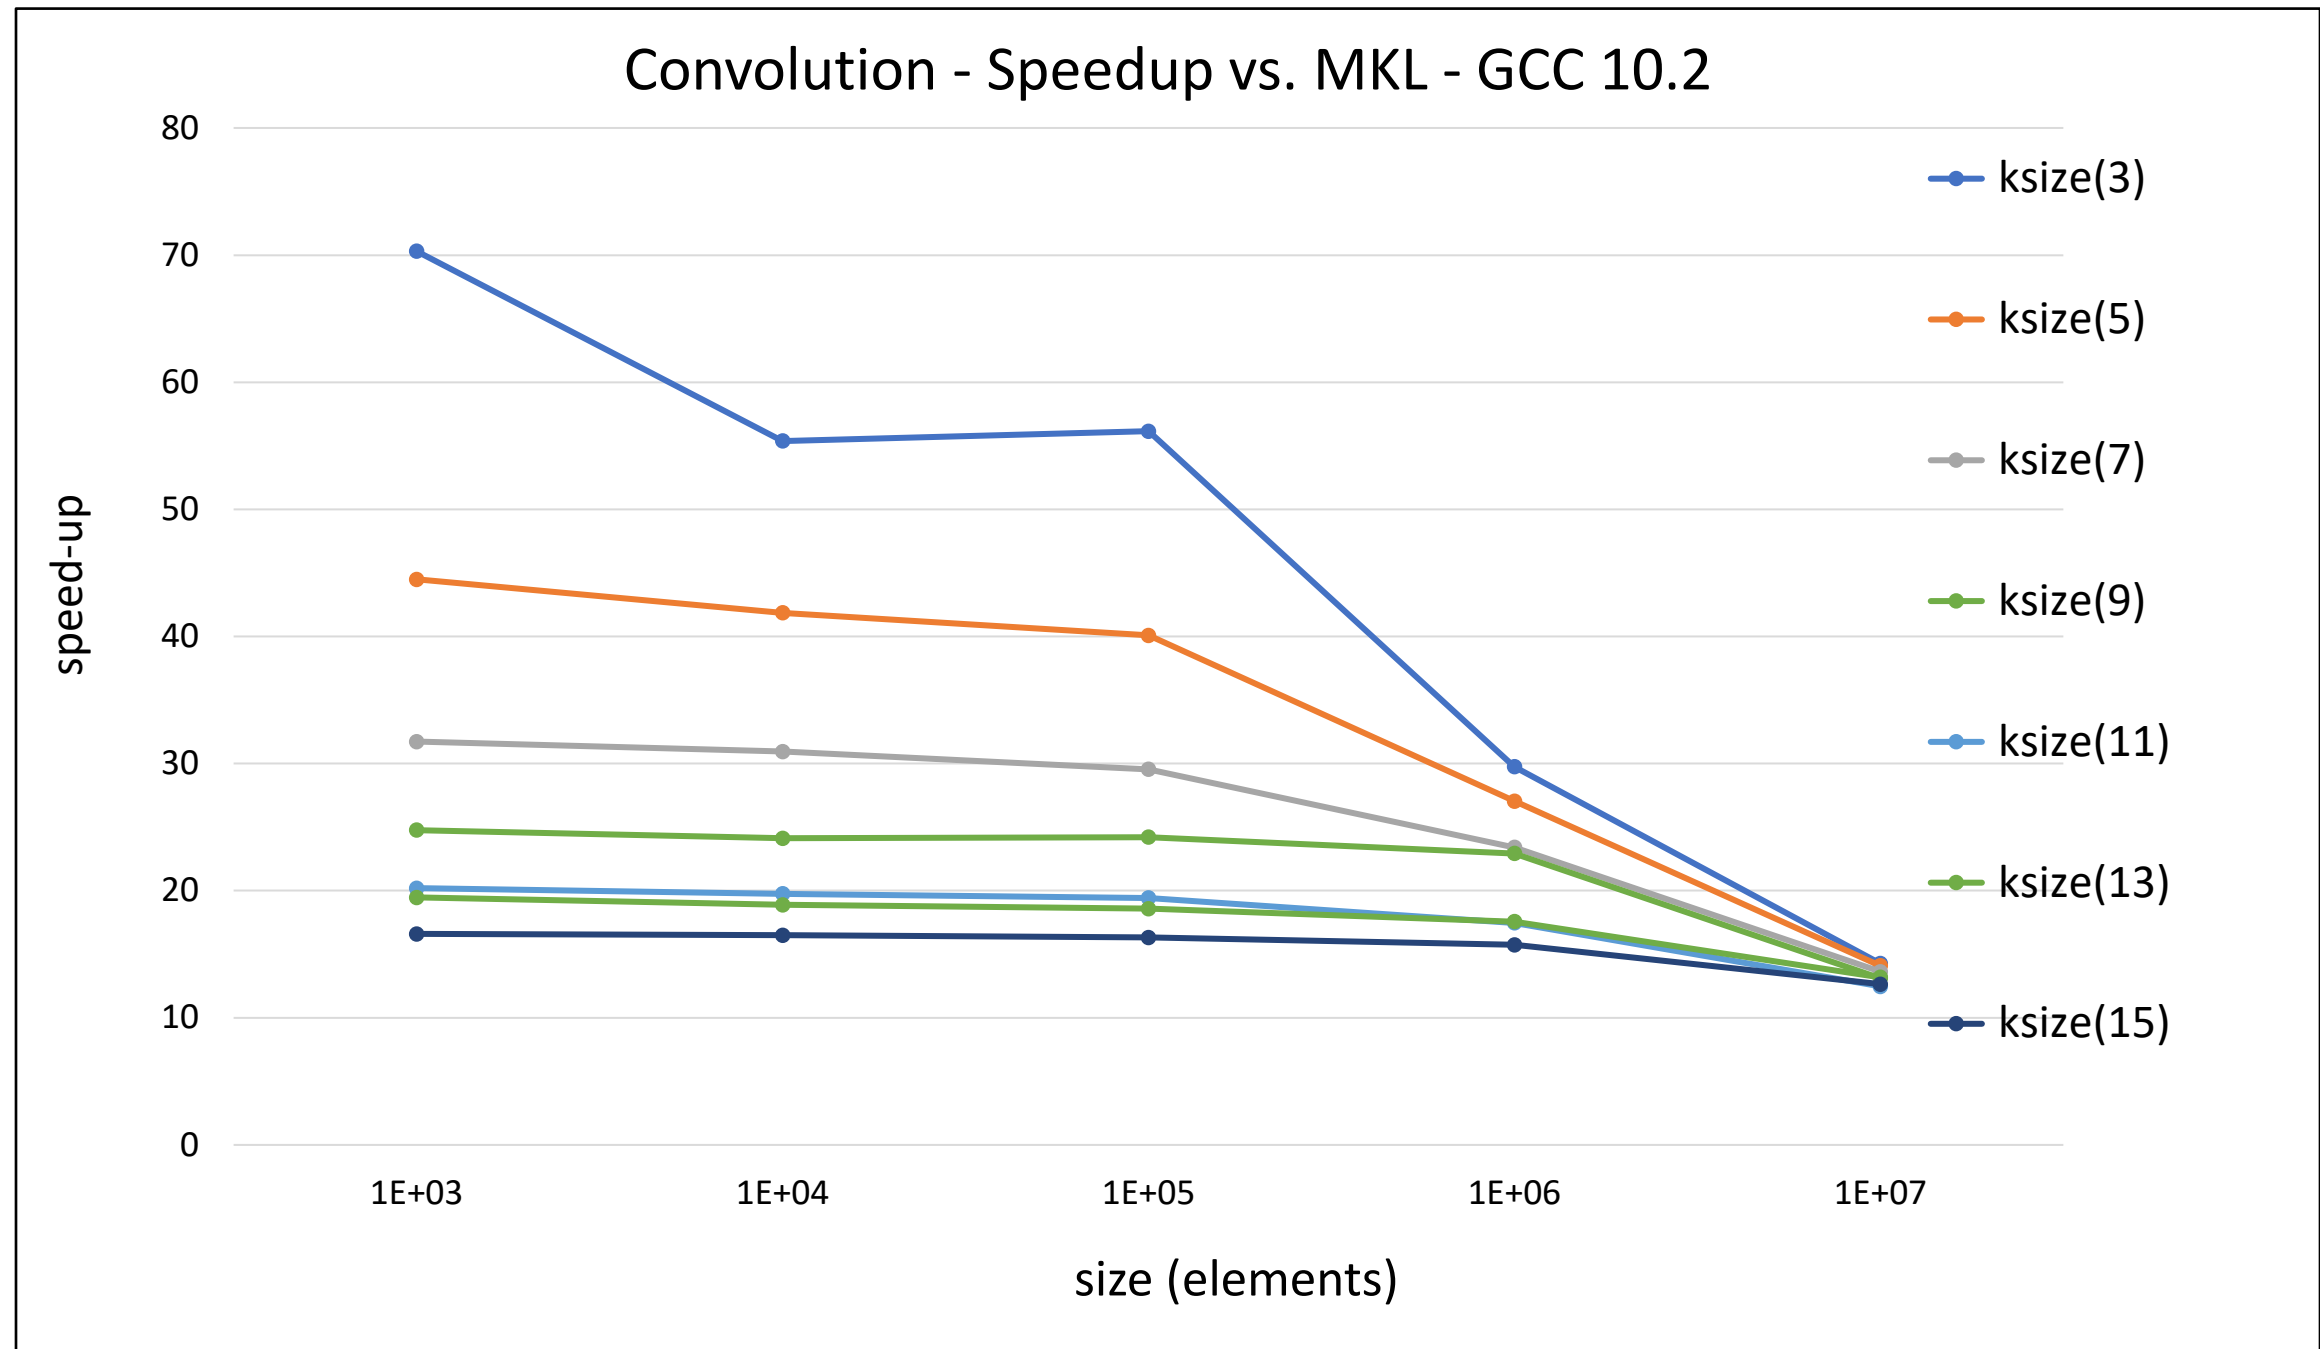
        for (int k = 0; k < KernelSize; ++k)
        {
            sum = fused_multiply_add(kcoeff[k], lo, sum); // - Update the accumulator
            in_place_shift_down_with_carry<1>(lo, hi);
        }
        store_to(pdst, sum);
    }
}
```

- Ubuntu 18.04 on Cascade Lake
 - **GCC 10.2**, all code compiled with `-O3 -mavx512 -march=skylake`
 - **Clang 10.0.1**, all code compiled with `-O3 -mavx512 -march=skylake`
- Intel MKL 2020.1.217
- Odd-number kernel sizes 3 – 15
- Element counts of 1E03 through 1E07 (by 10s)
- Collect timings for each combination using MKL and Small-Kernel AVX

Results – Convolution – Clang 10.0.1

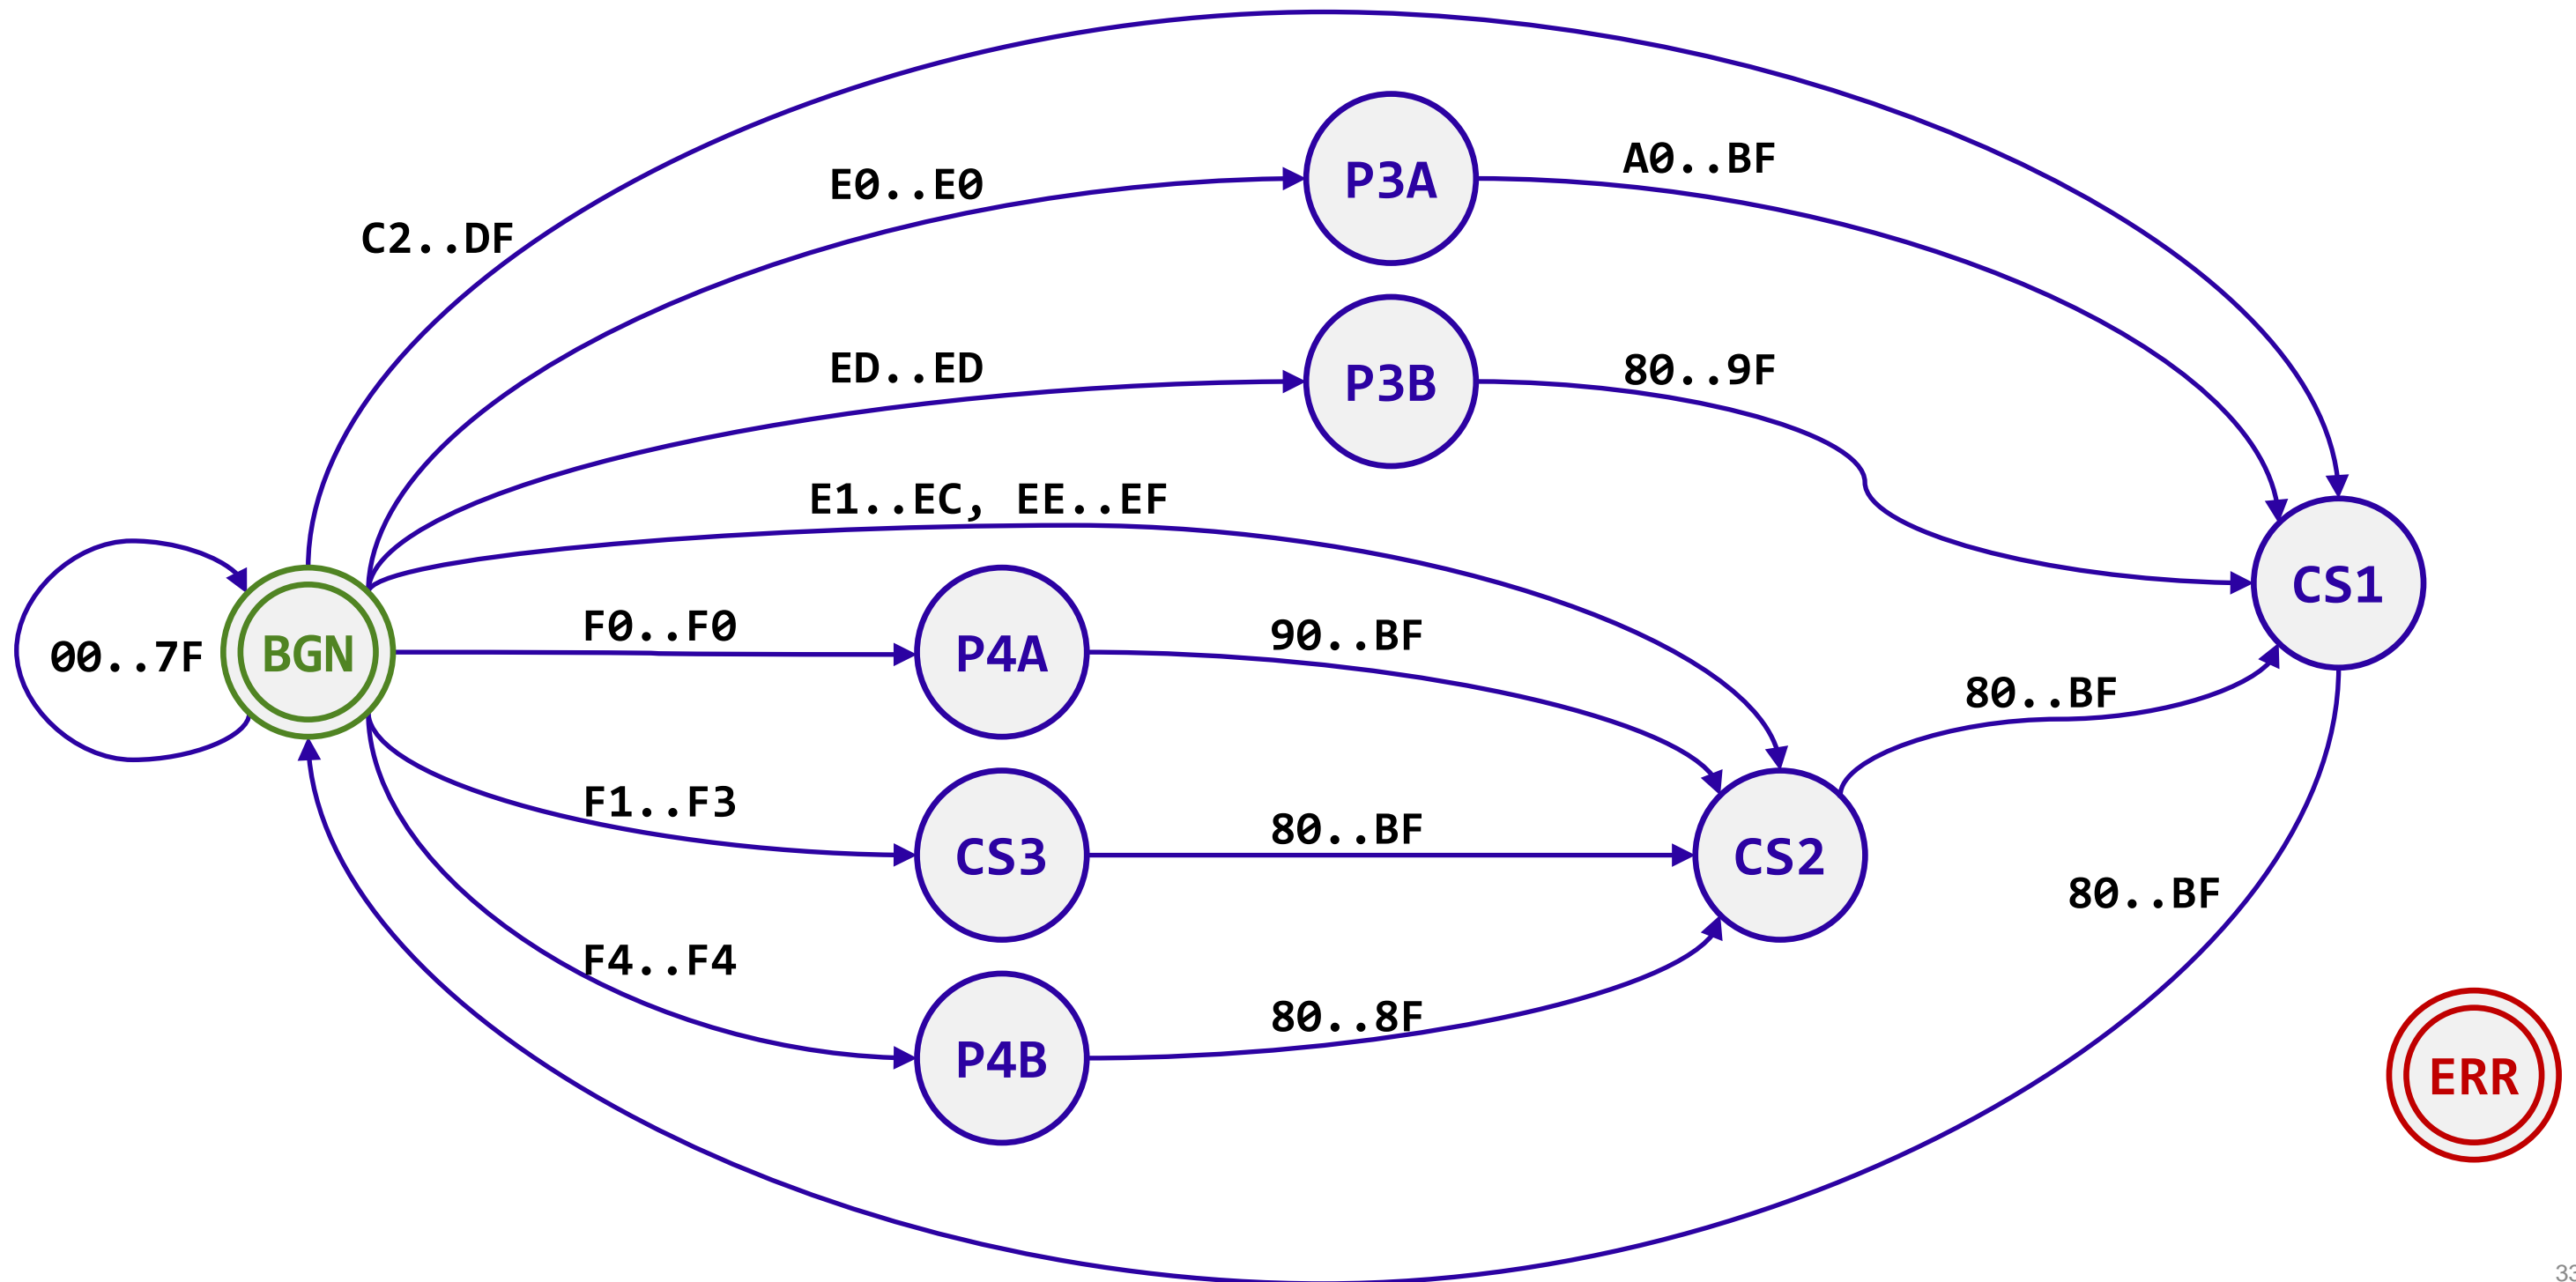


Results – Convolution – GCC 10.2



UTF-8 to UTF-32 Conversion Algorithm

The UTF-8 Decoding DFA



The Basic Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::BasicConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
        {
            *pDst++ = cdpt;
        }
        else
        {
            return -1;
        }
    }
    return pDst - pDstOrig;
}
```

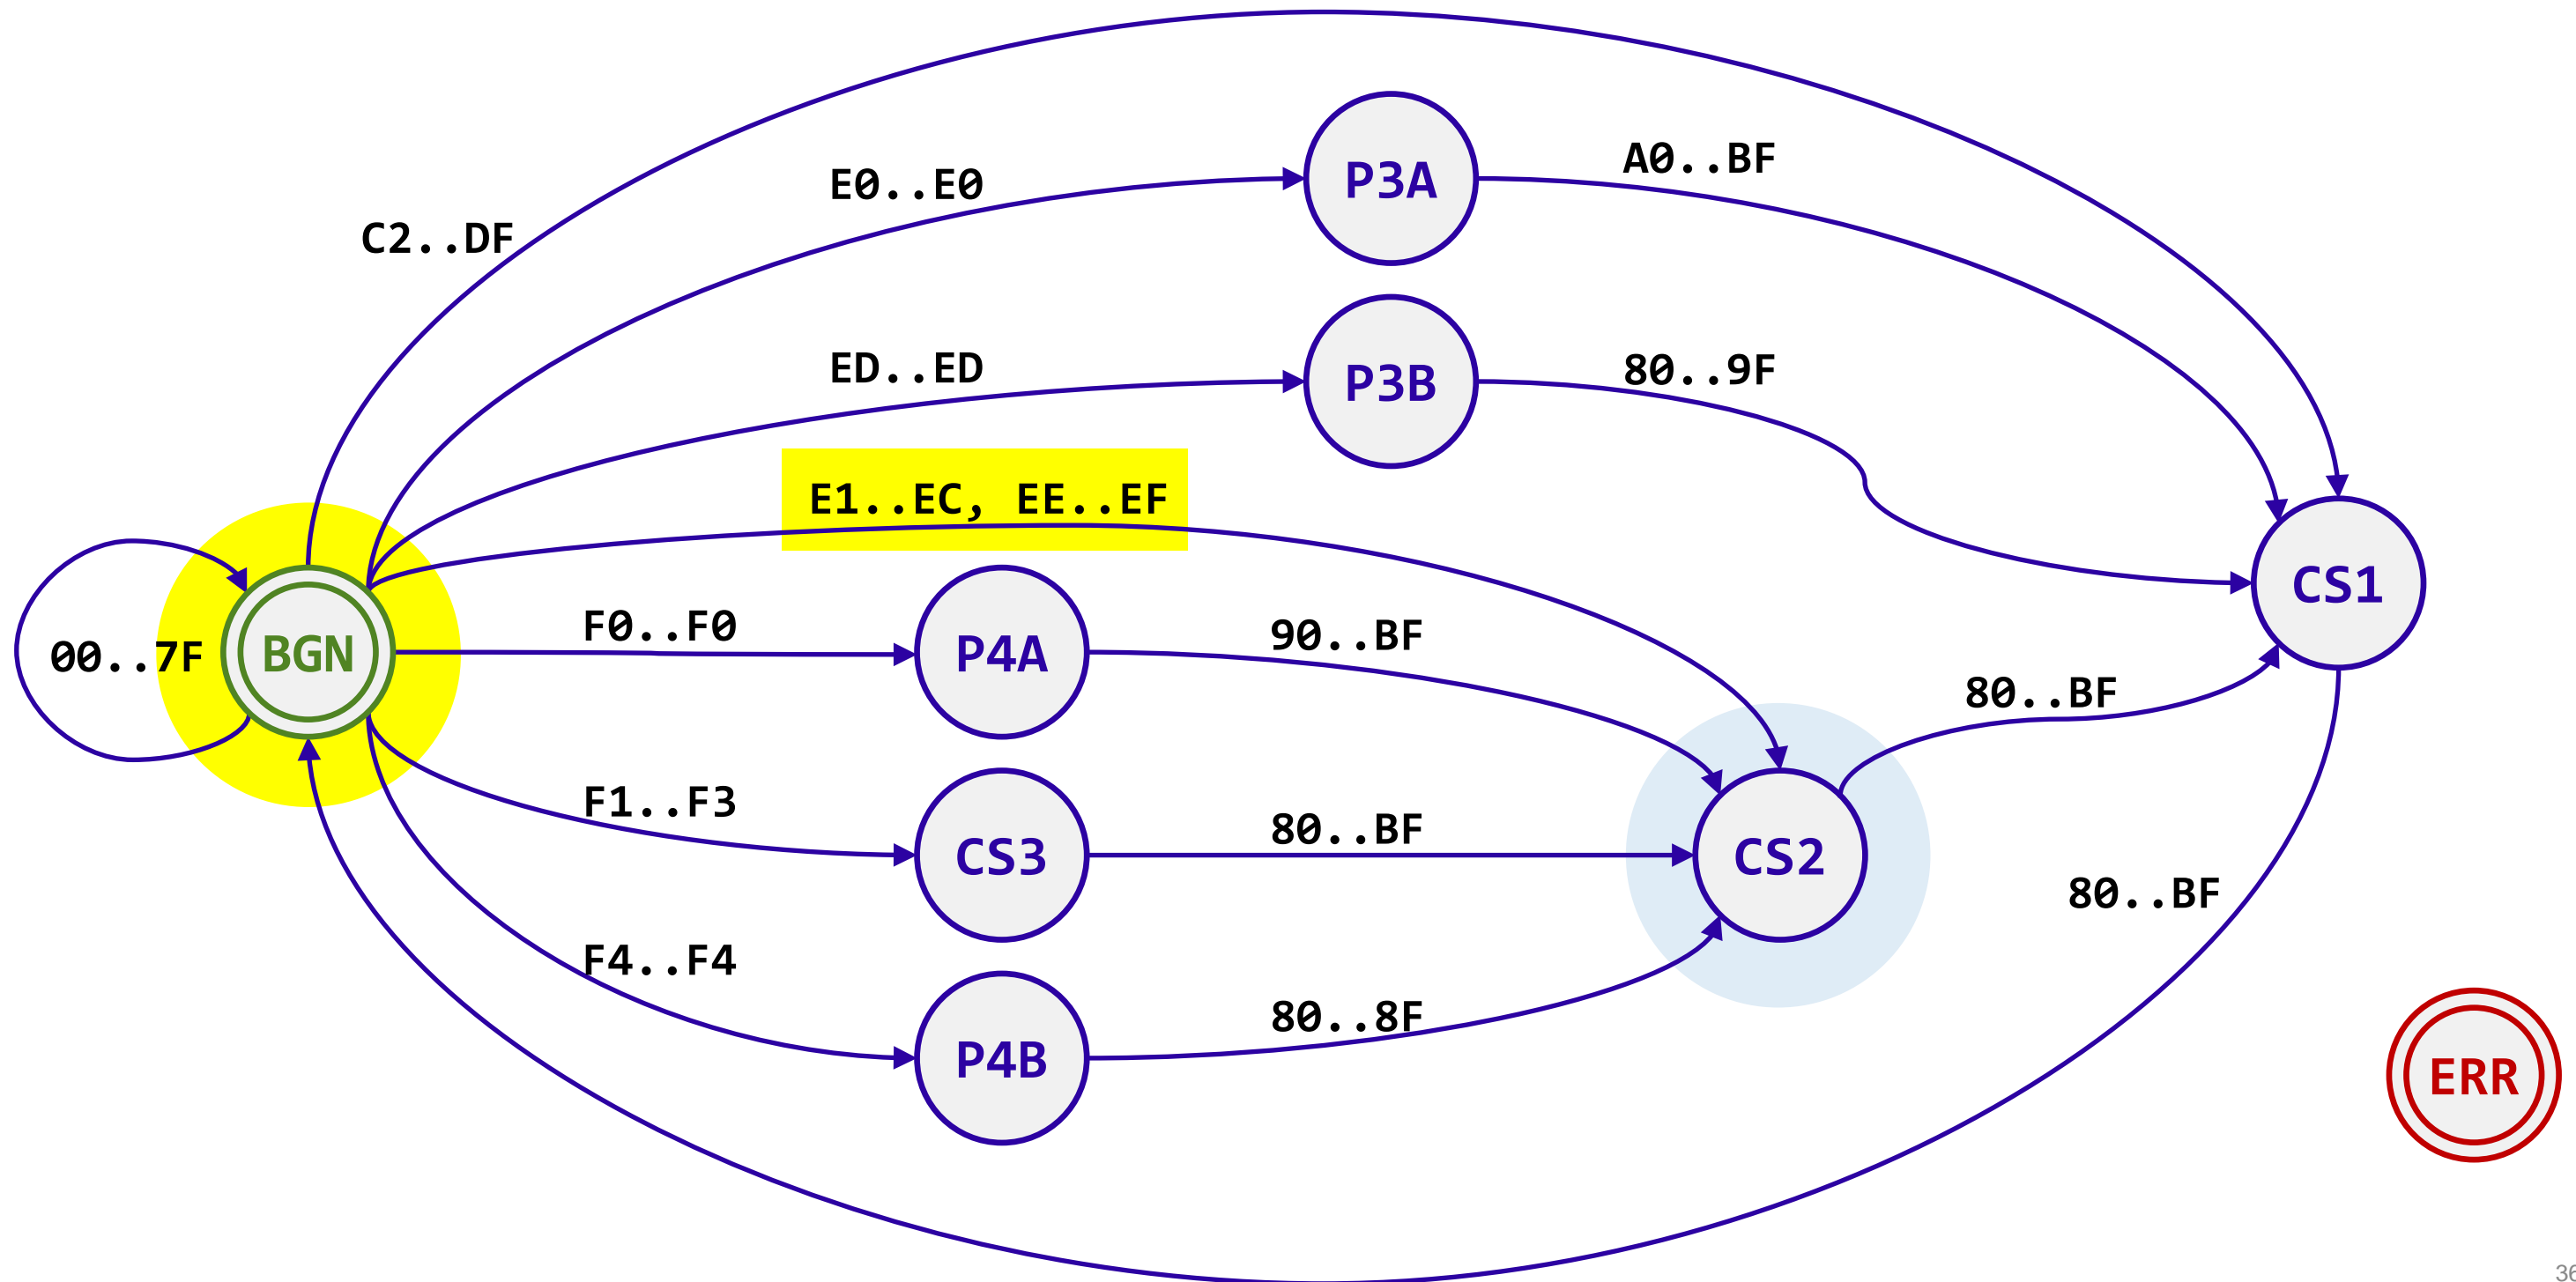
Converting a Single Code Point - Overview

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    FirstUnitInfo    info;    //- The descriptor for the first code unit
    char32_t          unit;    //- The current UTF-8 code unit
    int32_t           type;    //- The code unit's character class
    int32_t           curr;    //- The current DFA state

    info = smTables.maFirstUnitTable[*pSrc++];    //- Look up the first descriptor
    cdpt = info.mFirstOctet;    //- Get the initial code point value
    curr = info.mNextState;    //- Advance to the next state

    while (curr > ERR)    //- Loop over subsequent units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;    //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);    //- Adjust code point with continuation bits
            type = smTables.maOctetCategory[unit];    //- Look up the code unit's character class
            curr = smTables.maTransitions[curr + type];    //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

A Decoding Example – { .. **E2** 88 85 .. }



The Basic Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::BasicConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
        {
            *pDst++ = cdpt;
        }
        else
        {
            return -1;
        }
    }
    return pDst - pDstOrig;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    FirstUnitInfo    info;    //- The descriptor for the first code unit
    char32_t          unit;    //- The current UTF-8 code unit
    int32_t           type;    //- The code unit's character class
    int32_t           curr;    //- The current DFA state

    info = smTables.maFirstUnitTable[*pSrc++];    //- Look up the first code unit descriptor
    cdpt = info.mFirstOctet;                      //- Get the initial code point value
    curr = info.mNextState;                       //- Advance to the next state

    while (curr > ERR)                            //- Loop over subsequent code units
    {
        ...
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    FirstUnitInfo    info;    //- The descriptor for the first code unit
    char32_t          unit;    //- The current UTF-8 code unit
    int32_t           type;    //- The code unit's character class
    int32_t           curr;    //- The current DFA state

    info = smTables.maFirstUnitTable[*pSrc++];    //- Look up the first code unit descriptor
    cdpt = info.mFirstOctet;                    //- Get the initial code point value
    curr = info.mNextState;                      //- Advance to the next state

    while (curr > ERR)                            //- Loop over subsequent code units
    {
        ...
    }
    return curr;
}
```

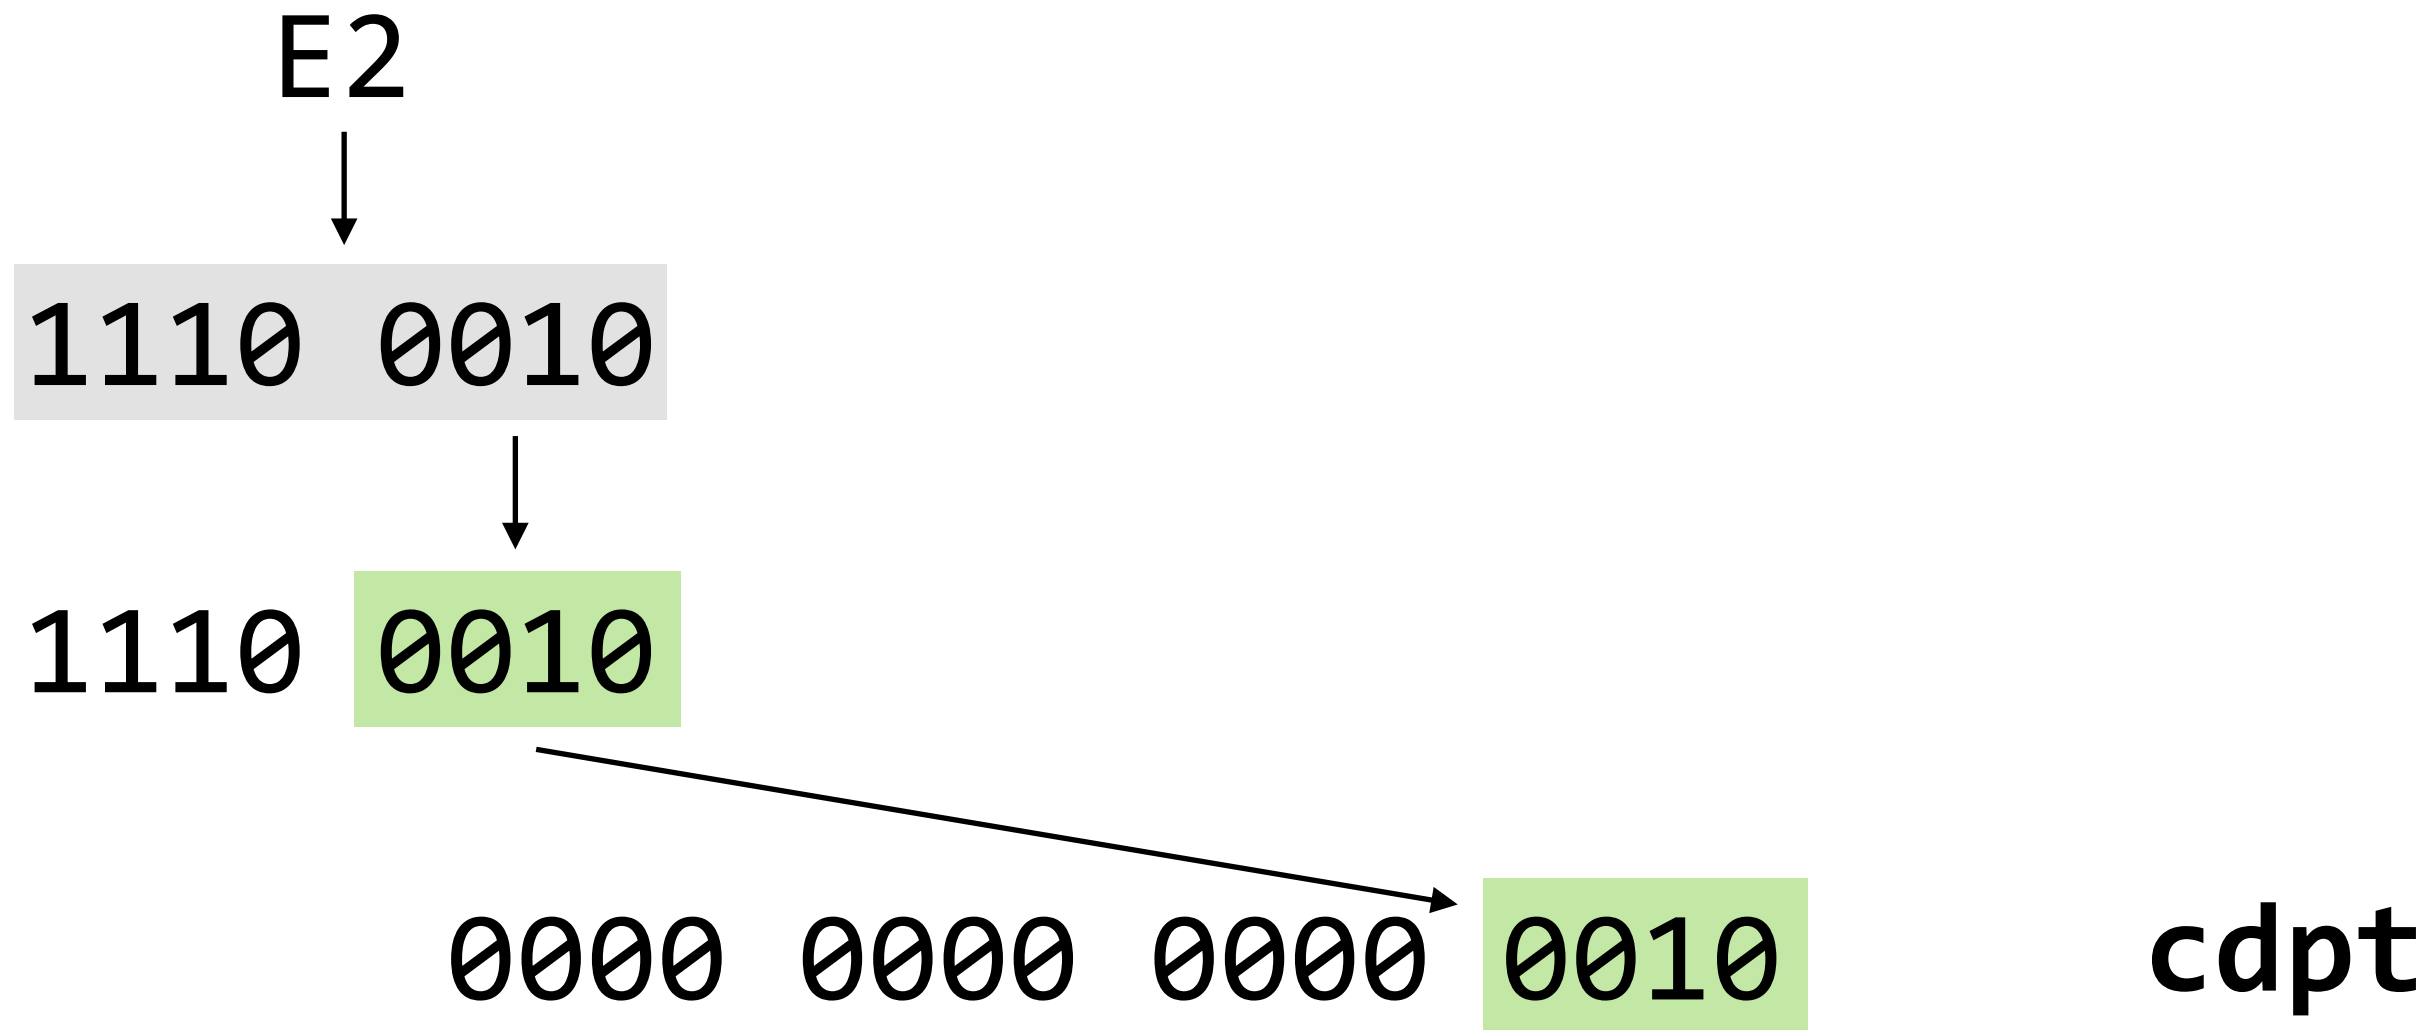
Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    FirstUnitInfo    info;    //- The descriptor for the first code unit
    char32_t          unit;    //- The current UTF-8 code unit
    int32_t           type;    //- The code unit's character class
    int32_t           curr;    //- The current DFA state

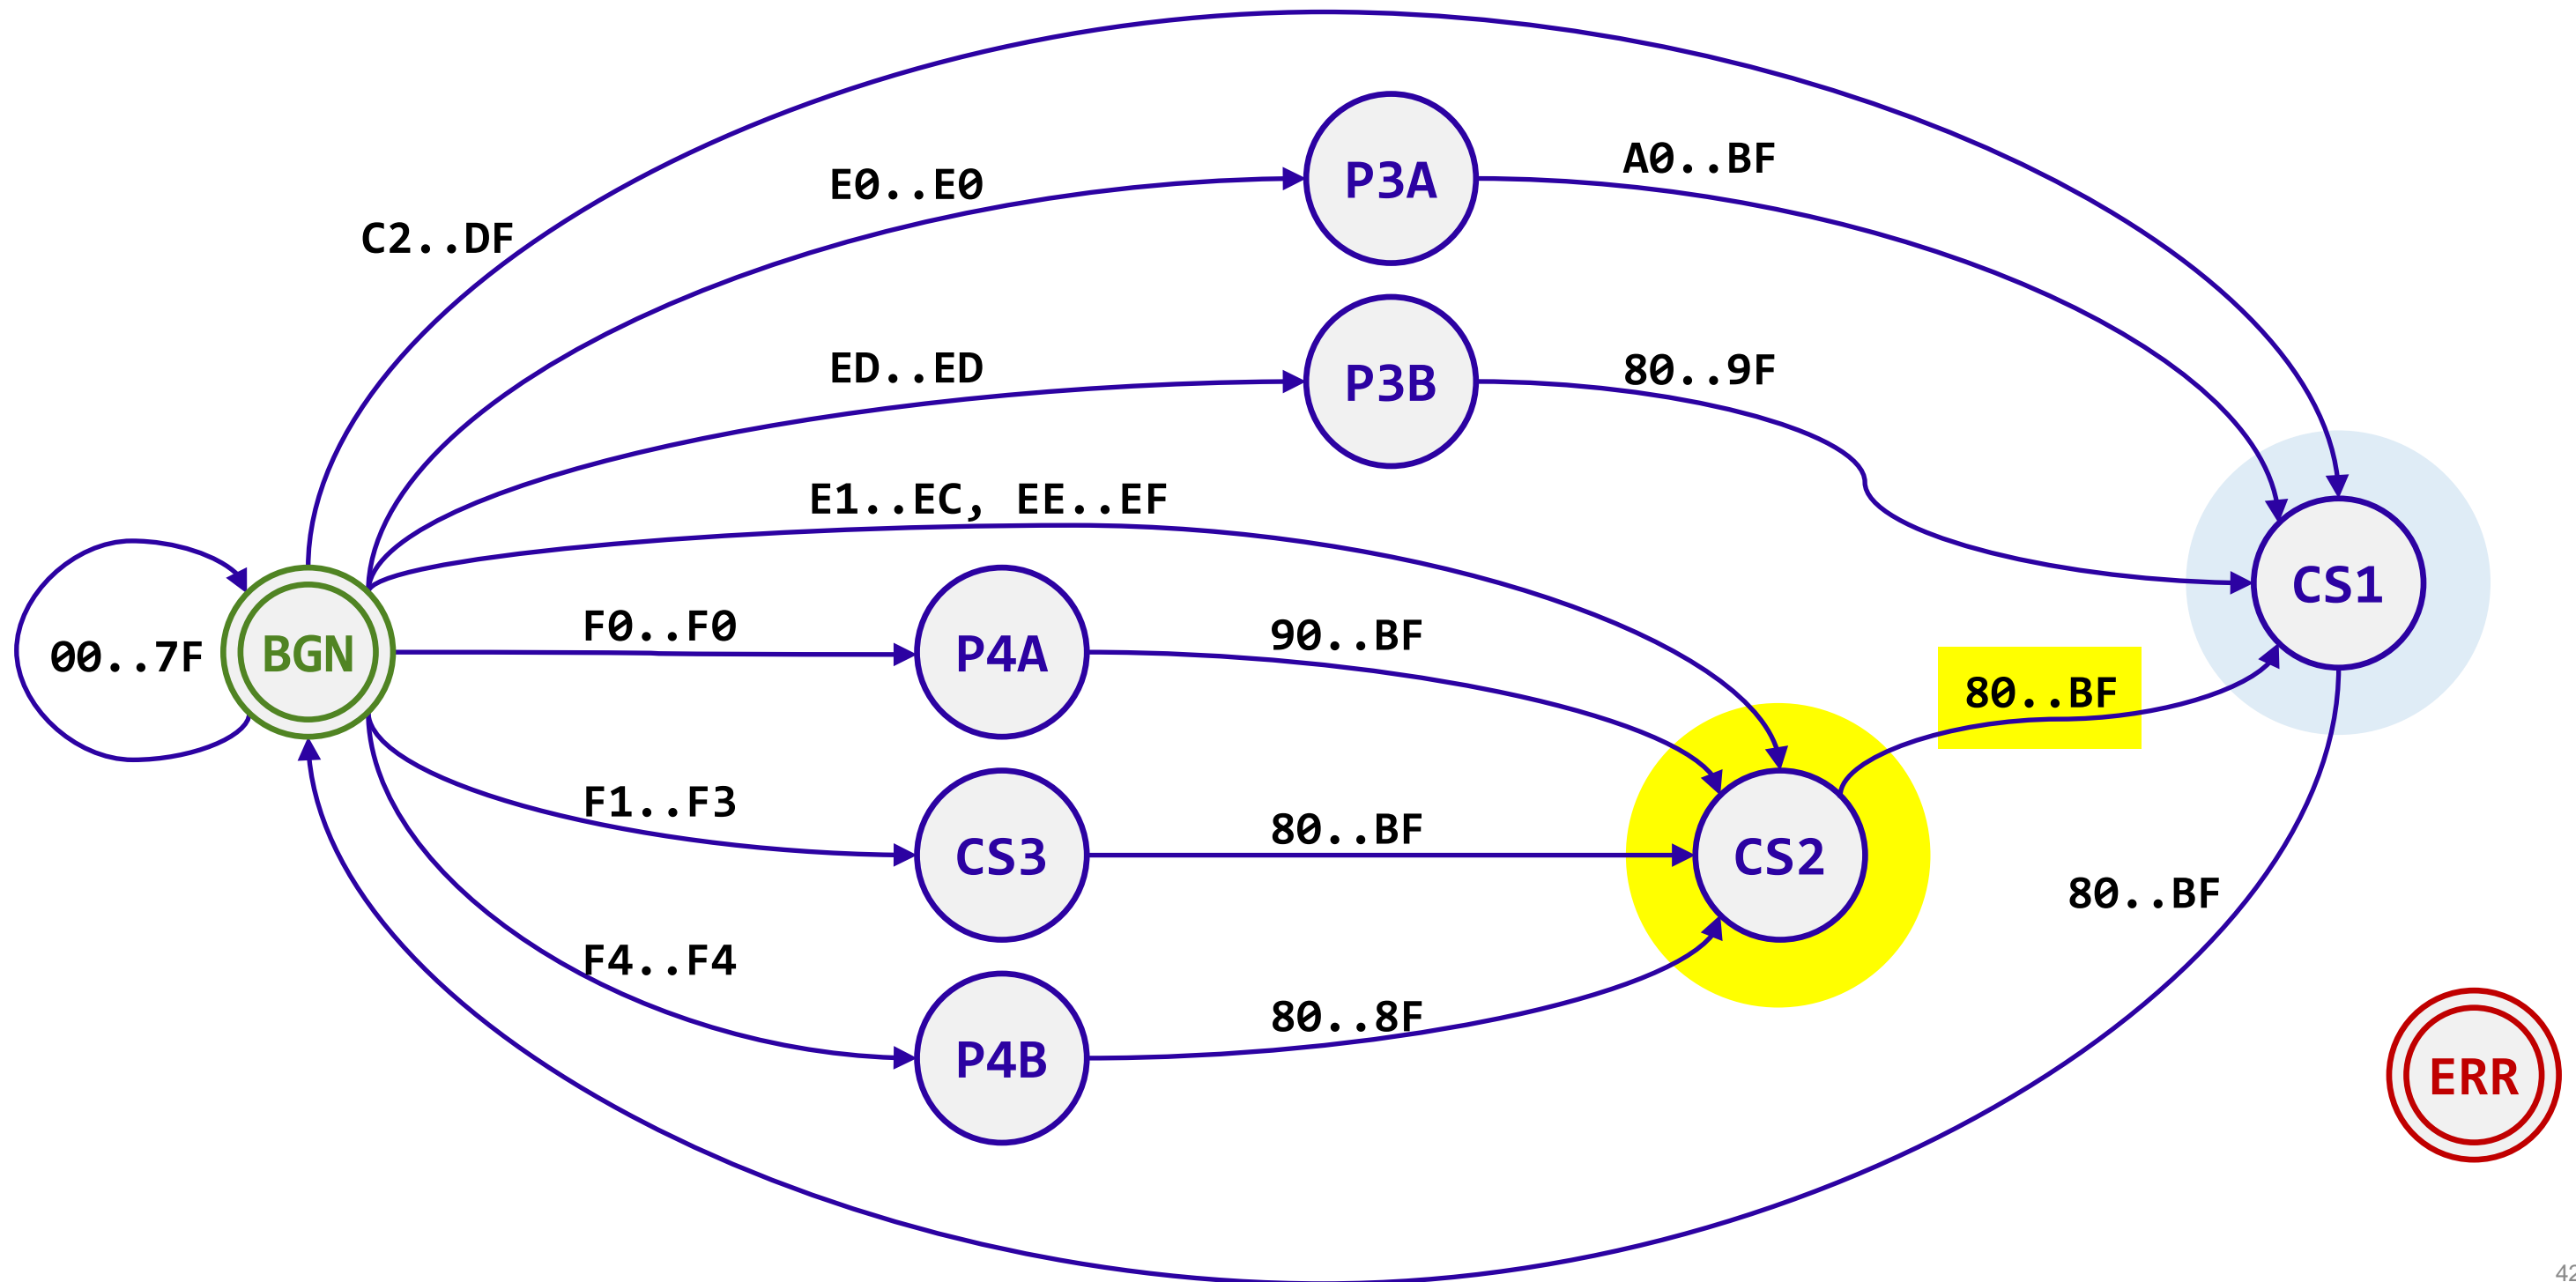
    info = smTables.maFirstUnitTable[*pSrc++];    //- Look up the first code unit descriptor
    cdpt = info.mFirstOctet;                      //- Get the initial code point value
    curr = info.mNextState;                       //- Advance to the next state

    while (curr > ERR)                            //- Loop over subsequent code units
    {
        ...
    }
    return curr;
}
```


A Decoding Example – { .. E2 88 85 .. }



A Decoding Example – { .. E2 88 85 .. }



Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                            //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);         //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];      //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type]; //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

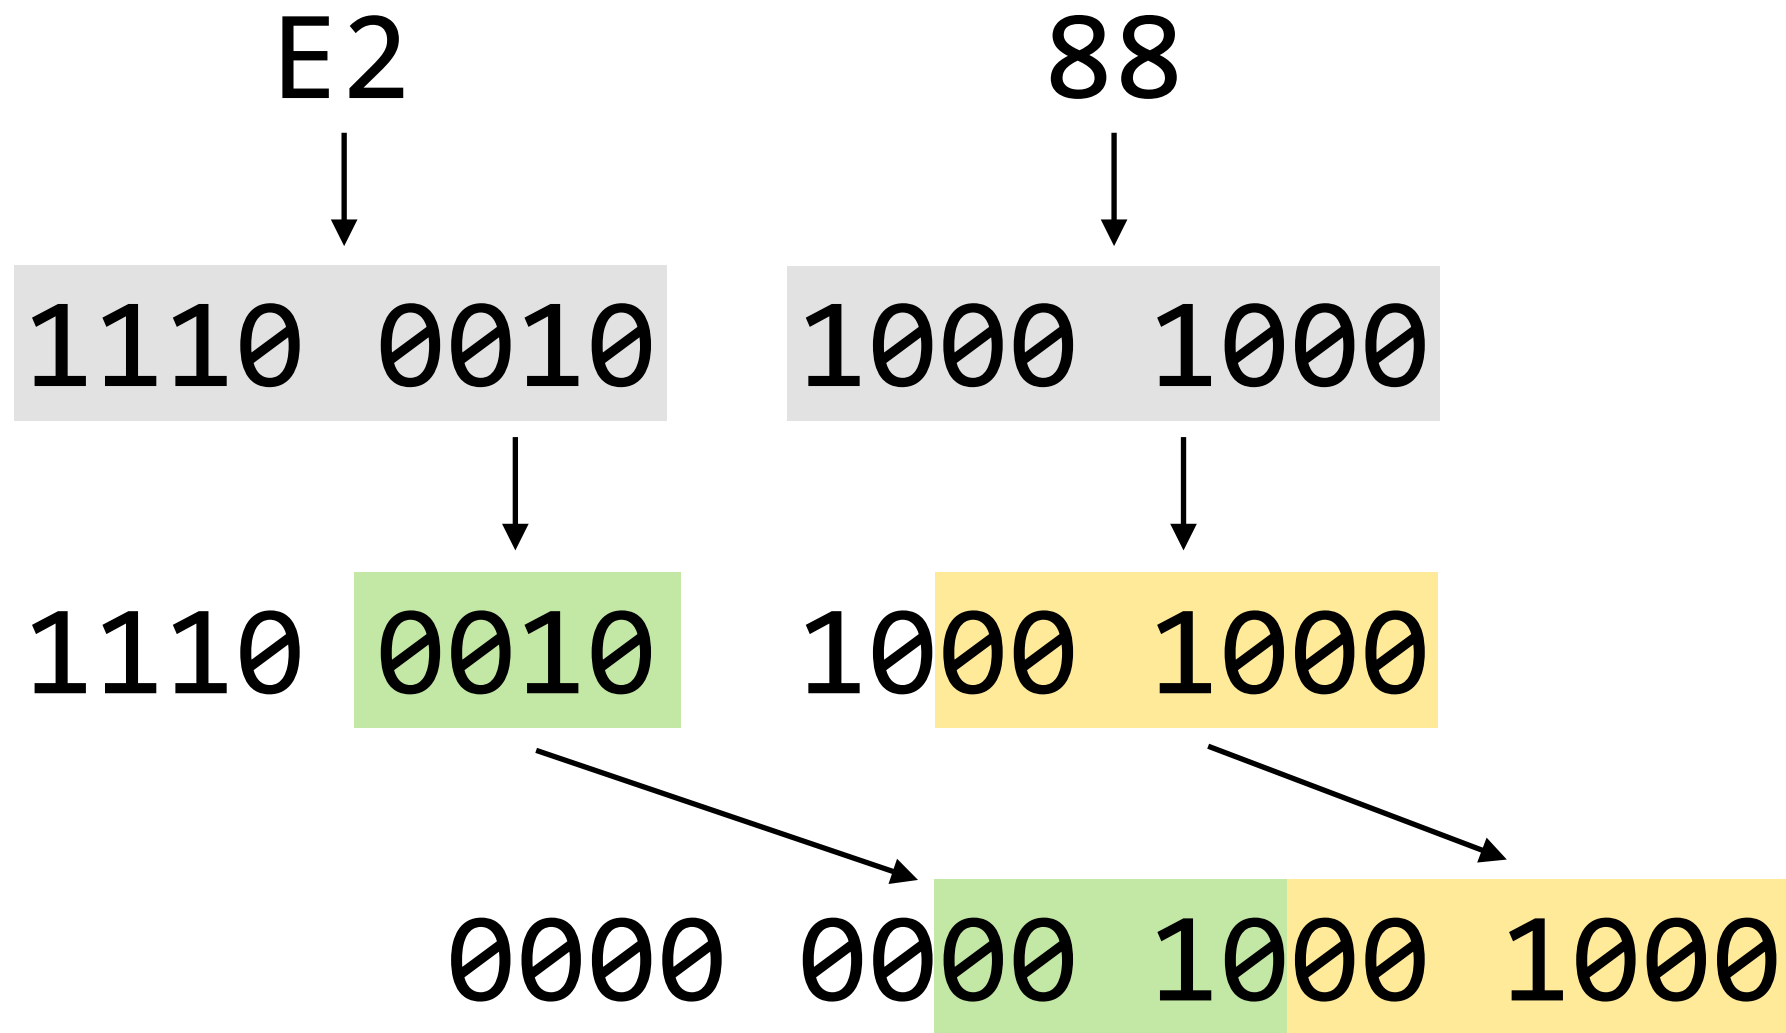
    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);            //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];         //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];    //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

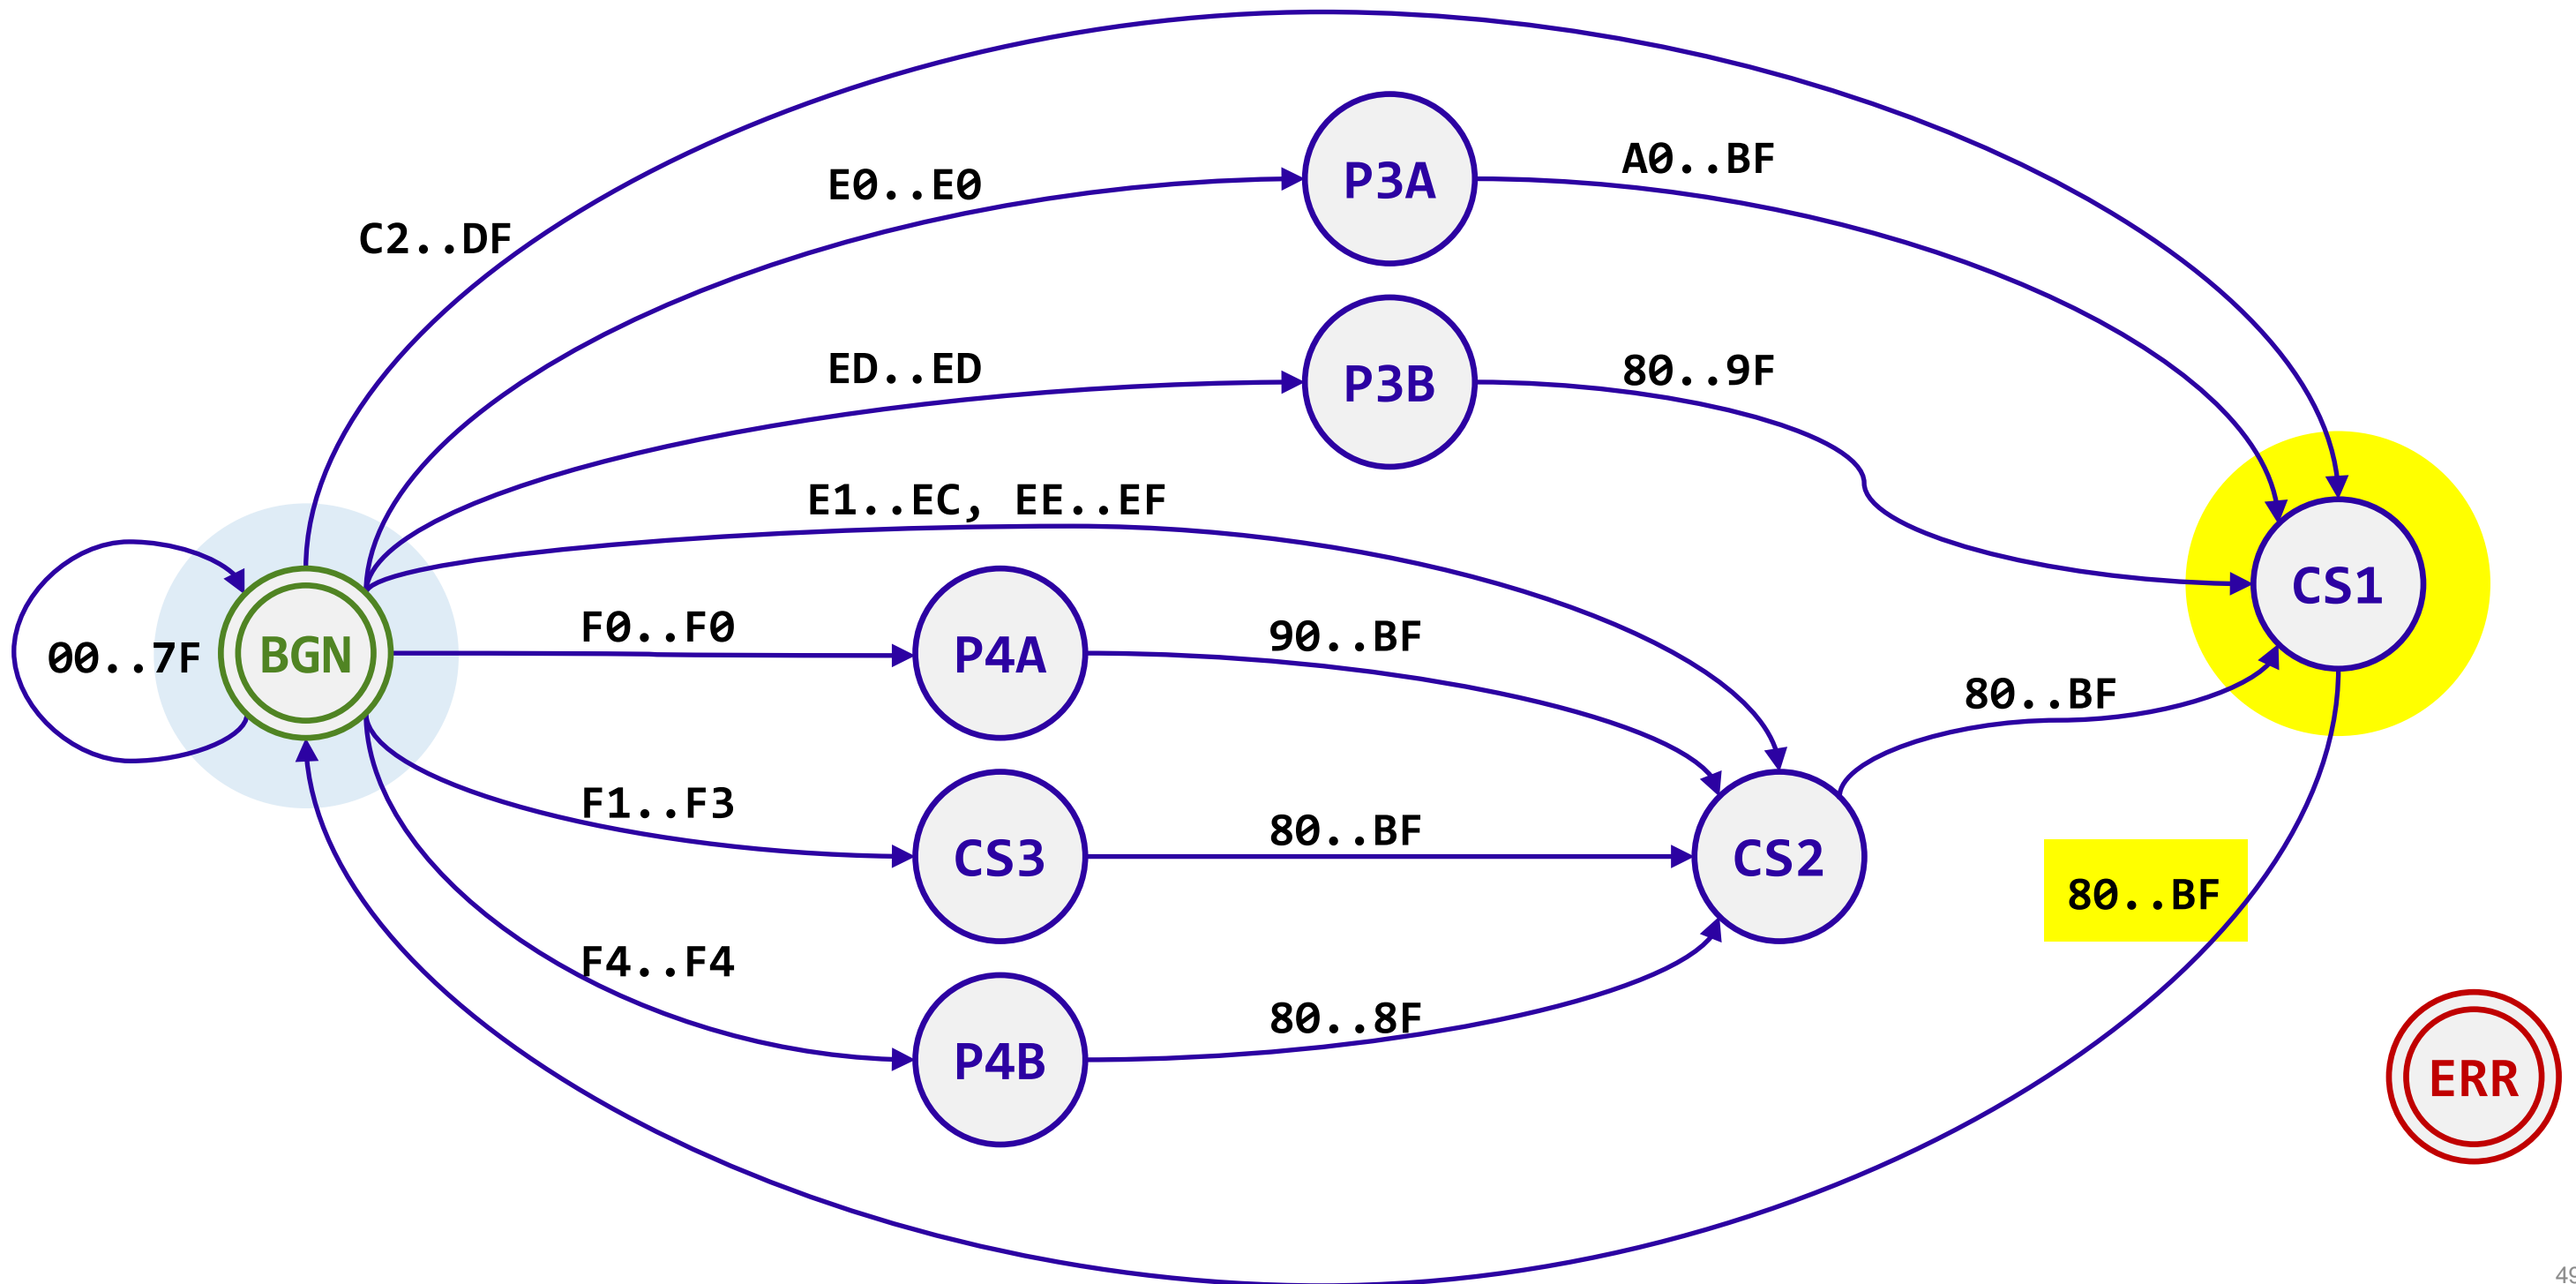
    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

A Decoding Example – { .. E2 88 85 .. }



cdpt

A Decoding Example – { .. E2 88 85 .. }



Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...


    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

Converting a Single Code Point

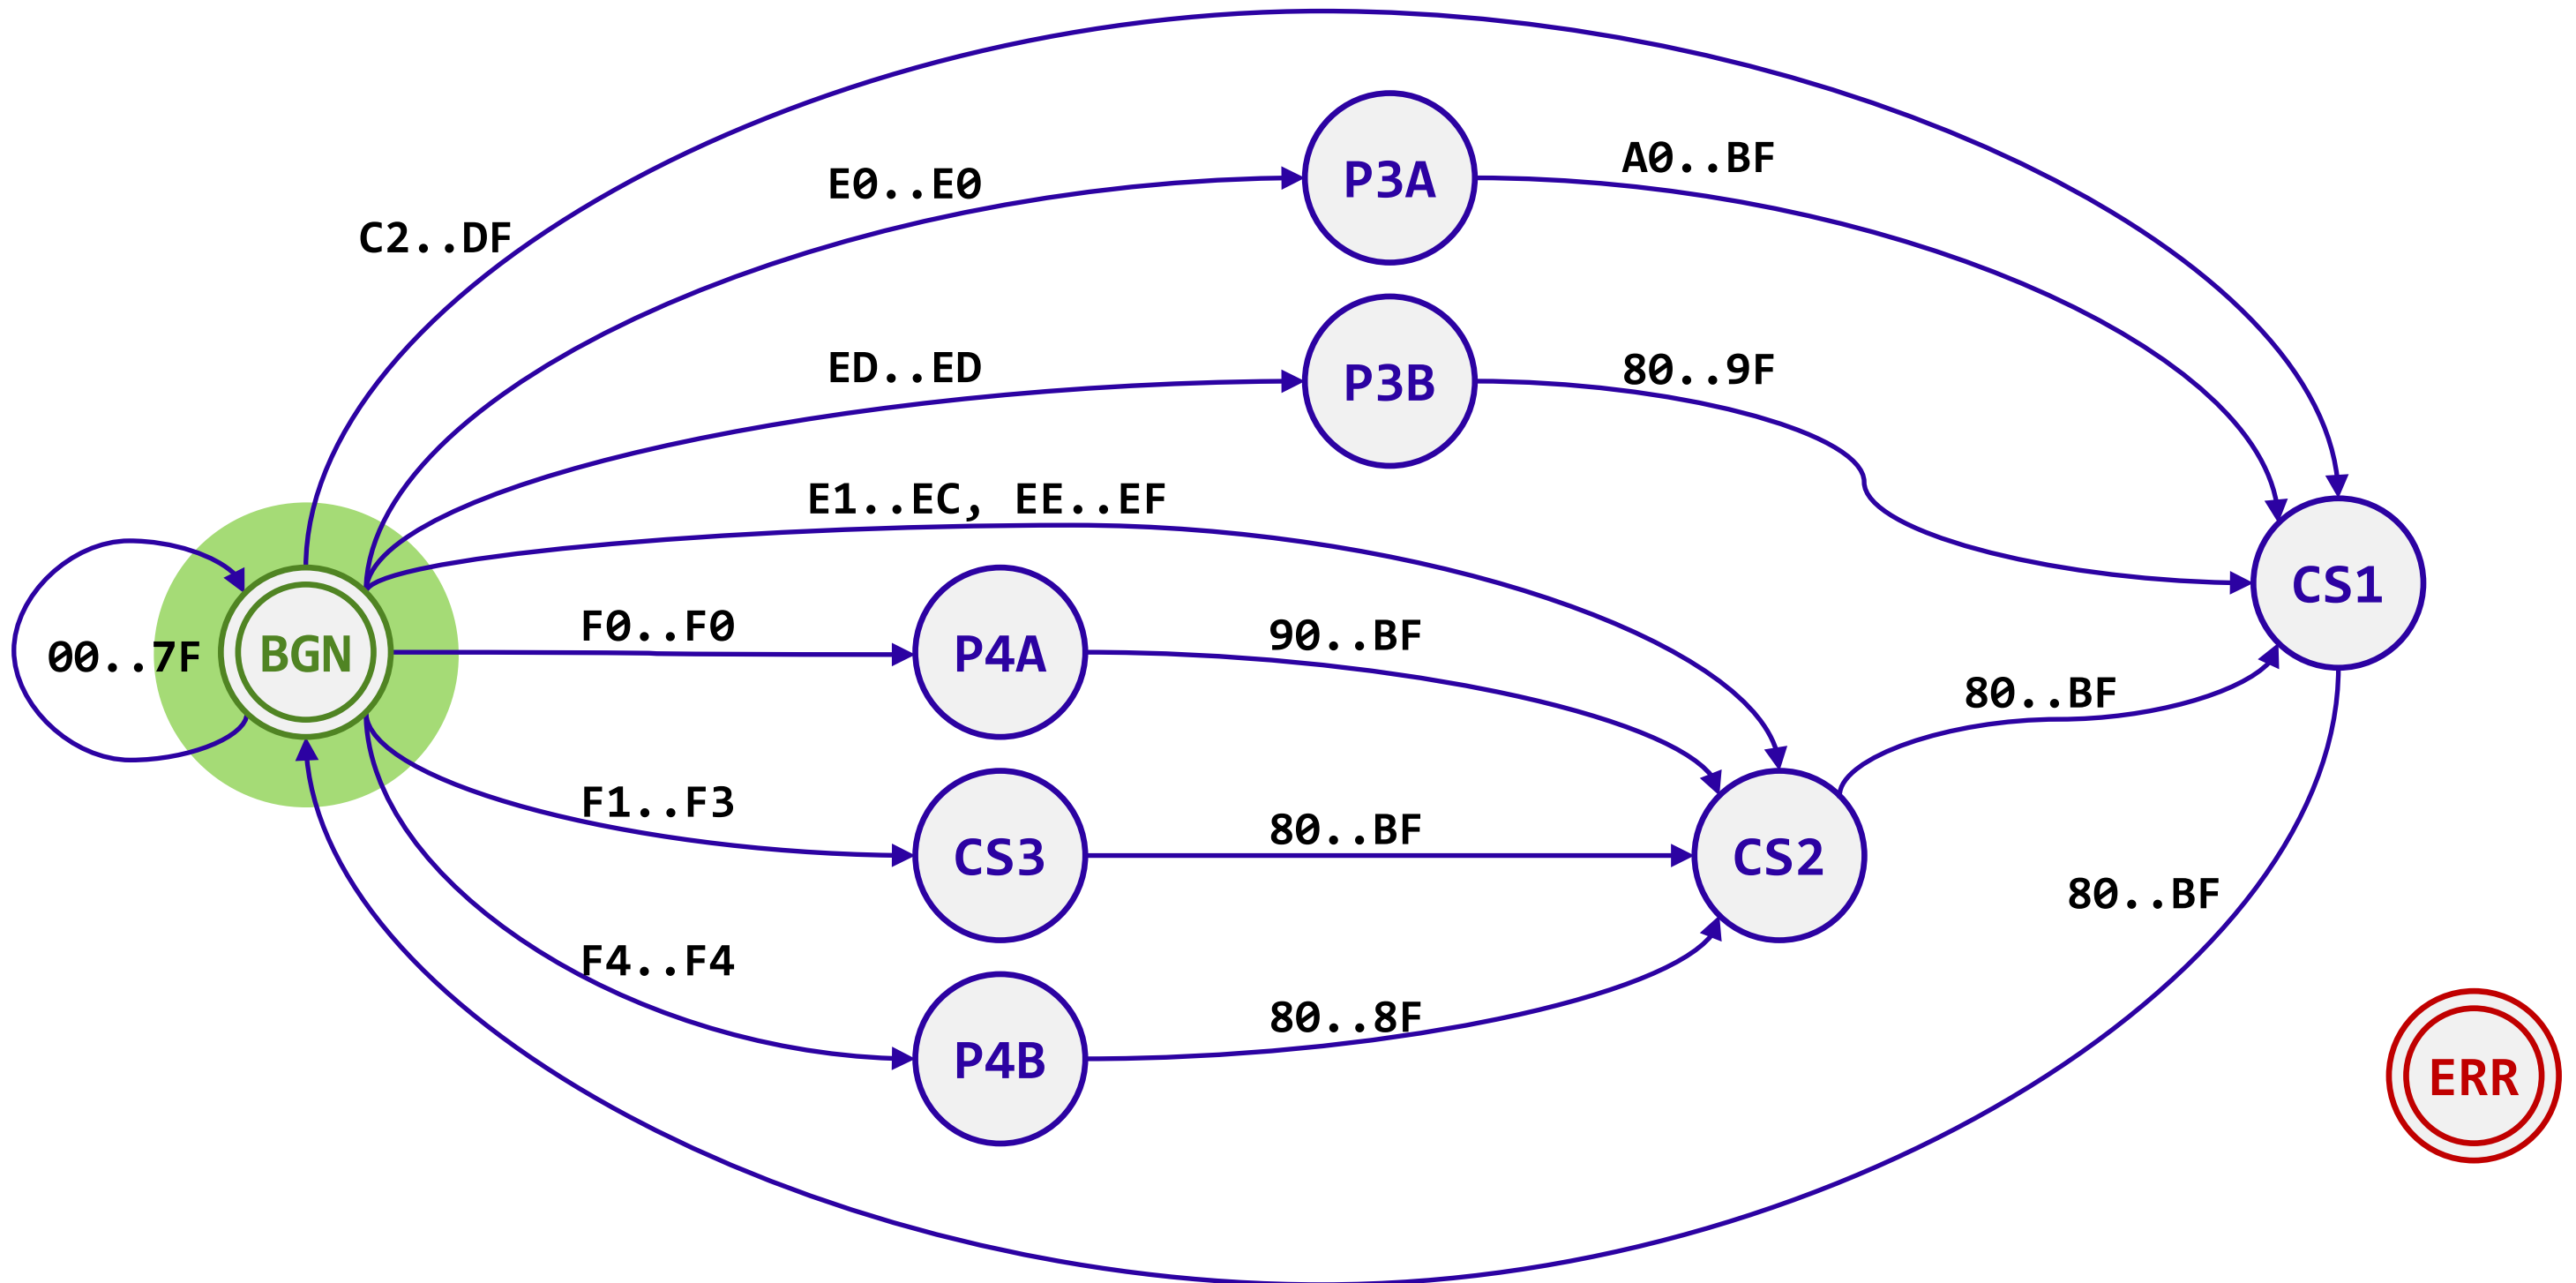
```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);             //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];          //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];     //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```

A Decoding Example – { .. E2 88 85 .. }



A Decoding Example – { .. E2 88 85 .. }

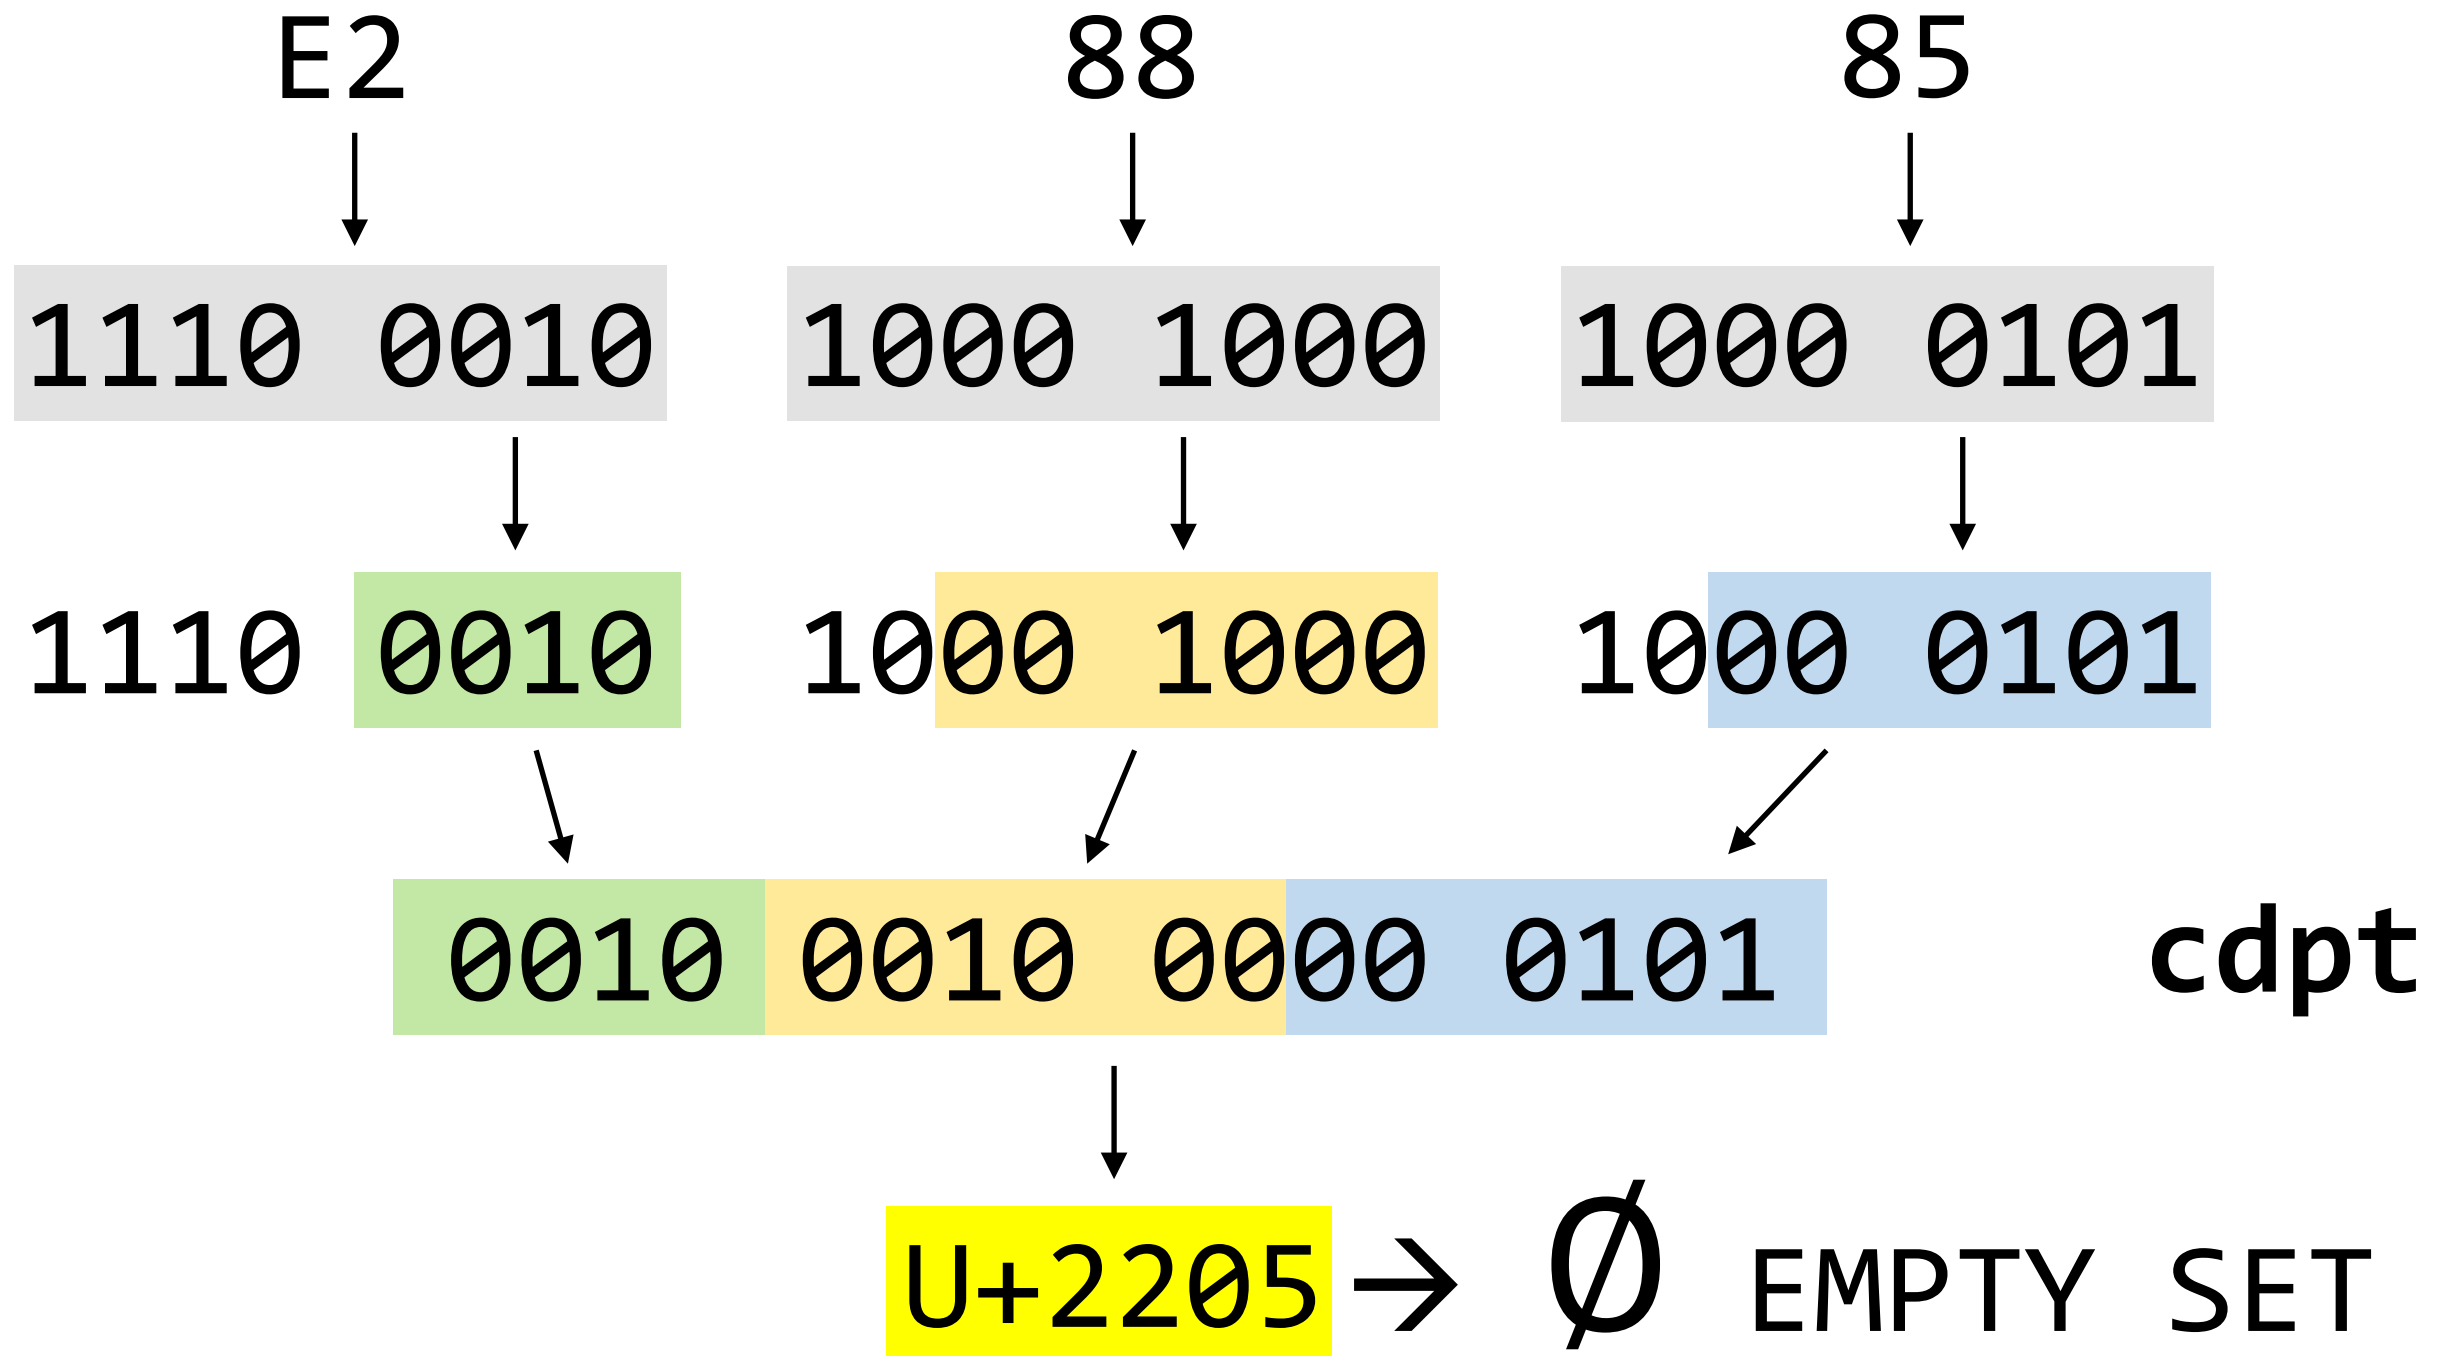


Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    ...

    while (curr > ERR)                                     //- Loop over subsequent code units
    {
        if (pSrc < pSrcEnd)
        {
            unit = *pSrc++;                                //- Cache the current code unit
            cdpt = (cdpt << 6) | (unit & 0x3F);            //- Adjust code point with new bits
            type = smTables.maOctetCategory[unit];         //- Look up the code unit's char class
            curr = smTables.maTransitions[curr + type];    //- Advance to the next state
        }
        else
        {
            return ERR;
        }
    }
    return curr;
}
```


A Decoding Example – { .. E2 88 85 .. }



The Basic Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::BasicConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
        {
            *pDst++ = cdpt;
        }
        else
        {
            return -1;
        }
    }
    return pDst - pDstOrig;
}
```

Optimizing for ASCII

The Basic Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::BasicConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
        {
            *pDst++ = cdpt;
        }
        else
        {
            return -1;
        }
    }
    return pDst - pDstOrig;
}
```

Converting a Single Code Point

```
KEWB_FORCE_INLINE int32_t
UtfUtils::Advance(char8_t const*& pSrc, char8_t const* pSrcEnd, char32_t& cdpt) noexcept
{
    FirstUnitInfo    info;    //- The descriptor for the first code unit
    char32_t          unit;    //- The current UTF-8 code unit
    int32_t           type;    //- The code unit's character class
    int32_t           curr;    //- The current DFA state

    info = smTables.maFirstUnitTable[*pSrc++];    //- Look up the first code unit descriptor
    cdpt = info.mFirstOctet;                      //- Get the initial code point value
    curr = info.mNextState;                      //- Get the second state

    while (curr > ERR)                            //- Loop over subsequent code units
    {
        ...
    }
    return curr;
}
```

The ASCII-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::FastConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (*pSrc < 0x80)
        {
            *pDst++ = *pSrc++;
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    return pDst - pDstOrig;
}
```

Optimizing for ASCII with SSE

The ASCII-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::FastConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < pSrcEnd)
    {
        if (*pSrc < 0x80)
        {
            *pDst++ = *pSrc++;
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    return pDst - pDstOrig;
}
```


The SSE-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::SseConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < (pSrcEnd - sizeof(__m128i)))
    {
        if (*pSrc < 0x80)
        {
            ConvertAsciiWithSse(pSrc, pDst);
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    ...
}
```

The SSE-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t  
UtfUtils::SseConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept  
{
```

```
...
```

```
while (pSrc < pSrcEnd)  
{  
    if (*pSrc < 0x80)  
    {  
        *pDst++ = *pSrc++;  
    }  
    else  
    {  
        if (Advance(pSrc, pSrcEnd, cdpt) != ERR)  
        {  
            *pDst++ = cdpt;  
        }  
        else  
        {  
            return -1;  
        }  
    }  
}  
return pDst - pDstOrig;  
}
```

Converting ASCII Character Runs - Overview

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                       //- ASCII bit mask and advancement

    zero = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask = _mm_movemask_epi8(chunk);              //- Find the octets with high bit set

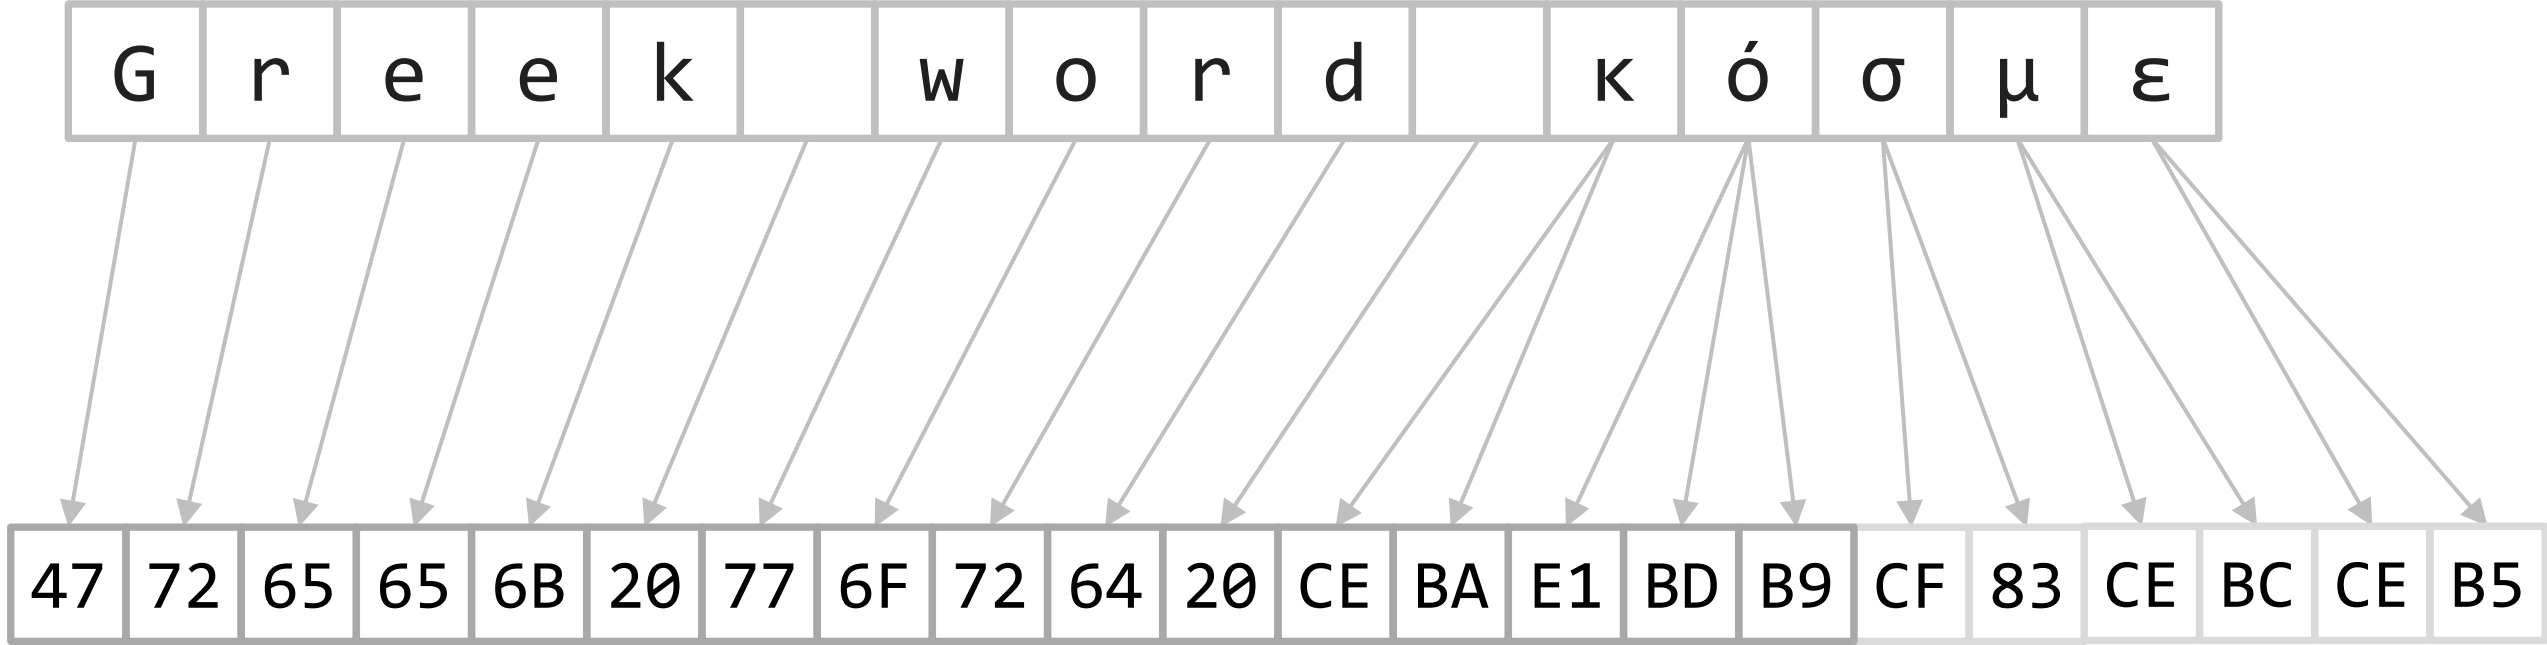
    half = _mm_unpacklo_epi8(chunk, zero);        //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);       //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);        //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);         //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr); //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);         //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr); //- Write to memory

    //- If no bits were set in the mask, then all 16 code units were ASCII.
    //-
    if (mask == 0)
    {
        pSrc += 16;
        pDst += 16;
    }

    //- Otherwise, the number of trailing (low-order) zero bits in the mask is
    //- the number of ASCII code units starting from the lowest byte address.
    else
    {
        incr = GetTrailingZeros(mask);
        pSrc += incr;
        pDst += incr;
    }
}
```

Converting ASCII Character Runs – SSE Example



LSB  MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);              //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;                //- SSE "registers"
    int32_t    mask, incr;                             //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                          //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc);      //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);                  //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);              //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);             //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);           //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);             //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);     //- Write to memory

    ...
}
```

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE “registers”
    int32_t    mask, incr;                        //- ASCII bit mask and advancement


    zero = _mm_set1_epi8(0);                      //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask = _mm_movemask_epi8(chunk);               //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);          //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);          //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

```
zero = _mm_set1_epi8(0)
```



zero

LSB



MSB

Converting ASCII Character Runs


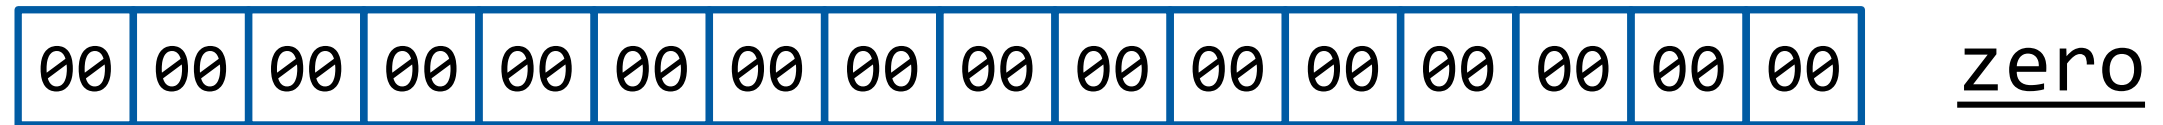
```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                    //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);             //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);        //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);        //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);      //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);        //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example



```
chunk = _mm_loadu_si128((__m128i const*) pSrc)
```



LSB → **MSB**

Converting ASCII Character Runs

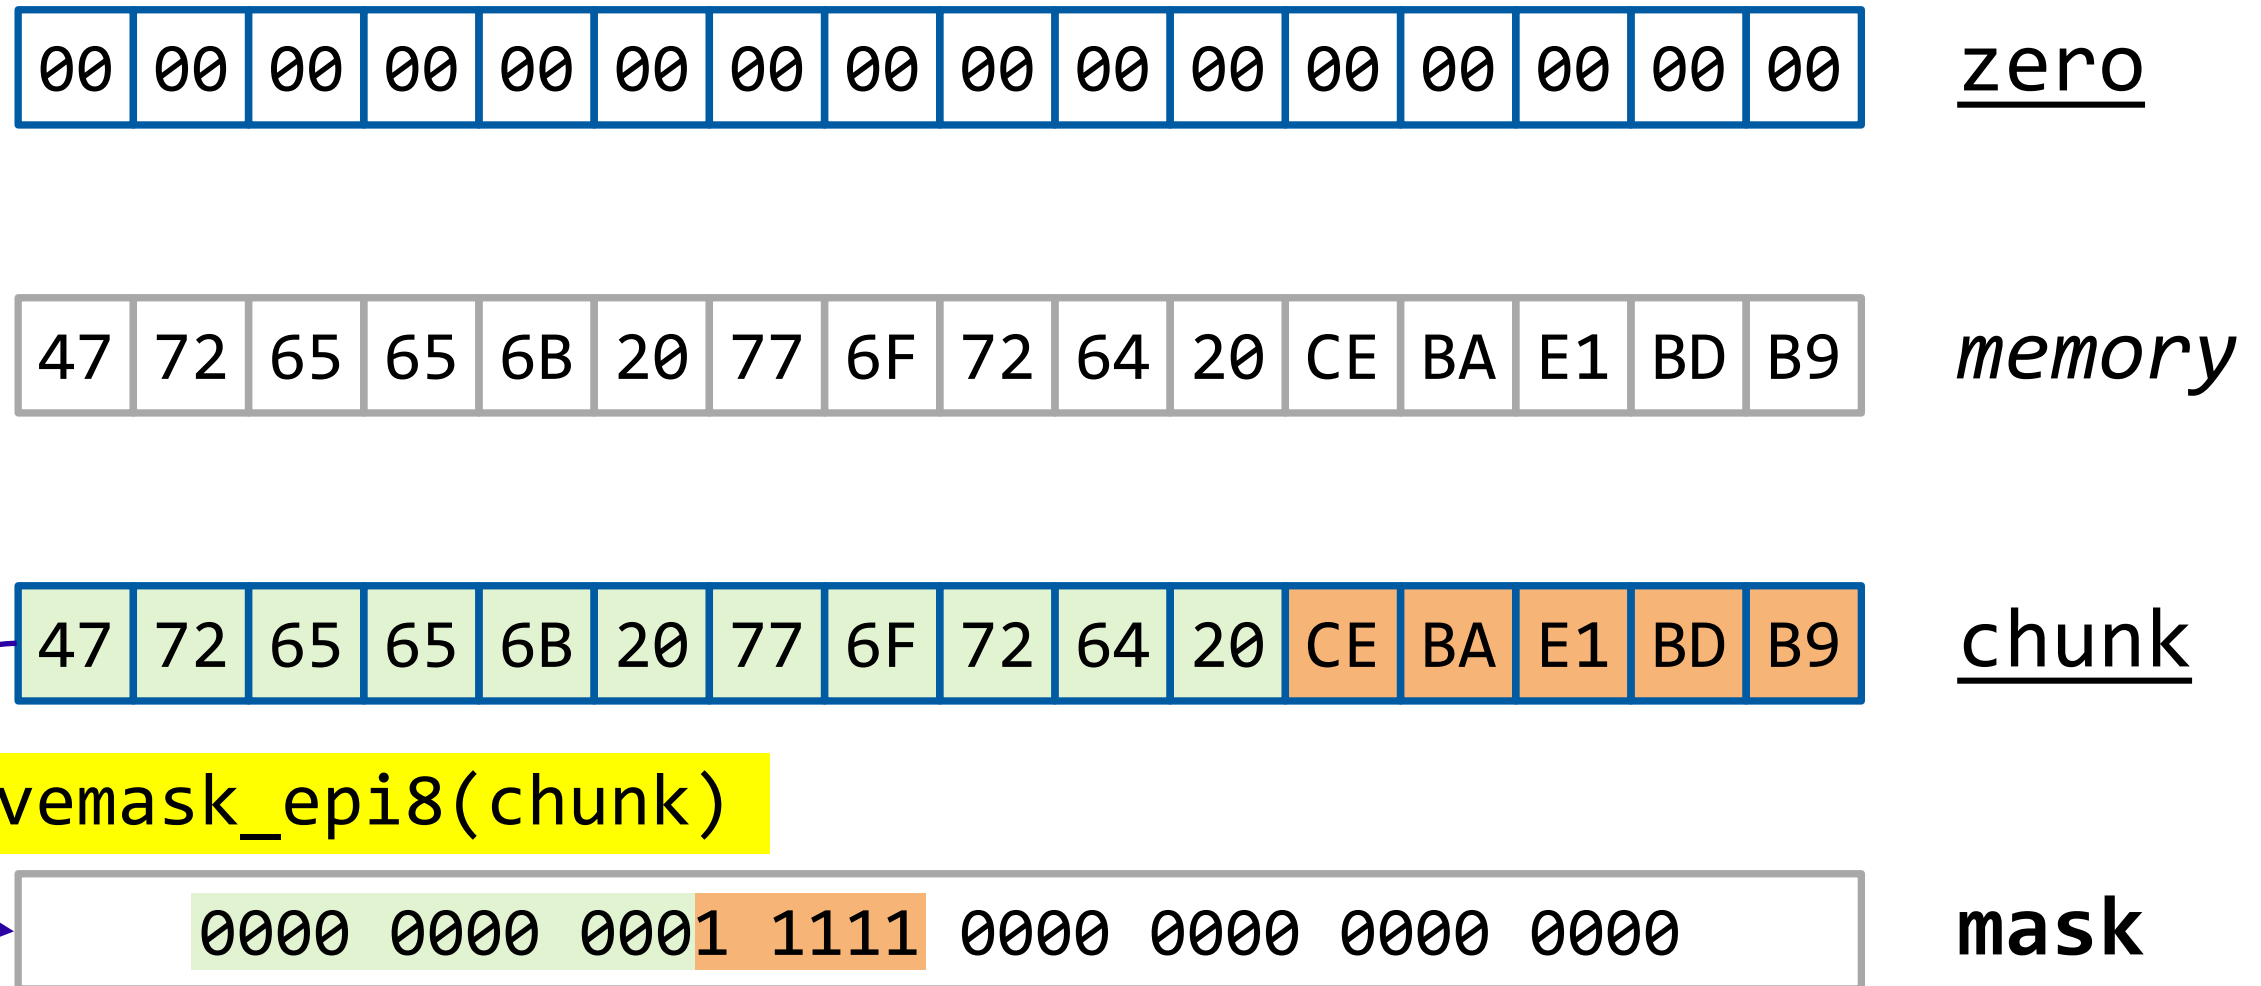
```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);              //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example



LSB



MSB

Converting ASCII Character Runs

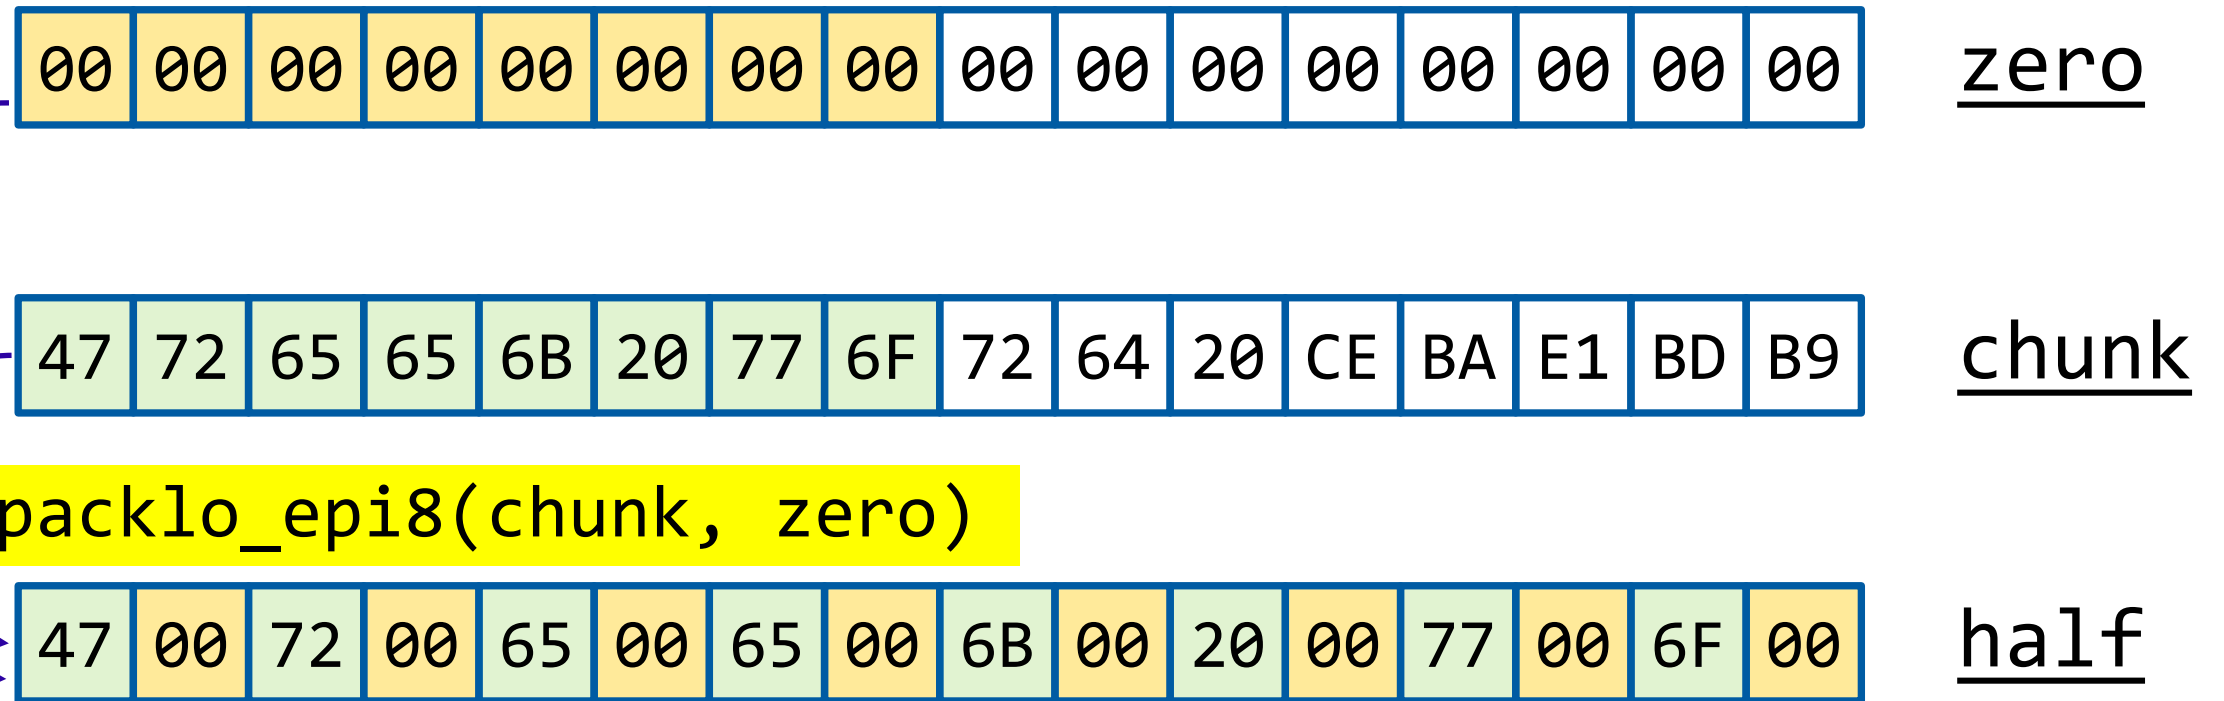
```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);              //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);  //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example



LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                    //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);             //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);        //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);        //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);      //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);        //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

47	00	72	00	65	00	65	00	6B	00	20	00	77	00	6F	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

```
qrtr = _mm_unpacklo_epi16(half, zero)
```

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);              //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

47	00	72	00	65	00	65	00	6B	00	20	00	77	00	6F	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

`_mm_storeu_si128((__m128i*) pDst, qrtr)`

pDst →

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

memory

LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                    //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);             //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);        //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);        //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);      //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);        //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

47	00	72	00	65	00	65	00	6B	00	20	00	77	00	6F	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

```
qrtr = _mm_unpackhi_epi16(half, zero)
```

6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    __m128i    chunk, half, qrtr, zero;           //- SSE "registers"
    int32_t    mask, incr;                        //- ASCII bit mask and advancement

    zero  = _mm_set1_epi8(0);                     //- Zero out the interleave register
    chunk = _mm_loadu_si128((__m128i const*) pSrc); //- Load a register with 8-bit values
    mask  = _mm_movemask_epi8(chunk);              //- Find octets with high bit set

    half = _mm_unpacklo_epi8(chunk, zero);         //- Unpack bytes 0-7 into 16-bit words
    qrtr  = _mm_unpacklo_epi16(half, zero);         //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);        //- Write to memory
    qrtr  = _mm_unpackhi_epi16(half, zero);         //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

47	00	72	00	65	00	65	00	6B	00	20	00	77	00	6F	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

```
_mm_storeu_si128((__m128i*) pDst + 4, qrtr)
```

pDst + 4 →

6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

memory

LSB



MSB

Converting ASCII Character Runs

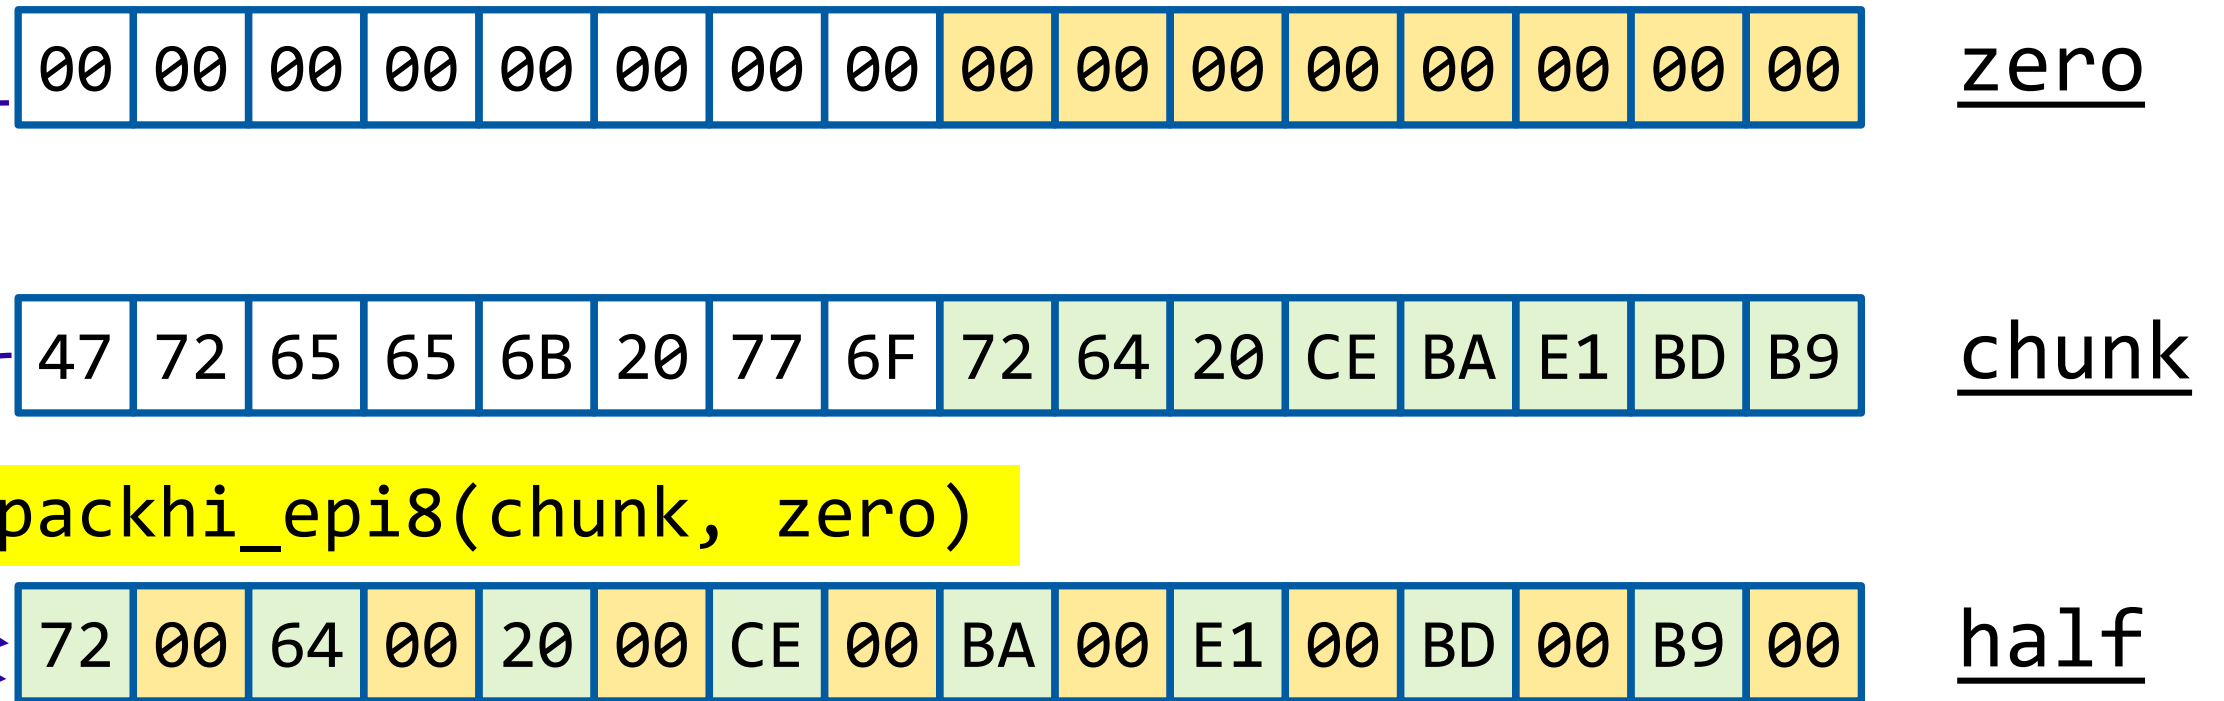
```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...

    half = _mm_unpacklo_epi8(chunk, zero);           //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);         //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);  //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);           //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr);  //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example



LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...

    half = _mm_unpacklo_epi8(chunk, zero);           //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);         //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);  //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);           //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr);  //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

72	00	64	00	20	00	CE	00	BA	00	E1	00	BD	00	B9	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

```
qrtr = _mm_unpacklo_epi16(half, zero)
```

72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...

    half = _mm_unpacklo_epi8(chunk, zero);           //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);         //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);   //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);           //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr);   //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr);  //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

72	00	64	00	20	00	CE	00	BA	00	E1	00	BD	00	B9	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

```
_mm_storeu_si128((__m128i*) pDst + 8, qrtr)
```

pDst + 8 →

72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

memory

LSB → MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...

    half = _mm_unpacklo_epi8(chunk, zero);           //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);         //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);  //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);           //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr);  //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr); //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

72	00	64	00	20	00	CE	00	BA	00	E1	00	BD	00	B9	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

```
qrtr = _mm_unpackhi_epi16(half, zero)
```

BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

LSB



MSB

Converting ASCII Character Runs

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...

    half = _mm_unpacklo_epi8(chunk, zero);           //- Unpack bytes 0-7 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 0-3 into 32-bit dwords
    _mm_storeu_si128((__m128i*) pDst, qrtr);         //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 4-7 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 4), qrtr);   //- Write to memory

    half = _mm_unpackhi_epi8(chunk, zero);           //- Unpack bytes 8-15 into 16-bit words
    qrtr = _mm_unpacklo_epi16(half, zero);           //- Unpack words 8-11 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 8), qrtr);   //- Write to memory
    qrtr = _mm_unpackhi_epi16(half, zero);           //- Unpack words 12-15 into 32-bit dwords
    _mm_storeu_si128((__m128i*) (pDst + 12), qrtr);  //- Write to memory

    ...
}
```

Converting ASCII Character Runs – SSE Example

00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

zero

72	00	64	00	20	00	CE	00	BA	00	E1	00	BD	00	B9	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

half

BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

qrtr

`_mm_storeu_si128((__m128i*) pDst + 12, qrtr)`

pDst + 12 →

BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

memory

LSB → MSB

Converting ASCII Character Runs – SSE Code

```
KEWB_FORCE_INLINE void
UtfUtils::ConvertAsciiWithSse(char8_t const*& pSrc, char32_t*& pDst) noexcept
{
    ...
    //- If no bits were set in the mask, then all 16 code units were ASCII.
    //
    if (mask == 0)
    {
        pSrc += 16;
        pDst += 16;
    }

    //- Otherwise, the number of trailing (low-order) zero bits in the mask is
    // the number of ASCII code units.
    //
    else
    {
        incr = GetTrailingZeros(mask);
        pSrc += incr;
        pDst += incr;
    }
}
```

Finding the Trailing Zero-Bit Count

```
#if defined KEWB_PLATFORM_LINUX  && (defined KEWB_COMPILER_CLANG  ||  defined KEWB_COMPILER_GCC)
```

```
    KEWB_FORCE_INLINE int32_t  
    UtfUtils::GetTrailingZeros(int32_t x) noexcept  
    {  
        return  __builtin_ctz((unsigned int) x);  
    }
```

```
#elif defined KEWB_PLATFORM_WINDOWS  &&  defined KEWB_COMPILER_MSVC
```

```
    KEWB_FORCE_INLINE int32_t  
    UtfUtils::GetTrailingZeros(int32_t x) noexcept  
    {  
        unsigned long    indx;  
        _BitScanForward(&indx, (unsigned long) x);  
        return (int32_t) indx;  
    }
```

```
#endif
```

Converting ASCII Character Runs – SSE Example

0000 0000 0001 1111 0000 0000 0000 0000

mask

`incr = GetTrailingZeros(mask)`

11 (eleven)

incr

pSrc →

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

pSrc += 11 →

pDst →

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00

pDst += 11 →

LSB

MSB

The SSE-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::SseConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < (pSrcEnd - sizeof(__m128i)))
    {
        if (*pSrc < 0x80)
        {
            ConvertAsciiWithSse(pSrc, pDst);
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    ...
}
```

Optimizing for ASCII with AVX

The AVX-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < (pSrcEnd - sizeof(__m128i)))
    {
        if (*pSrc < 0x80)
        {
            ConvertAsciiWithAvx(pSrc, pDst);
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    ...
}
```

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);             //- Store to memory

        mask   = _mm_movemask_epi8(chunk);                     //- Find bytes w/ high bit set
        chunk  = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half  = _mm256_cvtepu8_epi32(chunk);                    //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);             //- Store to memory

        mask  = _mm_movemask_epi8(chunk);                       //- Find bytes w/ high bit set
        chunk = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half  = _mm256_cvtepu8_epi32(chunk);                    //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```


Converting ASCII Character Runs – AVX Example

pSrc →

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

memory

```
chunk = _mm_loadu_si128((__m128i const*) pSrc)
```

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

LSB



MSB

The AVX-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half = _mm256_cvtepu8_epi32(chunk);                    //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);            //- Store to memory

        mask = _mm_movemask_epi8(chunk);                       //- Find bytes w/ high bit set
        chunk = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half = _mm256_cvtepu8_epi32(chunk);                    //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```

Converting ASCII Character Runs – AVX Example

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

`half = _mm256_cvtepu8_epi32(chunk)`

47	00	00	00	72	00	00	00	65	00	00	00	65	00	6F	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00

half

LSB



MSB

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half = _mm256_cvtepu8_epi32(chunk);                     //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);             //- Store to memory

        mask = _mm_movemask_epi8(chunk);                        //- Find bytes w/ high bit set
        chunk = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half = _mm256_cvtepu8_epi32(chunk);                     //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```

Converting ASCII Character Runs – AVX Example

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00

half

```
_mm256_storeu_si256((__m256i*) pDst, half)
```

pDst →

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00
00	00	00	00	00	00	00	00	00	00	00	00	00	00	00	00



Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

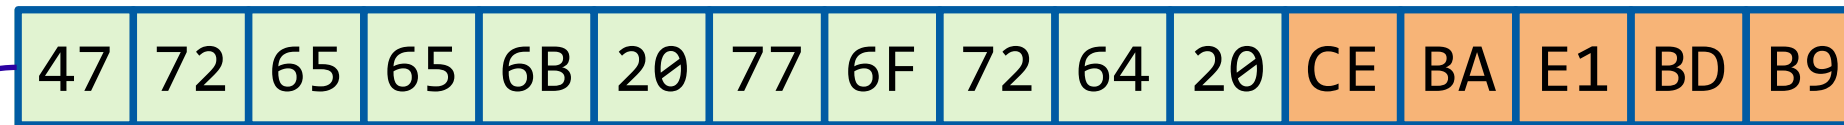
    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);            //- Store to memory

        mask   = _mm_movemask_epi8(chunk);                     //- Find bytes w/ high bit set
        chunk  = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory


        ...
    }
}
```

Converting ASCII Character Runs – AVX Example



chunk

```
mask = _mm_movemask_epi8(chunk)
```



mask

LSB



MSB

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half  = _mm256_cvtepu8_epi32(chunk);                   //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);            //- Store to memory

        mask  = _mm_movemask_epi8(chunk);                     //- Find bytes w/ high bit set
        chunk = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half  = _mm256_cvtepu8_epi32(chunk);                   //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);     //- Store to memory

        ...
    }
}
```


Converting ASCII Character Runs – AVX Example

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

```
chunk = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2))
```

72	64	20	CE	BA	E1	BD	B9	47	72	65	65	6B	20	77	6F
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

LSB



MSB

The AVX-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);            //- Store to memory

        mask   = _mm_movemask_epi8(chunk);                     //- Find bytes w/ high bit set
        chunk   = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half    = _mm256_cvtepu8_epi32(chunk);                 //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```

Converting ASCII Character Runs – AVX Example

72	64	20	CE	BA	E1	BD	B9	47	72	65	65	6B	20	77	6F
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

chunk

```
half = _mm256_cvtepu8_epi32(chunk)
```

72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00

half

LSB



MSB

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        int32_t      mask, incr;
        __m128i      chunk;                                     //- SSE register
        __m256i      half;                                     //- AVX register

        chunk = _mm_loadu_si128((__m128i const*) pSrc);         //- Load a register with bytes
        half   = _mm256_cvtepu8_epi32(chunk);                  //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) pDst, half);            //- Store to memory

        mask   = _mm_movemask_epi8(chunk);                     //- Find bytes w/ high bit set
        chunk   = _mm_shuffle_epi32(chunk, _MM_SHUFFLE(1,0,3,2)); //- Swap upper and lower
        half    = _mm256_cvtepu8_epi32(chunk);                 //- Zero-extend lower half
        _mm256_storeu_si256((__m256i*) (pDst + 8), half);      //- Store to memory

        ...
    }
}
```

Converting ASCII Character Runs – SSE Example

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9	<u>chunk</u>
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	--------------

72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00	<u>half</u>
BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00	

`_mm256_storeu_si256((__m256i*)(pDst + 8), half)`

pDst + 8 →

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00

LSB

MSB

Converting ASCII Character Runs – AVX Example

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    ...

    if (*pSrc < 0x80)
    {
        ...

        if (mask == 0)
        {
            pSrc += 16;
            pDst += 16;
        }
        else
        {
            incr = GetTrailingZeros(mask);
            pSrc += incr;
            pDst += incr;
        }
    }
    ...
}
```

Converting ASCII Character Runs – AVX Example

0000 0000 0001 1111 0000 0000 0000 0000

mask

`incr = GetTrailingZeros(mask)`

11 (eleven)

incr

pSrc →

47	72	65	65	6B	20	77	6F	72	64	20	CE	BA	E1	BD	B9
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

pSrc += 11 →

pDst →

47	00	00	00	72	00	00	00	65	00	00	00	65	00	00	00
6B	00	00	00	20	00	00	00	77	00	00	00	6F	00	00	00
72	00	00	00	64	00	00	00	20	00	00	00	CE	00	00	00
BA	00	00	00	E1	00	00	00	BD	00	00	00	B9	00	00	00

pDst += 11 →

LSB

MSB

The AVX-Optimized Conversion Algorithm (UTF-8 to UTF-32)

```
KEWB_ALIGN_FN std::ptrdiff_t
UtfUtils::AvxConvert(char8_t const* pSrc, char8_t const* pSrcEnd, char32_t* pDst) noexcept
{
    char32_t*    pDstOrig = pDst;
    char32_t     cdpt;

    while (pSrc < (pSrcEnd - sizeof(__m128i)))
    {
        if (*pSrc < 0x80)
        {
            ConvertAsciiWithAvx(pSrc, pDst);
        }
        else
        {
            if (Advance(pSrc, pSrcEnd, cdpt) != ERR)
            {
                *pDst++ = cdpt;
            }
            else
            {
                return -1;
            }
        }
    }
    ...
}
```


Testing and Benchmarks

- Ubuntu 18.04 Core i9
 - **GCC 9.3**, all code compiled with `-O3 -mavx2 -march=skylake`
 - **Clang 10.0.1**, all code compiled with `-O3 -mavx2 -march=skylake`
- Windows 10 Core i9
 - **Visual Studio 16.7.4**, all code compiled with `/O2 /Ob2 /Oi /Ot`

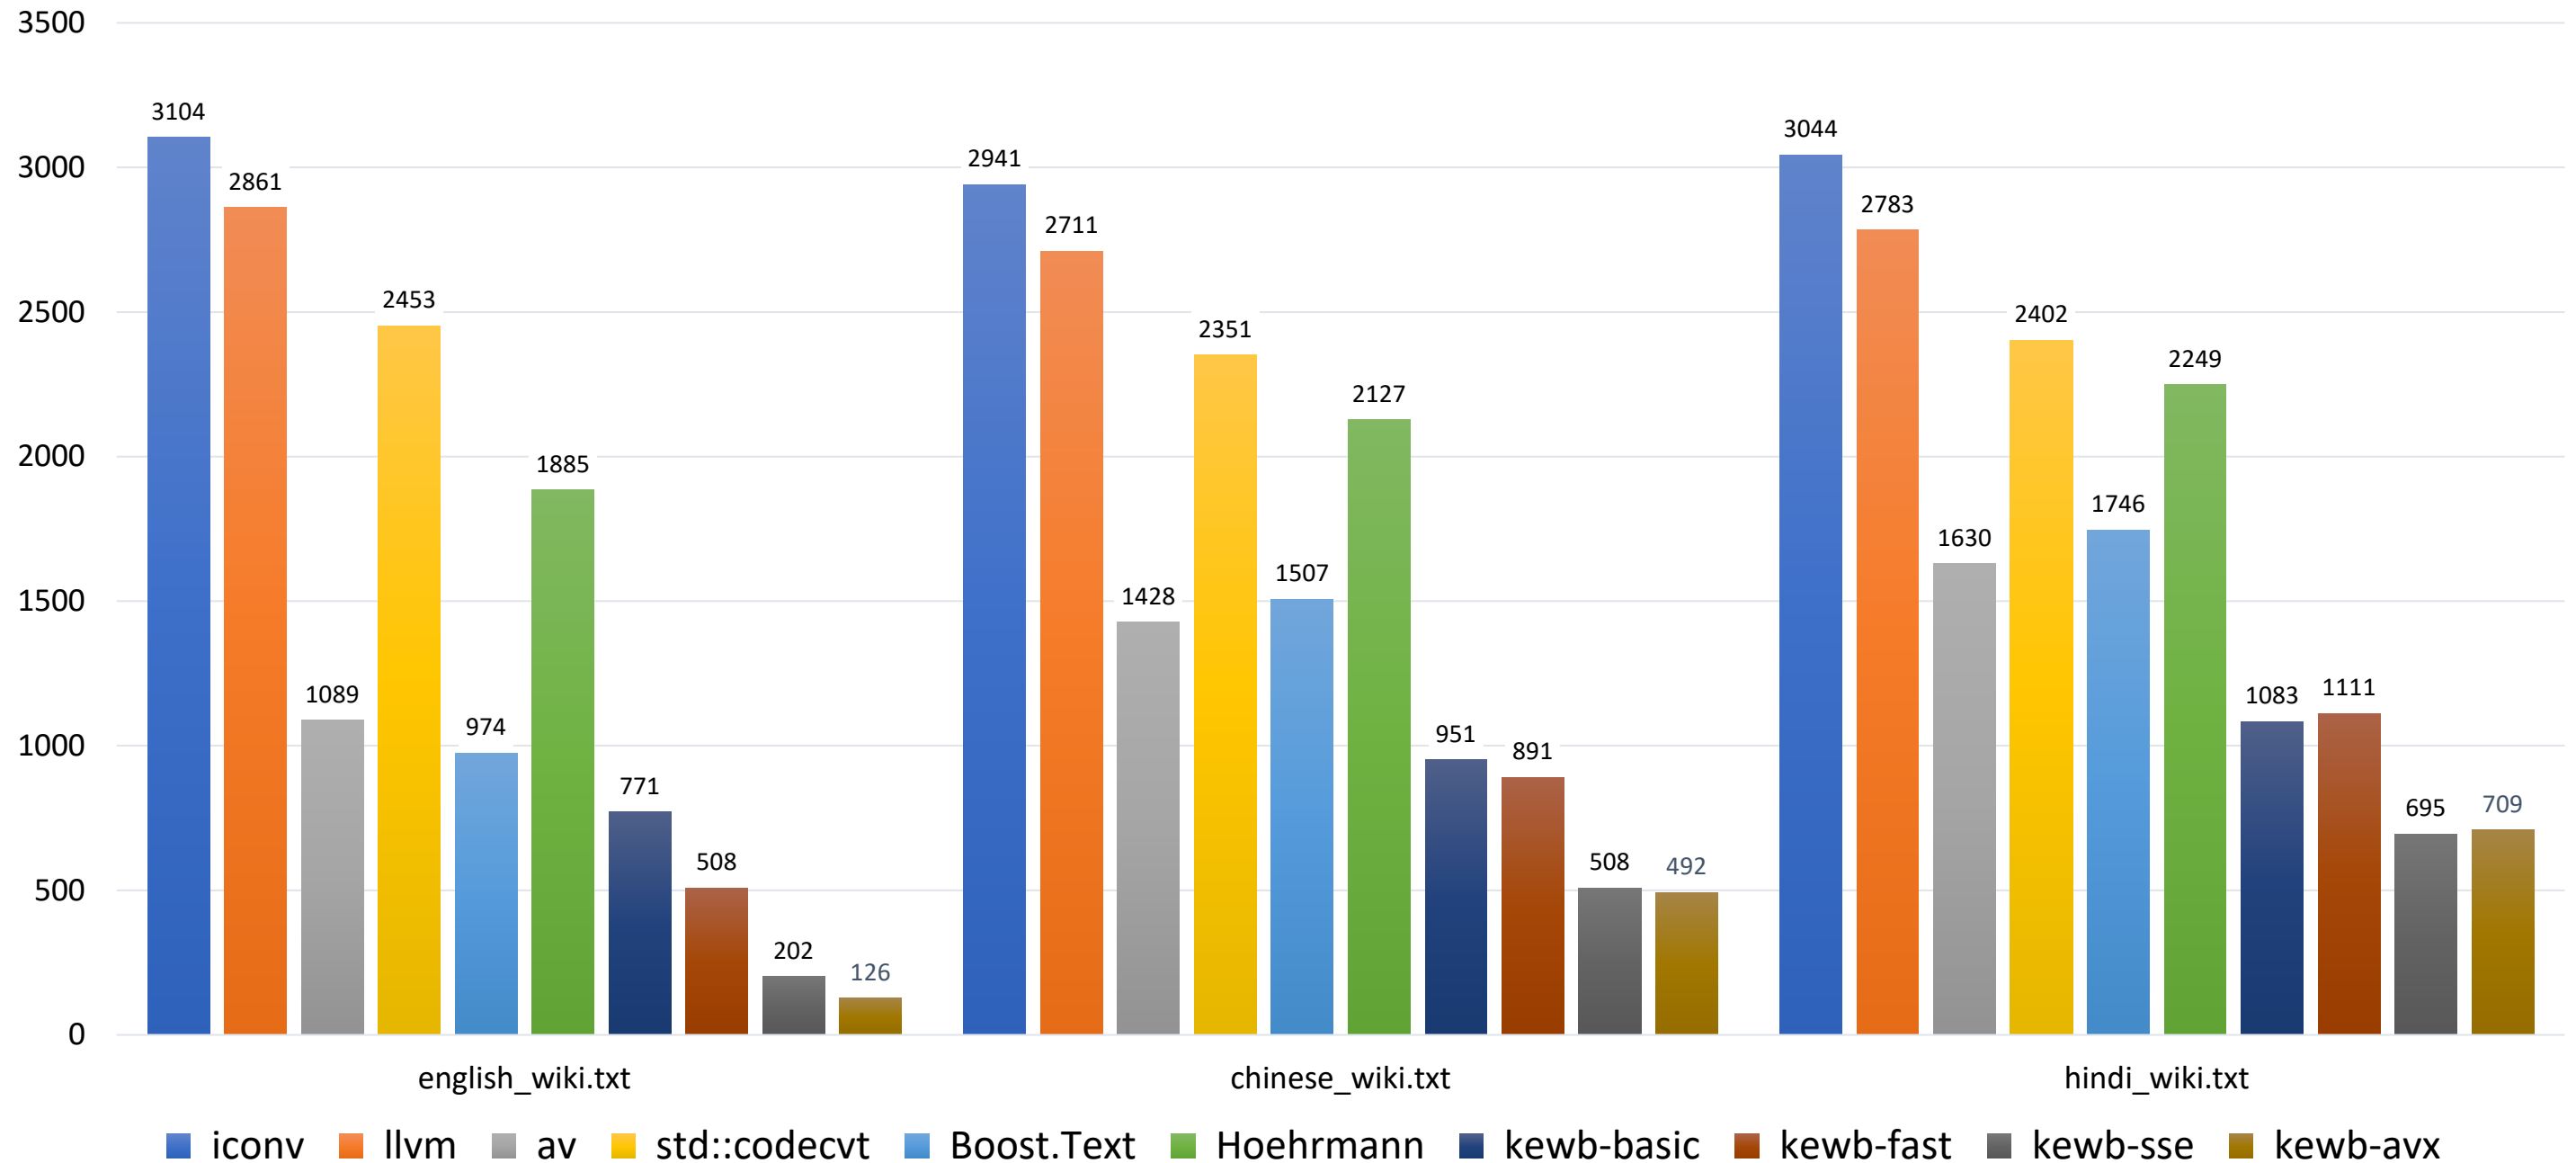
- Nine input files
 - english_wiki.txt
 - chinese_wiki.txt
 - hindi_wiki.txt
 - portuguese_wiki.txt
 - russian_wiki.txt
 - swedish_wiki.txt
 - stress_test_0.txt – 100K ASCII code points (100K code units)
 - stress_test_1.txt – 100K Chinese code points (300K code units)
 - stress_test_2.txt – 50K Chinese code points interleaved with 50K ASCII code points (200K code units)
- Taken directly from wikipedia.org

Testing Methodology – Reference Libraries

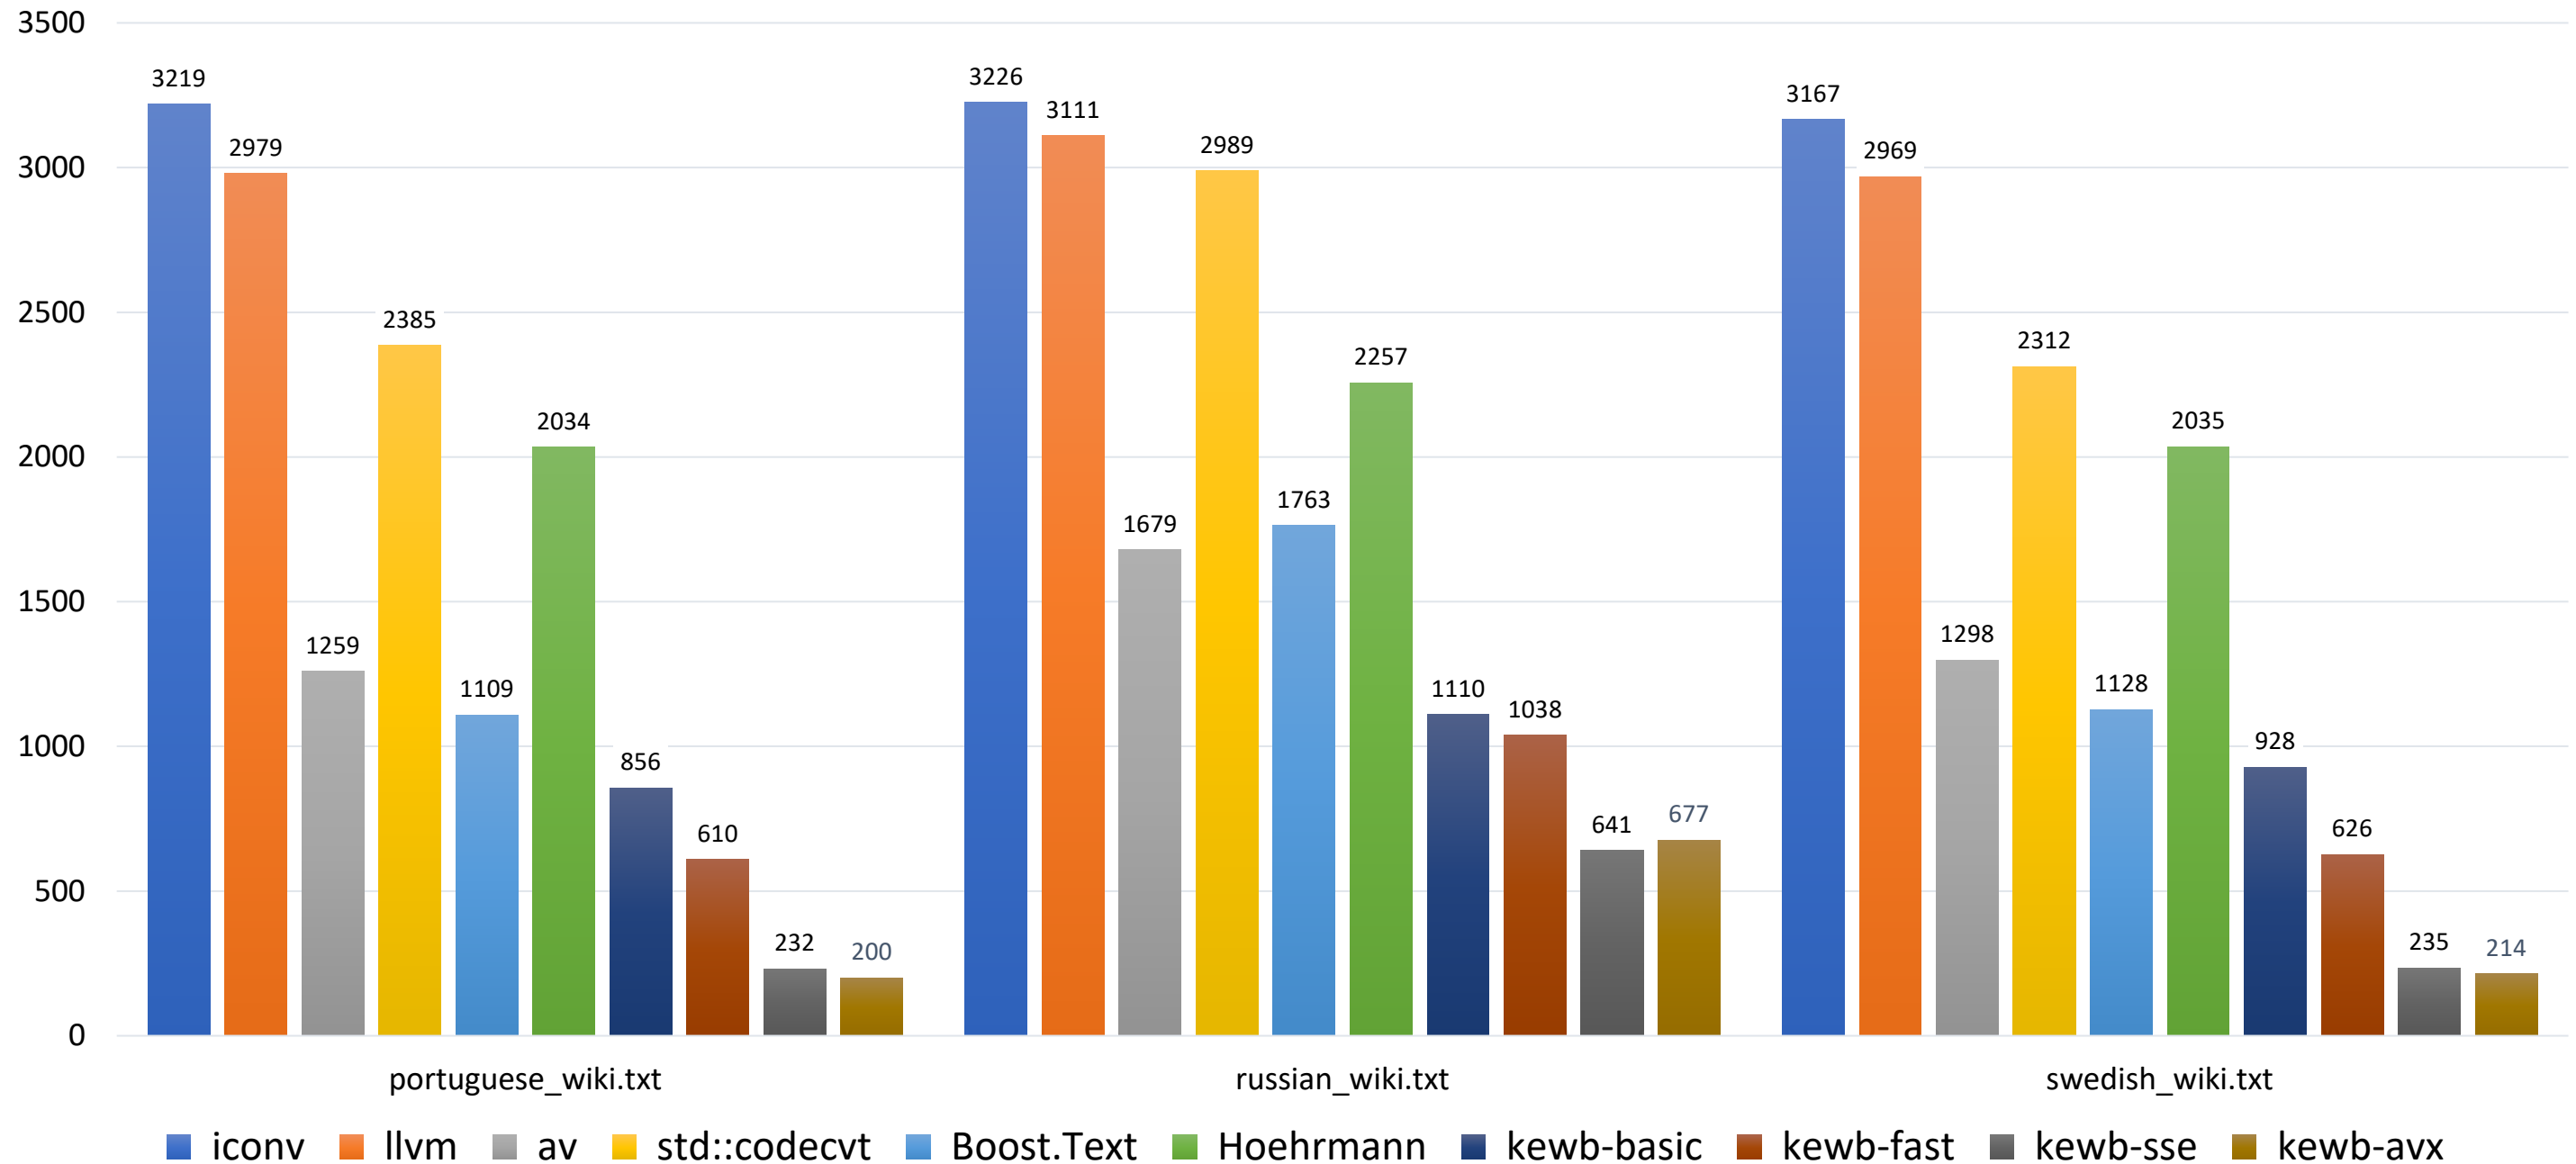
- `iconv` – GNU libiconv, used here as the “gold standard”
- `LLVM` – UTF conversion functions from the LLVM distribution
- `AV` – UTF-8 to UTF-32 conversion by Alexey Vatchenko
- `std::codecvt` – Standard library’s UTF conversion
- `Boost.Text` – Iterator-based interface to UTF conversion by Zach Laine
- `BH` – Alternative DFA-based conversion by Bjoern Hoehrmann

- Timings for each file were obtained by:
 1. Reading the input file
 2. Creating an oversized output buffer
 3. Starting the timer
 4. Entering the timing loop
 5. Performing conversion of the input buffer multiple times
 - The number of repetitions was such that 1GB of input text was processed
 6. Exiting the timing loop
 7. Stopping the timer
 8. Collecting and collating results
- To pass, a library's result had to agree with `iconv()`


Conversion Time (msec)



Conversion Time (msec)

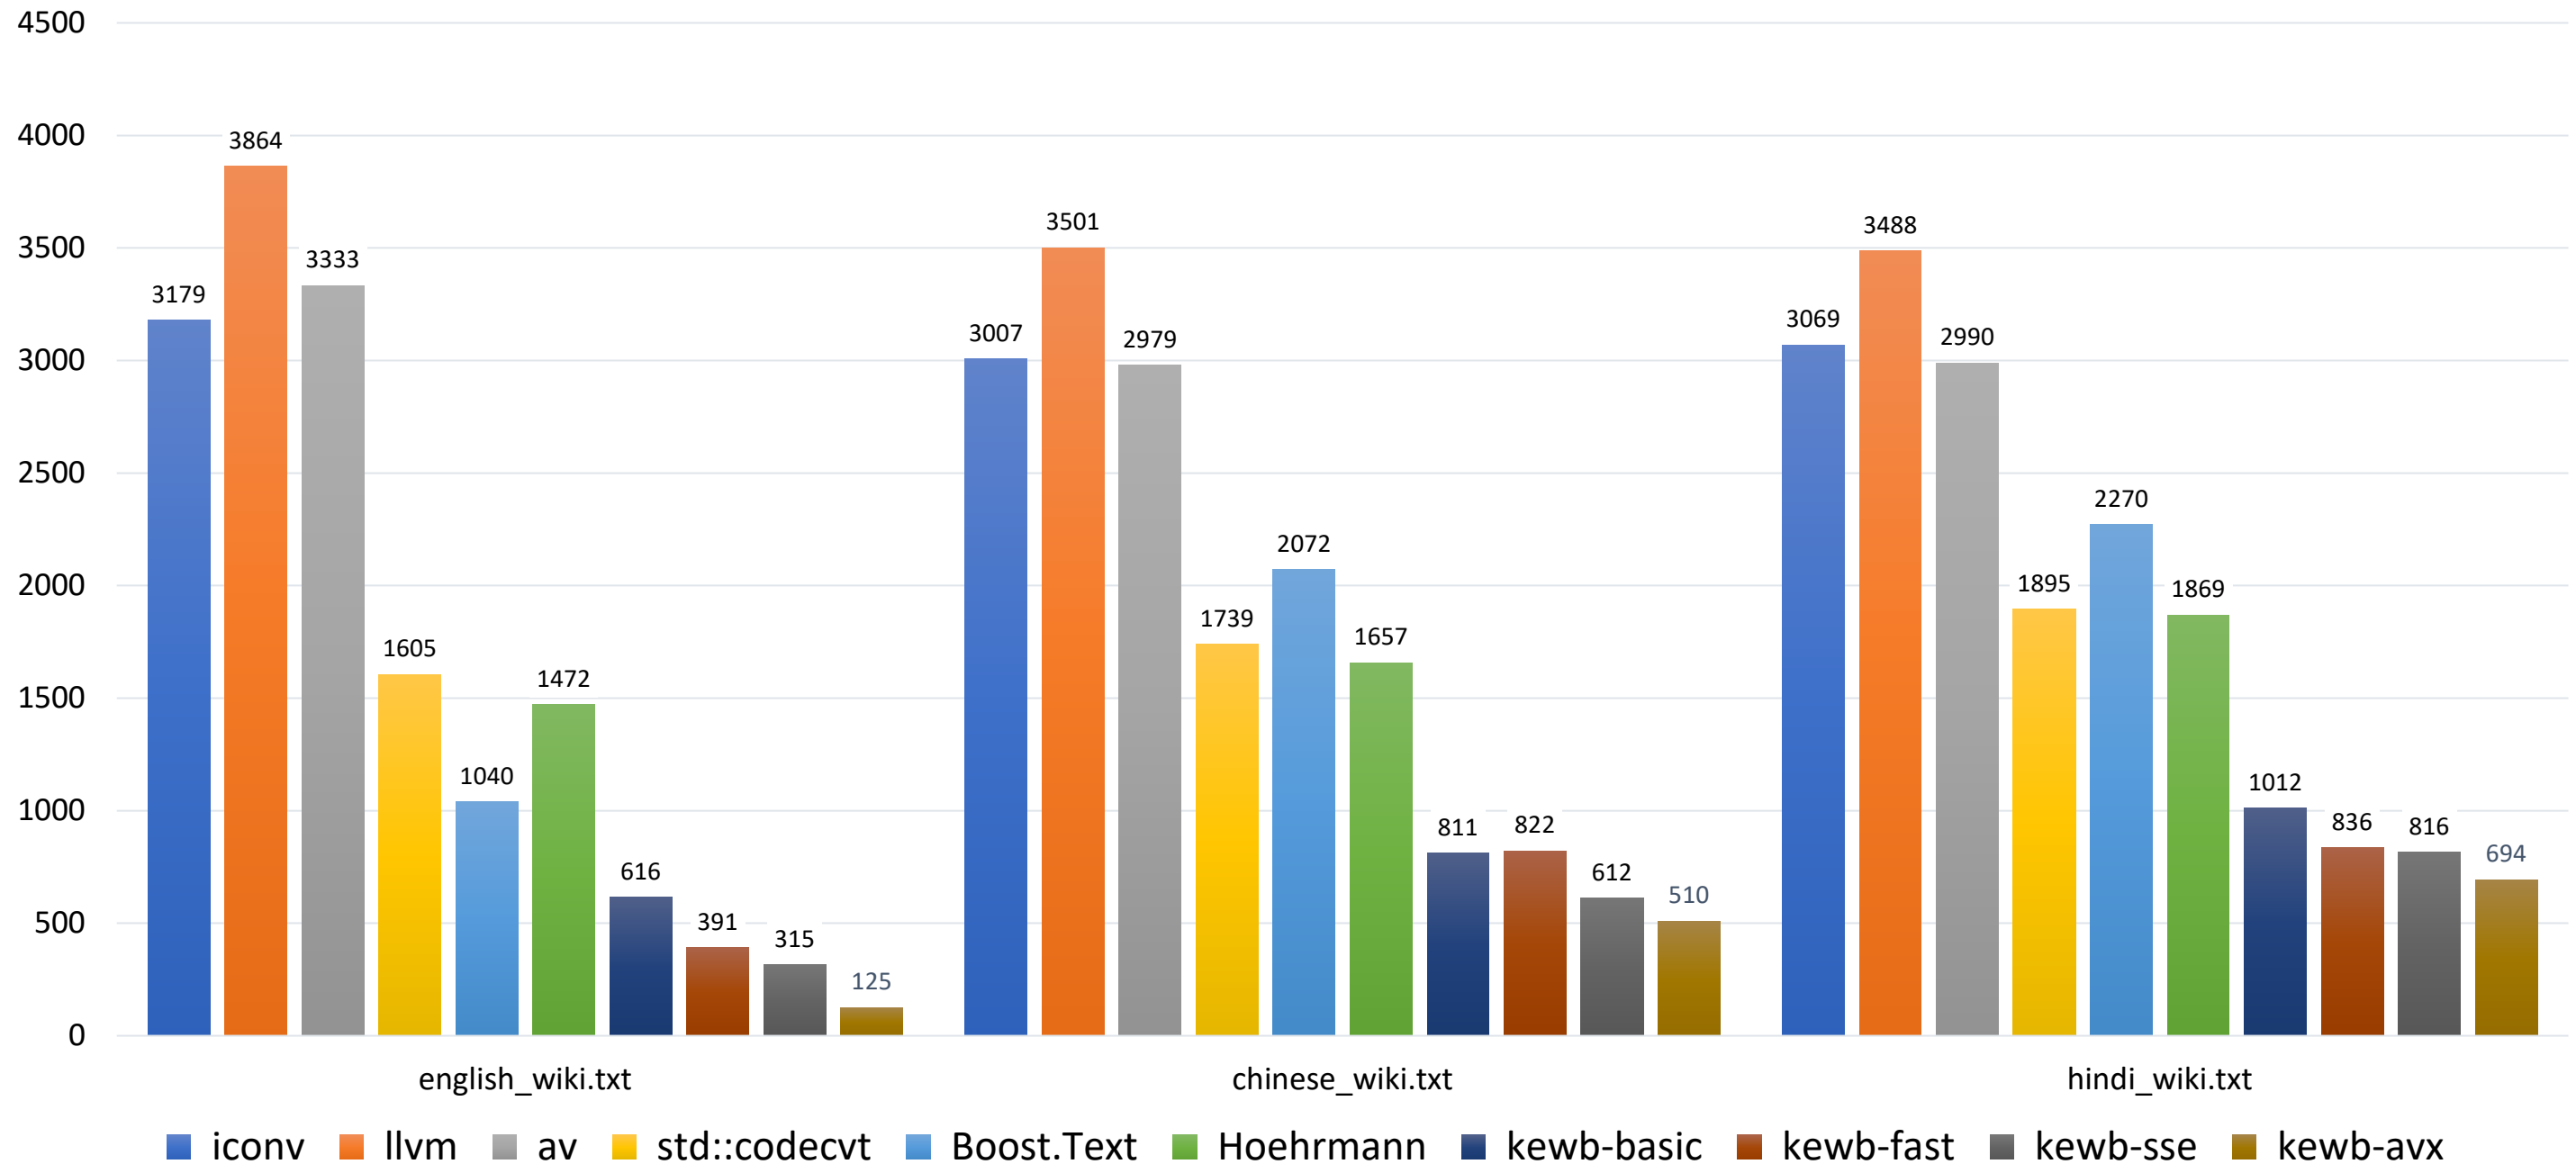


Conversion Time (msec)



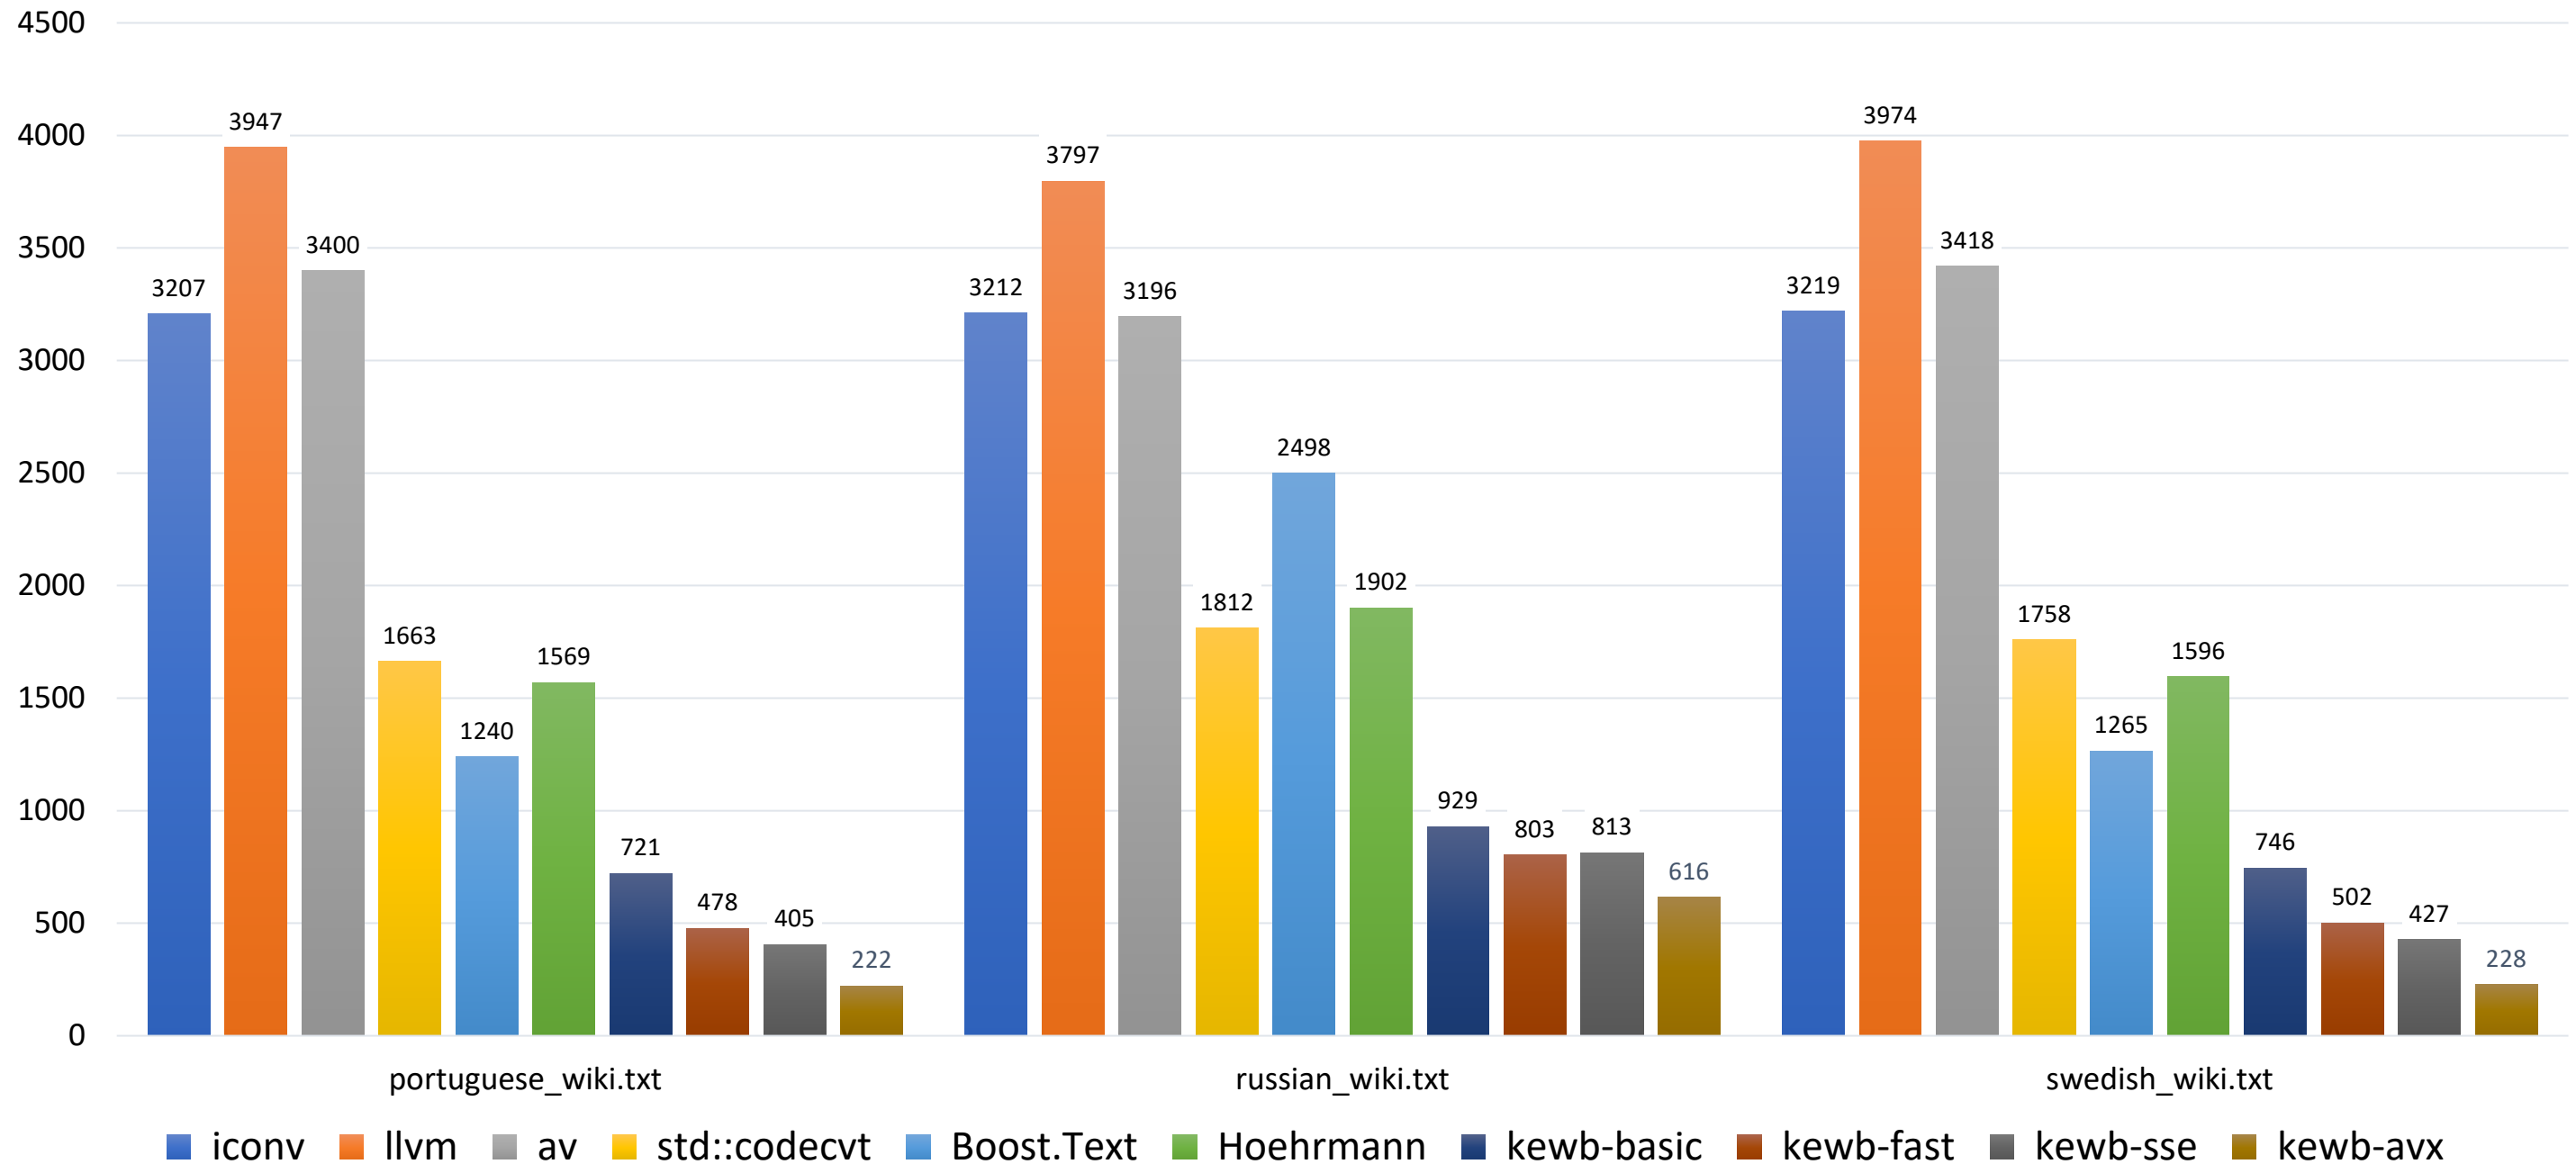
Clang 10.0.1 – Ubuntu 18.04 – Core i9 – UTF-32

Conversion Time (msec)



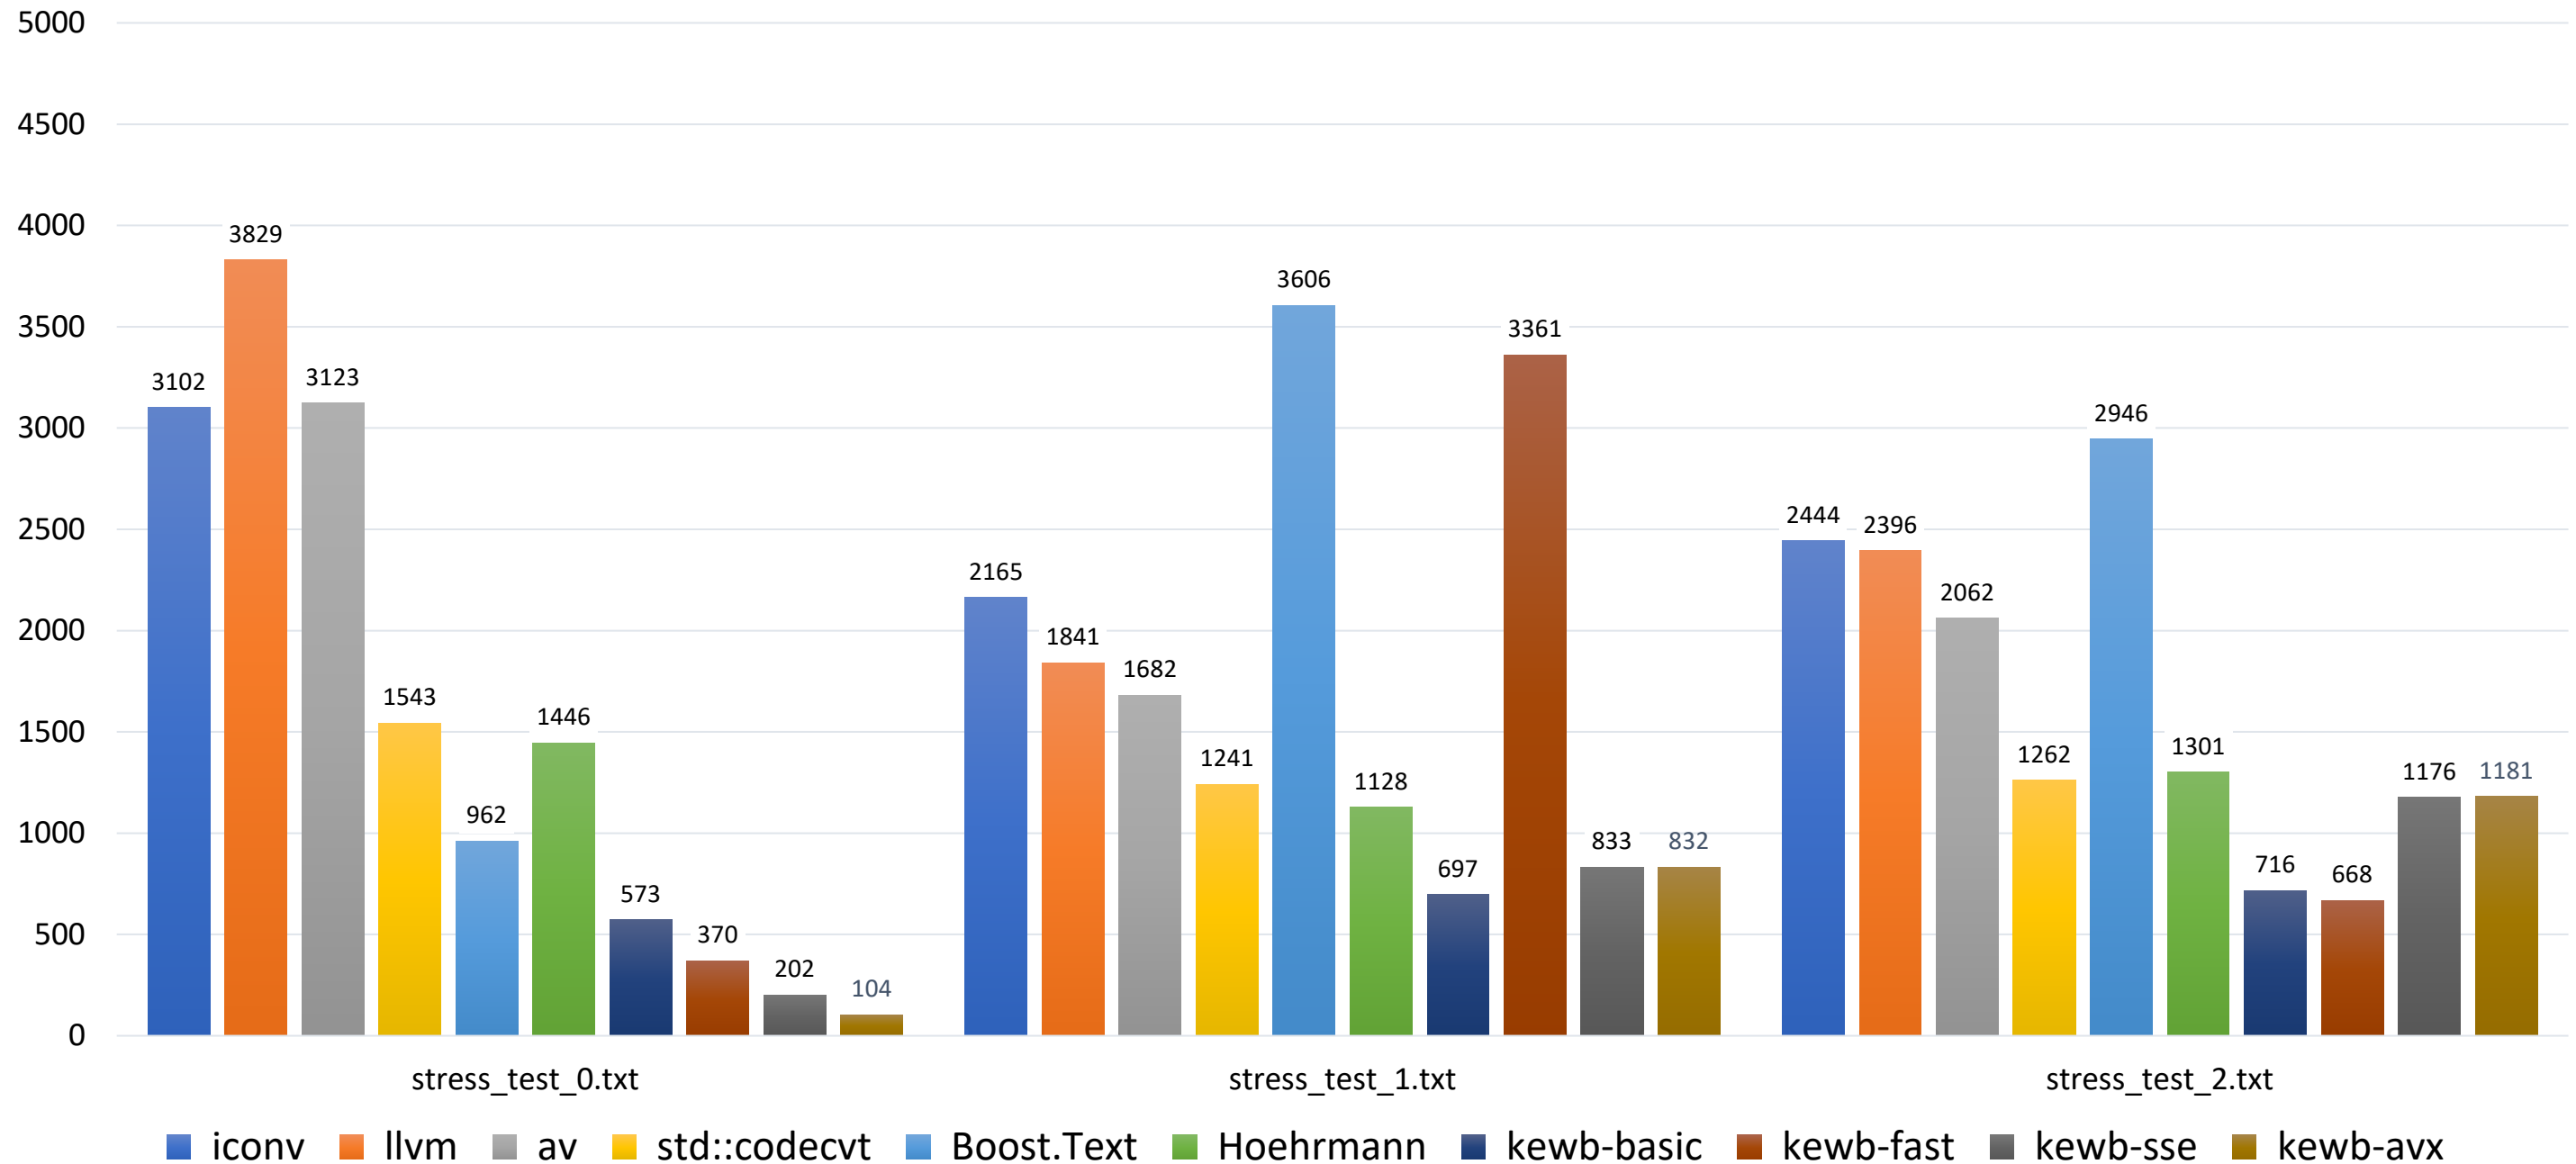
Clang 10.0.1 – Ubuntu 18.04 – Core i9 – UTF-32

Conversion Time (msec)

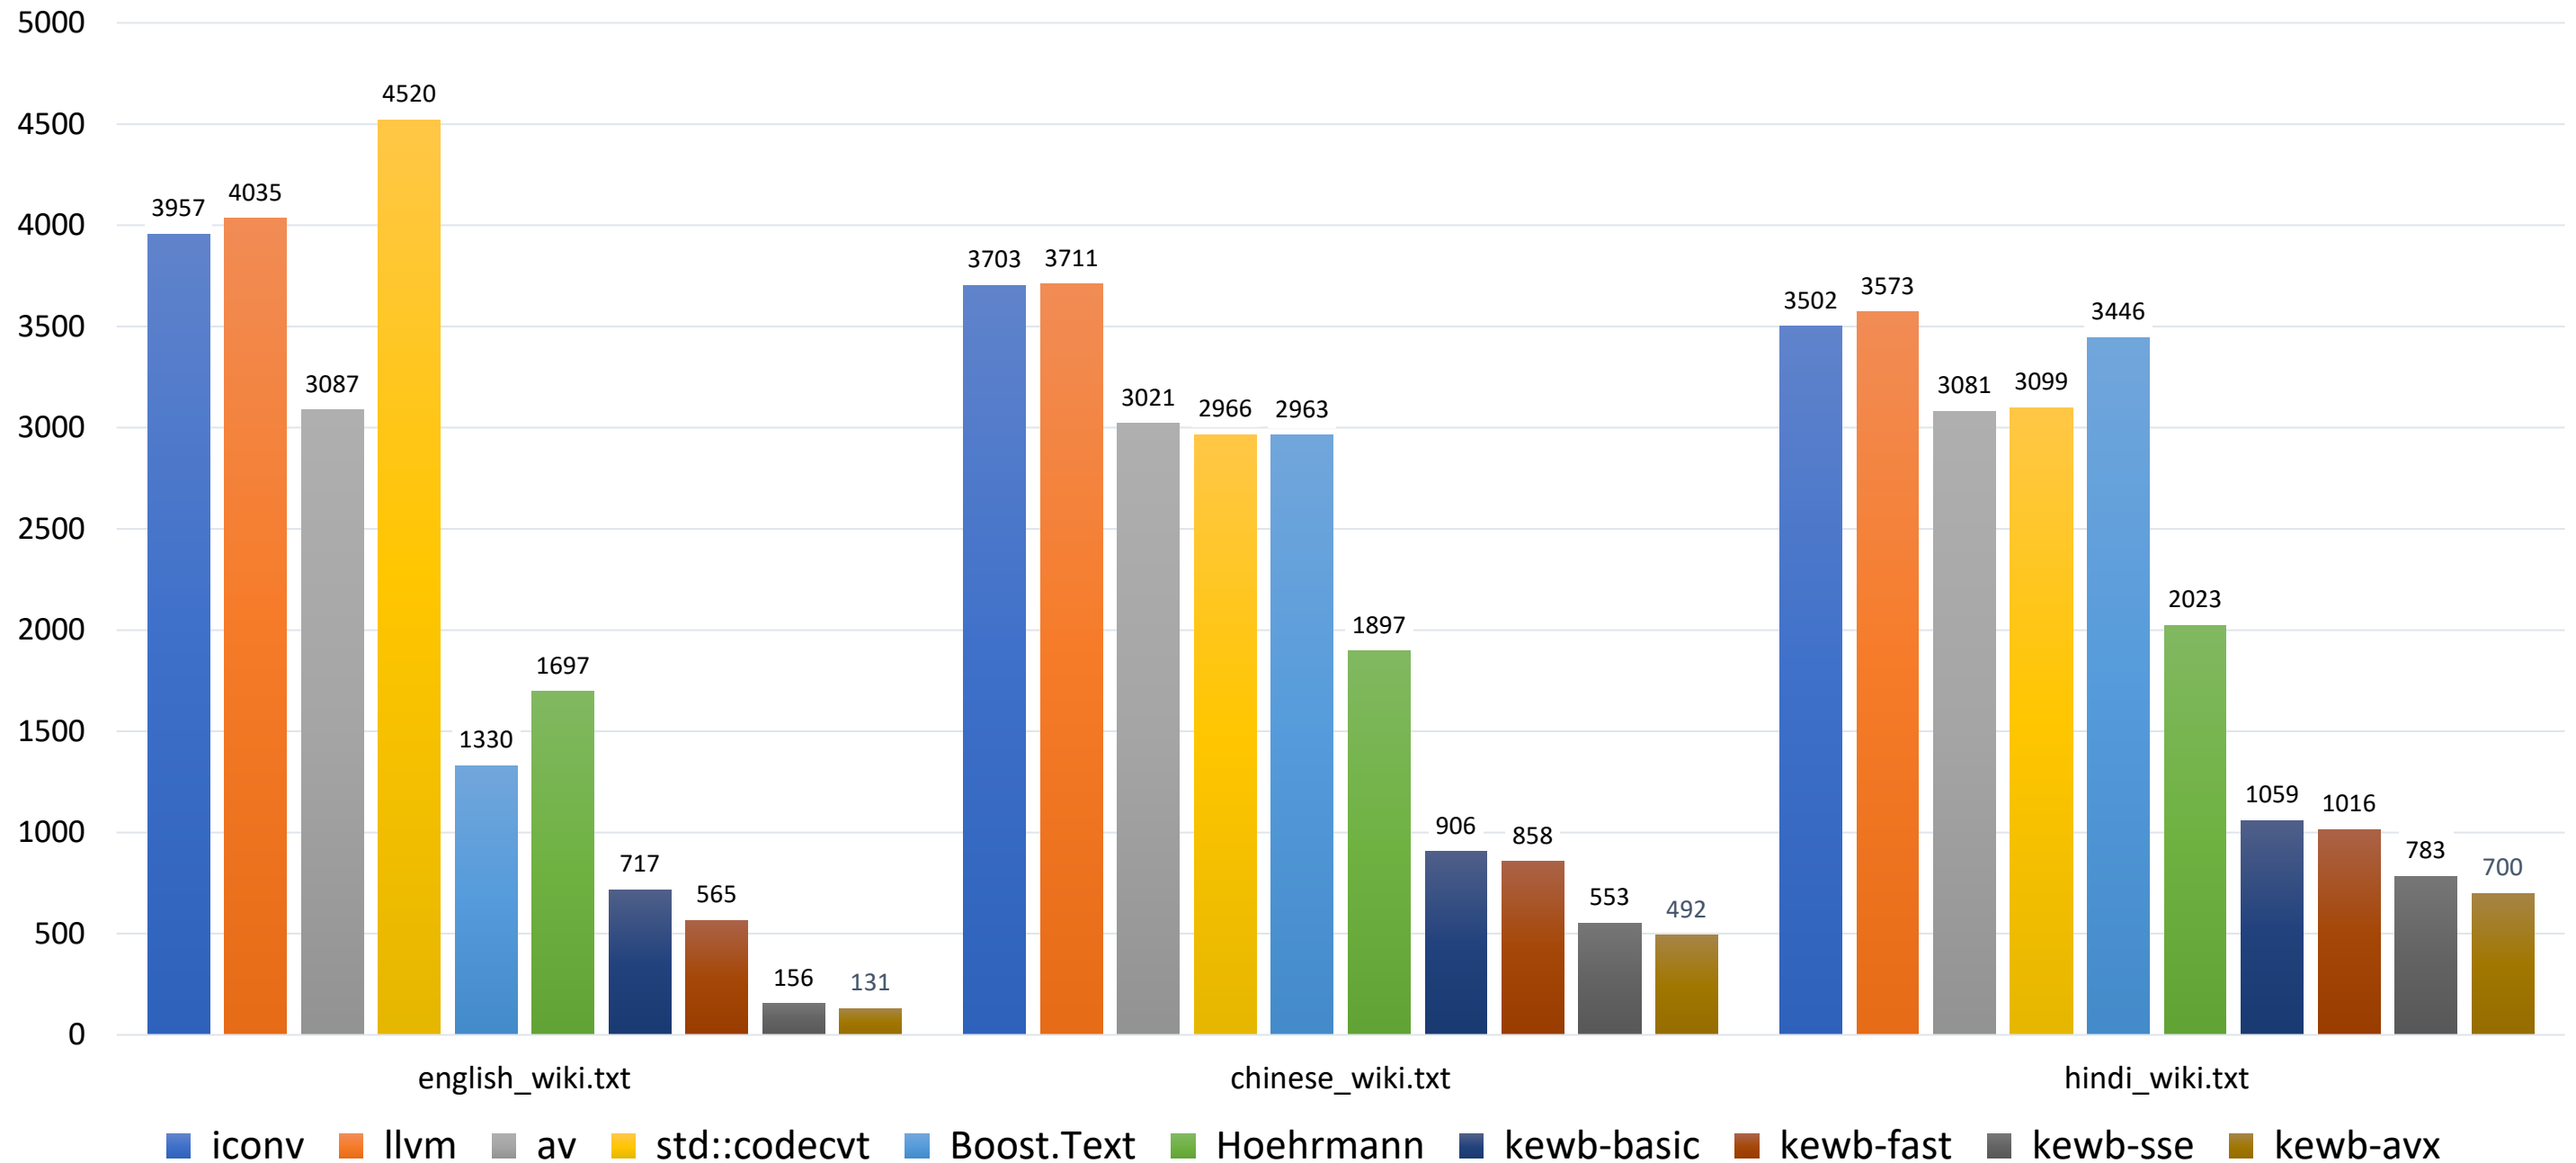


Clang 10.0.1 – Ubuntu 18.04 – Core i9 – UTF-32

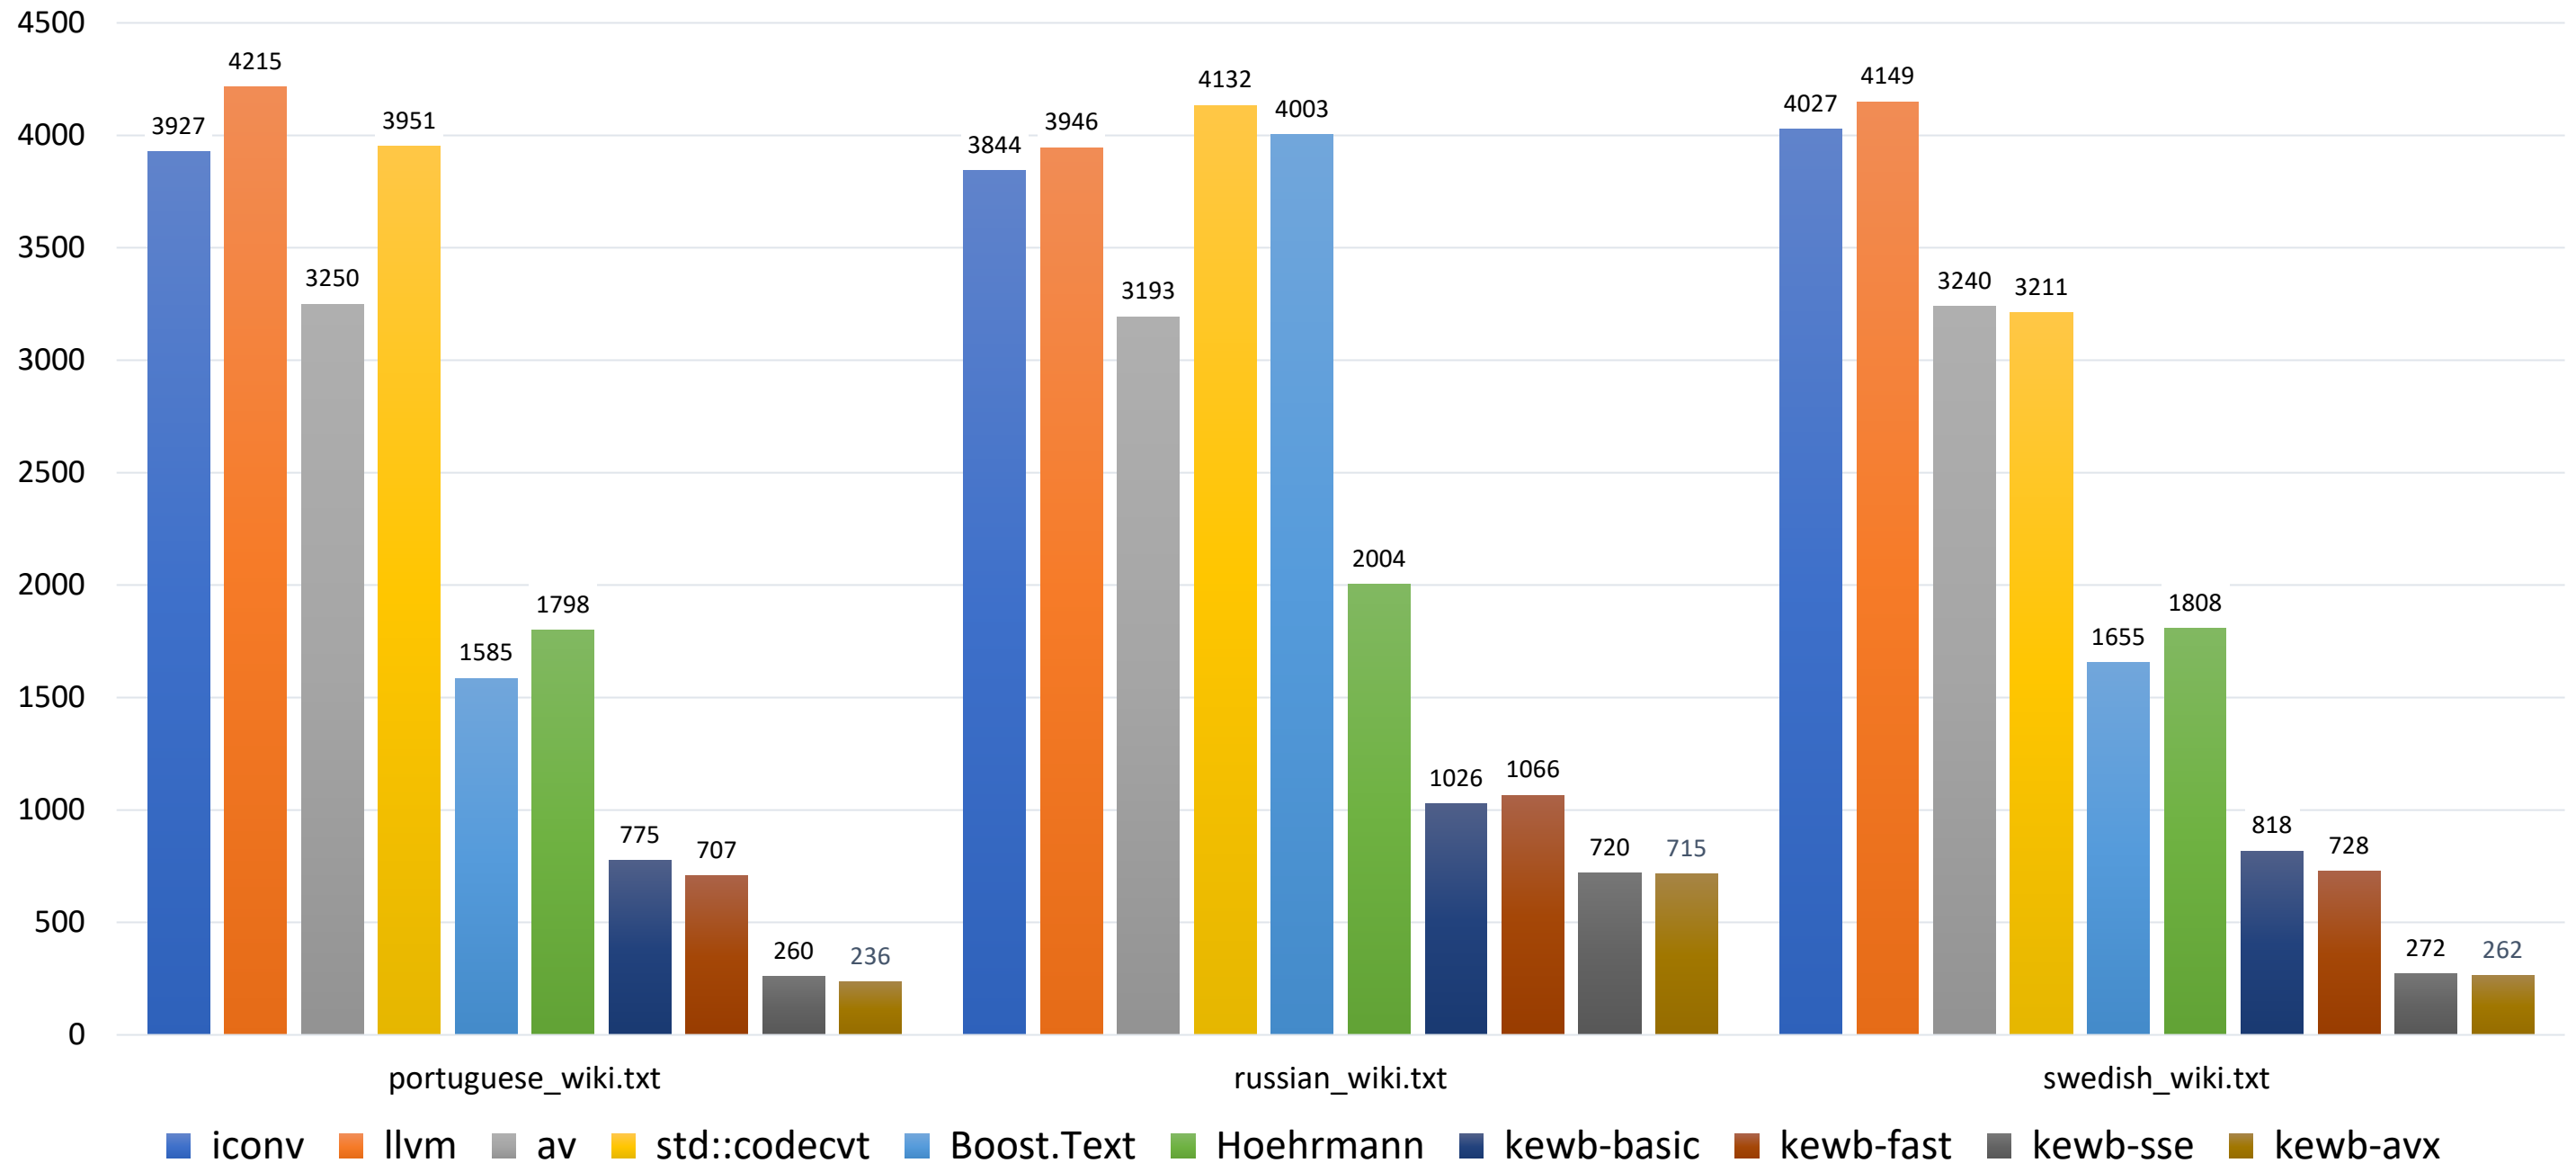
Conversion Time (msec)



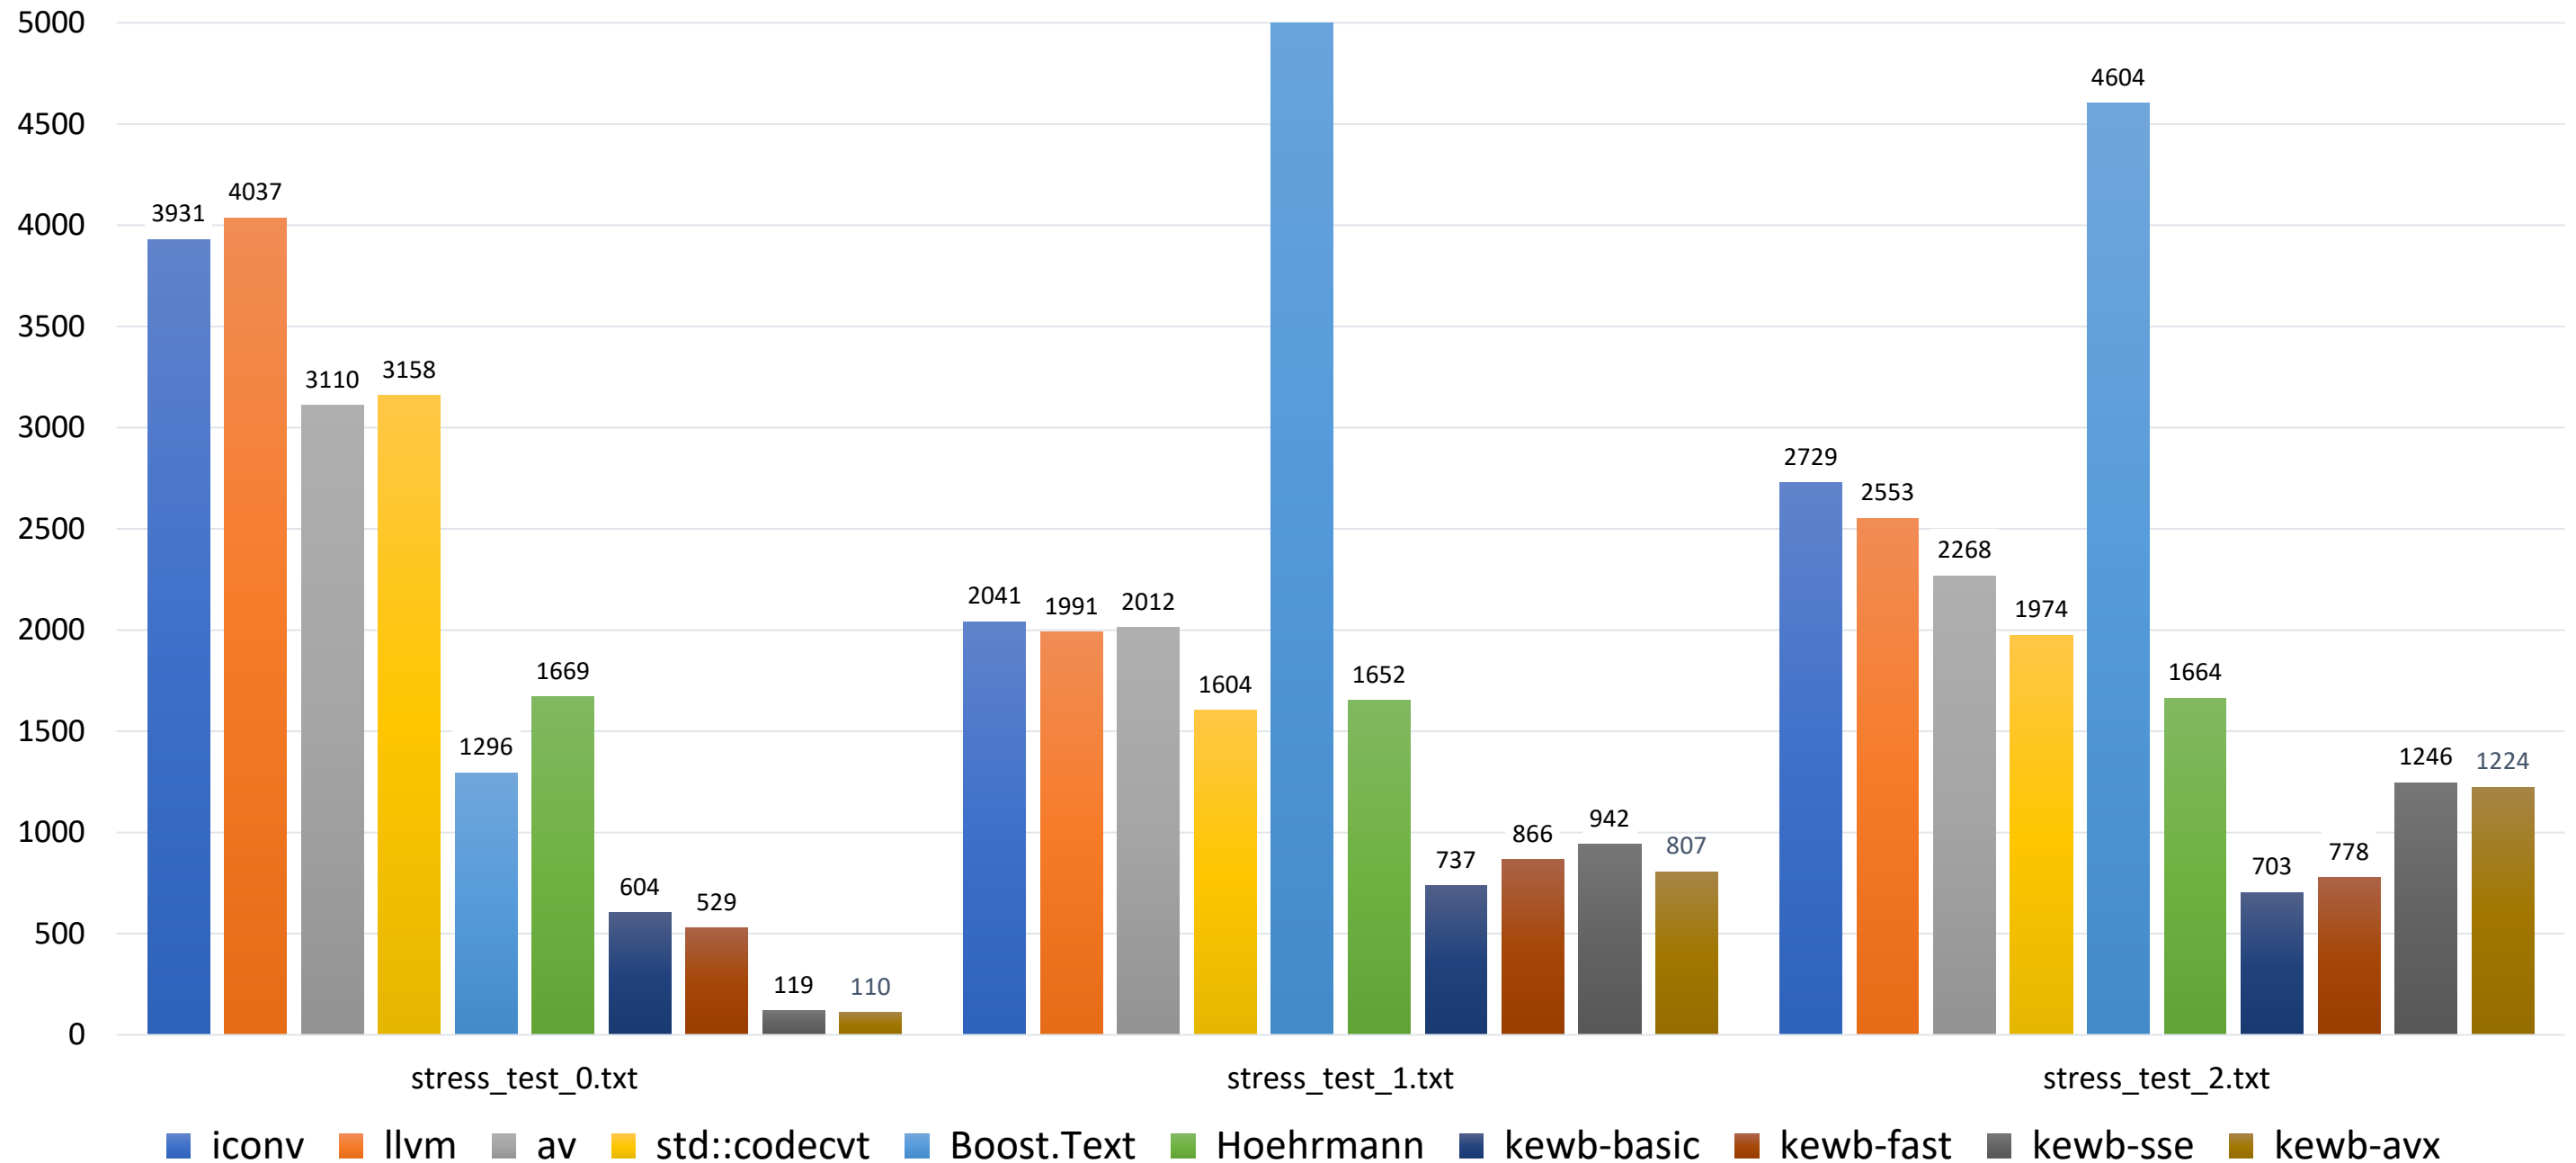
Conversion Time (msec)



Conversion Time (msec)



Conversion Time (msec)



Thank You for Attending!

Talk: github.com/BobSteagall/CppCon2020

Blog: bobsteagall.com