

A Multithreaded, Transaction-Based Locking Strategy for Containers

Bob Steagall
CppCon 2020



- Sharing a container among multiple threads
- A motivating problem
- Some possible solutions
- A solution based on strict timestamp ordering (STO)
- Testing the STO-based solution
- Summary

- Sometimes a container must be shared between threads
- We desire to avoid race conditions during writes
 - Assume elements are themselves unsynchronized – i.e., susceptible to races
 - Exactly one thread may update an element at any given time
 - No other thread may read an element while the writer is updating it
 - More than one thread may read an element when no update is in progress
- We now have a wealth of concurrency tools at our disposal
 - Writing multi-threaded applications is easier (?) than ever

Sharing a Container – Avoiding Race Conditions


- If all threads are readers...
 - No locking is required
- If the number of reads is much larger than the number of writes...
 - We might be able to use a readers/writer lock (`std::shared_mutex`)
- What about the case where most operations are writes?
 - A per-element mutex strategy might work...
 - ... if a given write operation requires locking exactly one element

Sharing a Container – Avoiding Race Conditions

- What about the case where **all** operations are writes...
-- *and* --
- Each element E to be updated is related to a set R_E of other elements
 - Let's call this set ***E 's related group***
 - Let's define $U_E = \{E \cup R_E\}$ and call it ***E 's update group***

Sharing a Container – Avoiding Race Conditions

T_n



$$E = \{e_8\}$$

$$R_E = \{e_1, e_2, e_5, e_{11}\}$$

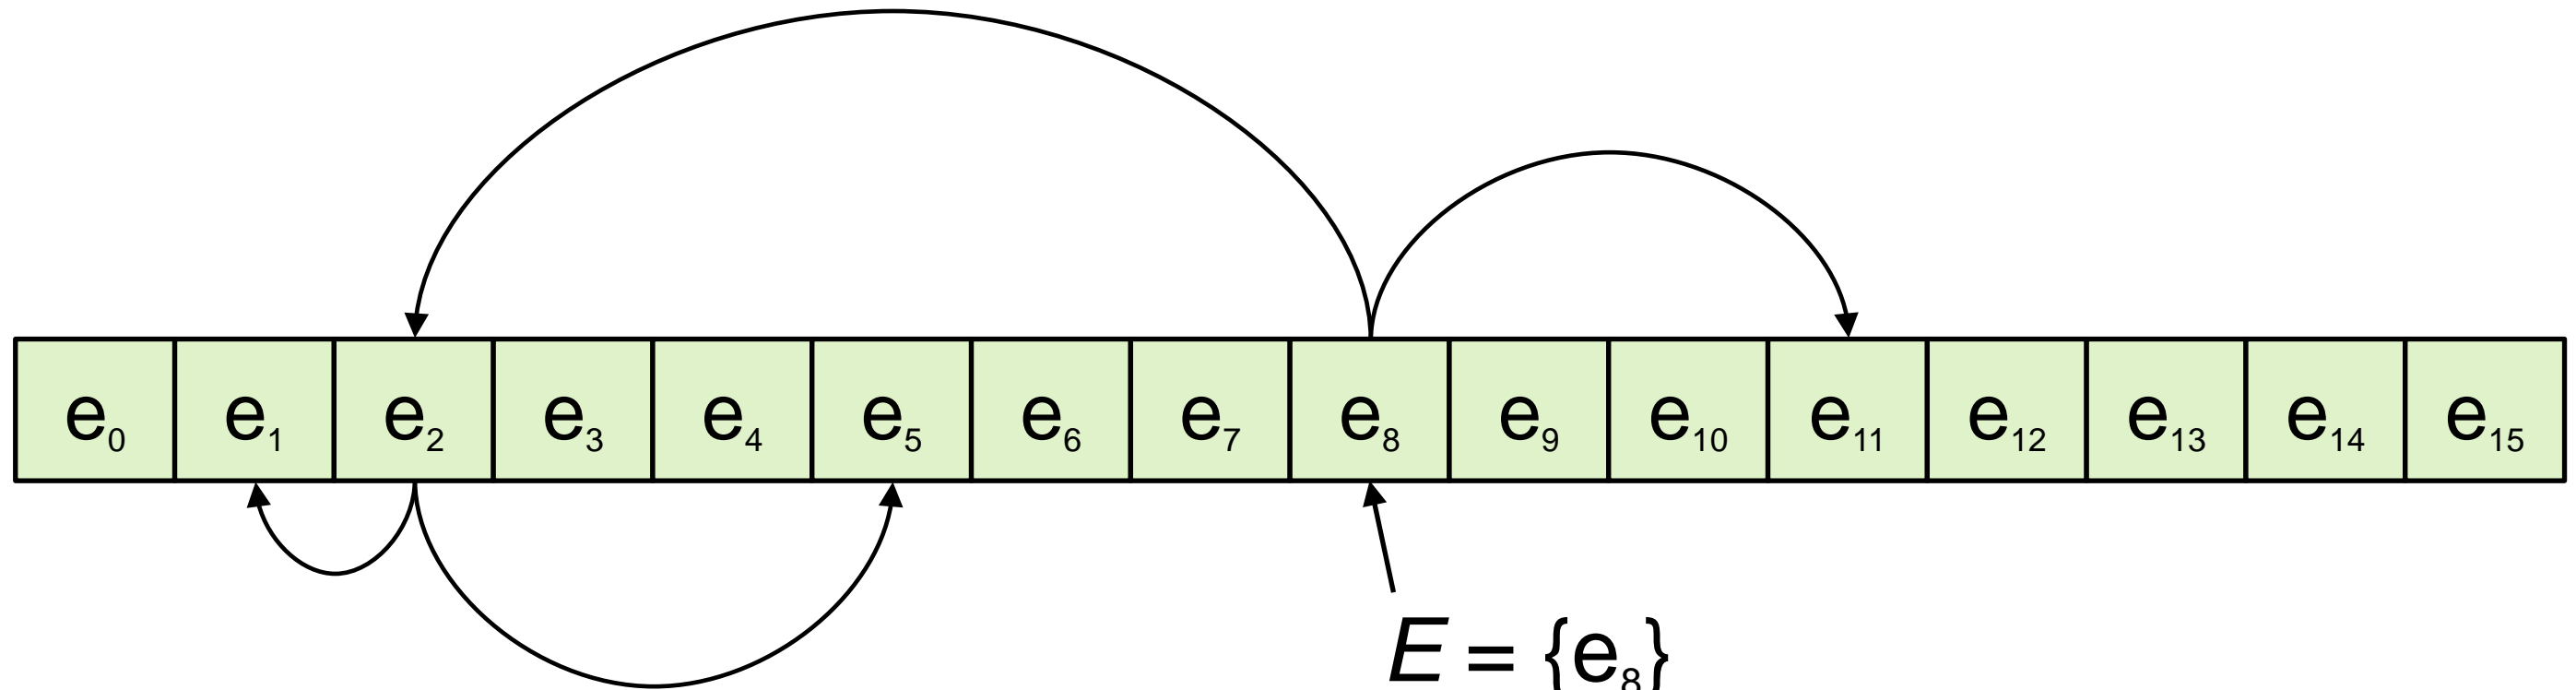
$$U_E = \{e_1, e_2, e_5, e_8, e_{11}\}$$

Sharing a Container – Avoiding Race Conditions

- What about the case where **all** operations are writes...
-- *and* --
- Each element E to be updated is related to a set R_E of other elements
 - Let's call this set **E 's related group**
 - Let's define $U_E = \{E \cup R_E\}$ and call it **E 's update group**-- *and* --
- One or more elements in R_E must also be updated at the same time as, and consistently, with E
-- *and* --
- The number of elements in R_E to be updated varies
-- *and* --
- The membership of R_E varies?

Sharing a Container – Avoiding Race Conditions

T_n



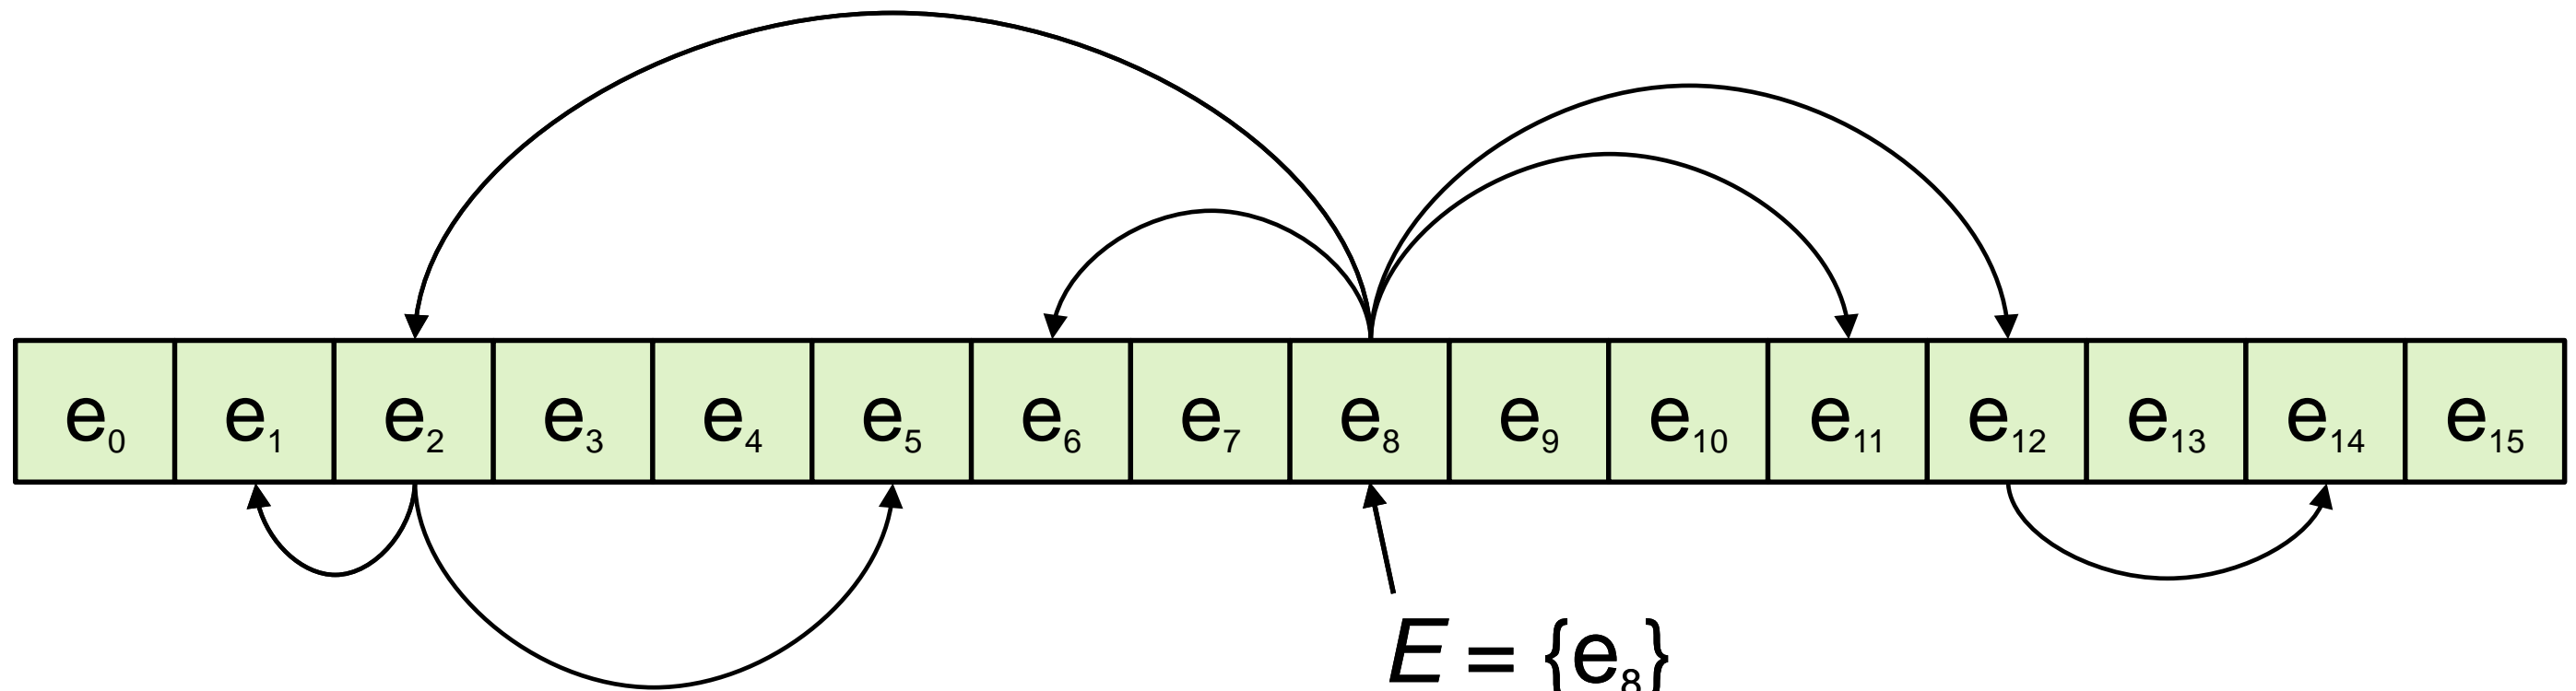
$$E = \{e_8\}$$

$$R_E = \{e_1, e_2, e_5, e_{11}\}$$

$$U_E = \{e_1, e_2, e_5, e_8, e_{11}\}$$

Sharing a Container – Avoiding Race Conditions

T_{n+1}



$$E = \{e_8\}$$

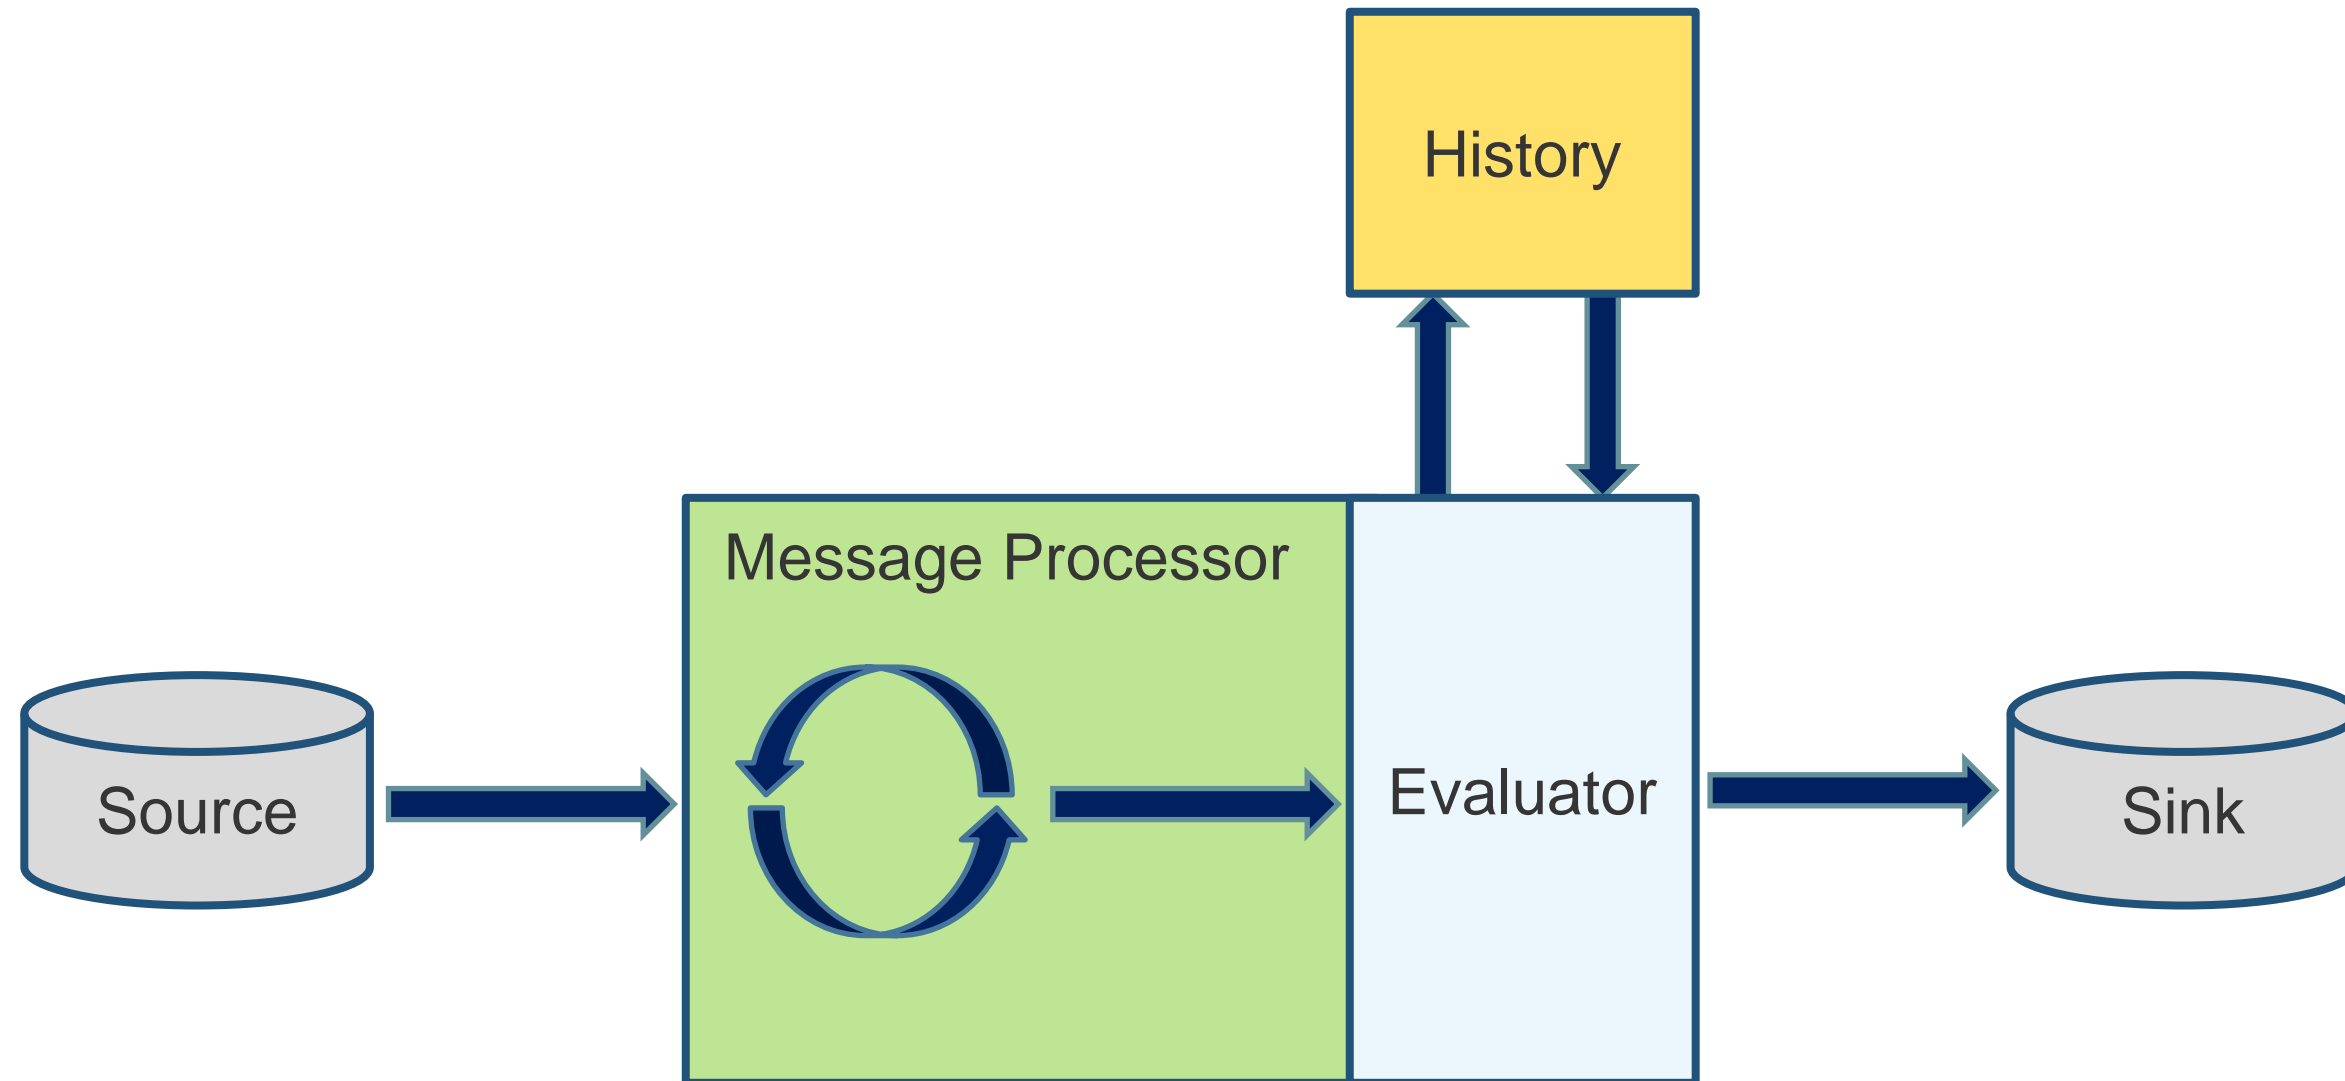
$$R_E = \{e_1, e_2, e_5, e_6, e_{11}, e_{12}, e_{14}\}$$

$$U_E = \{e_1, e_2, e_5, e_6, e_8, e_{11}, e_{12}, e_{14}\}$$

Motivating Problem – Reactive Message Processor

- Receives a continuous stream of input messages
- Generates output only when something interesting happens
 - A **history** must be maintained to detect relevant changes
- Reading the state of the history is never required
- Every message input requires a write operation to the history
 - Which uses one or more containers

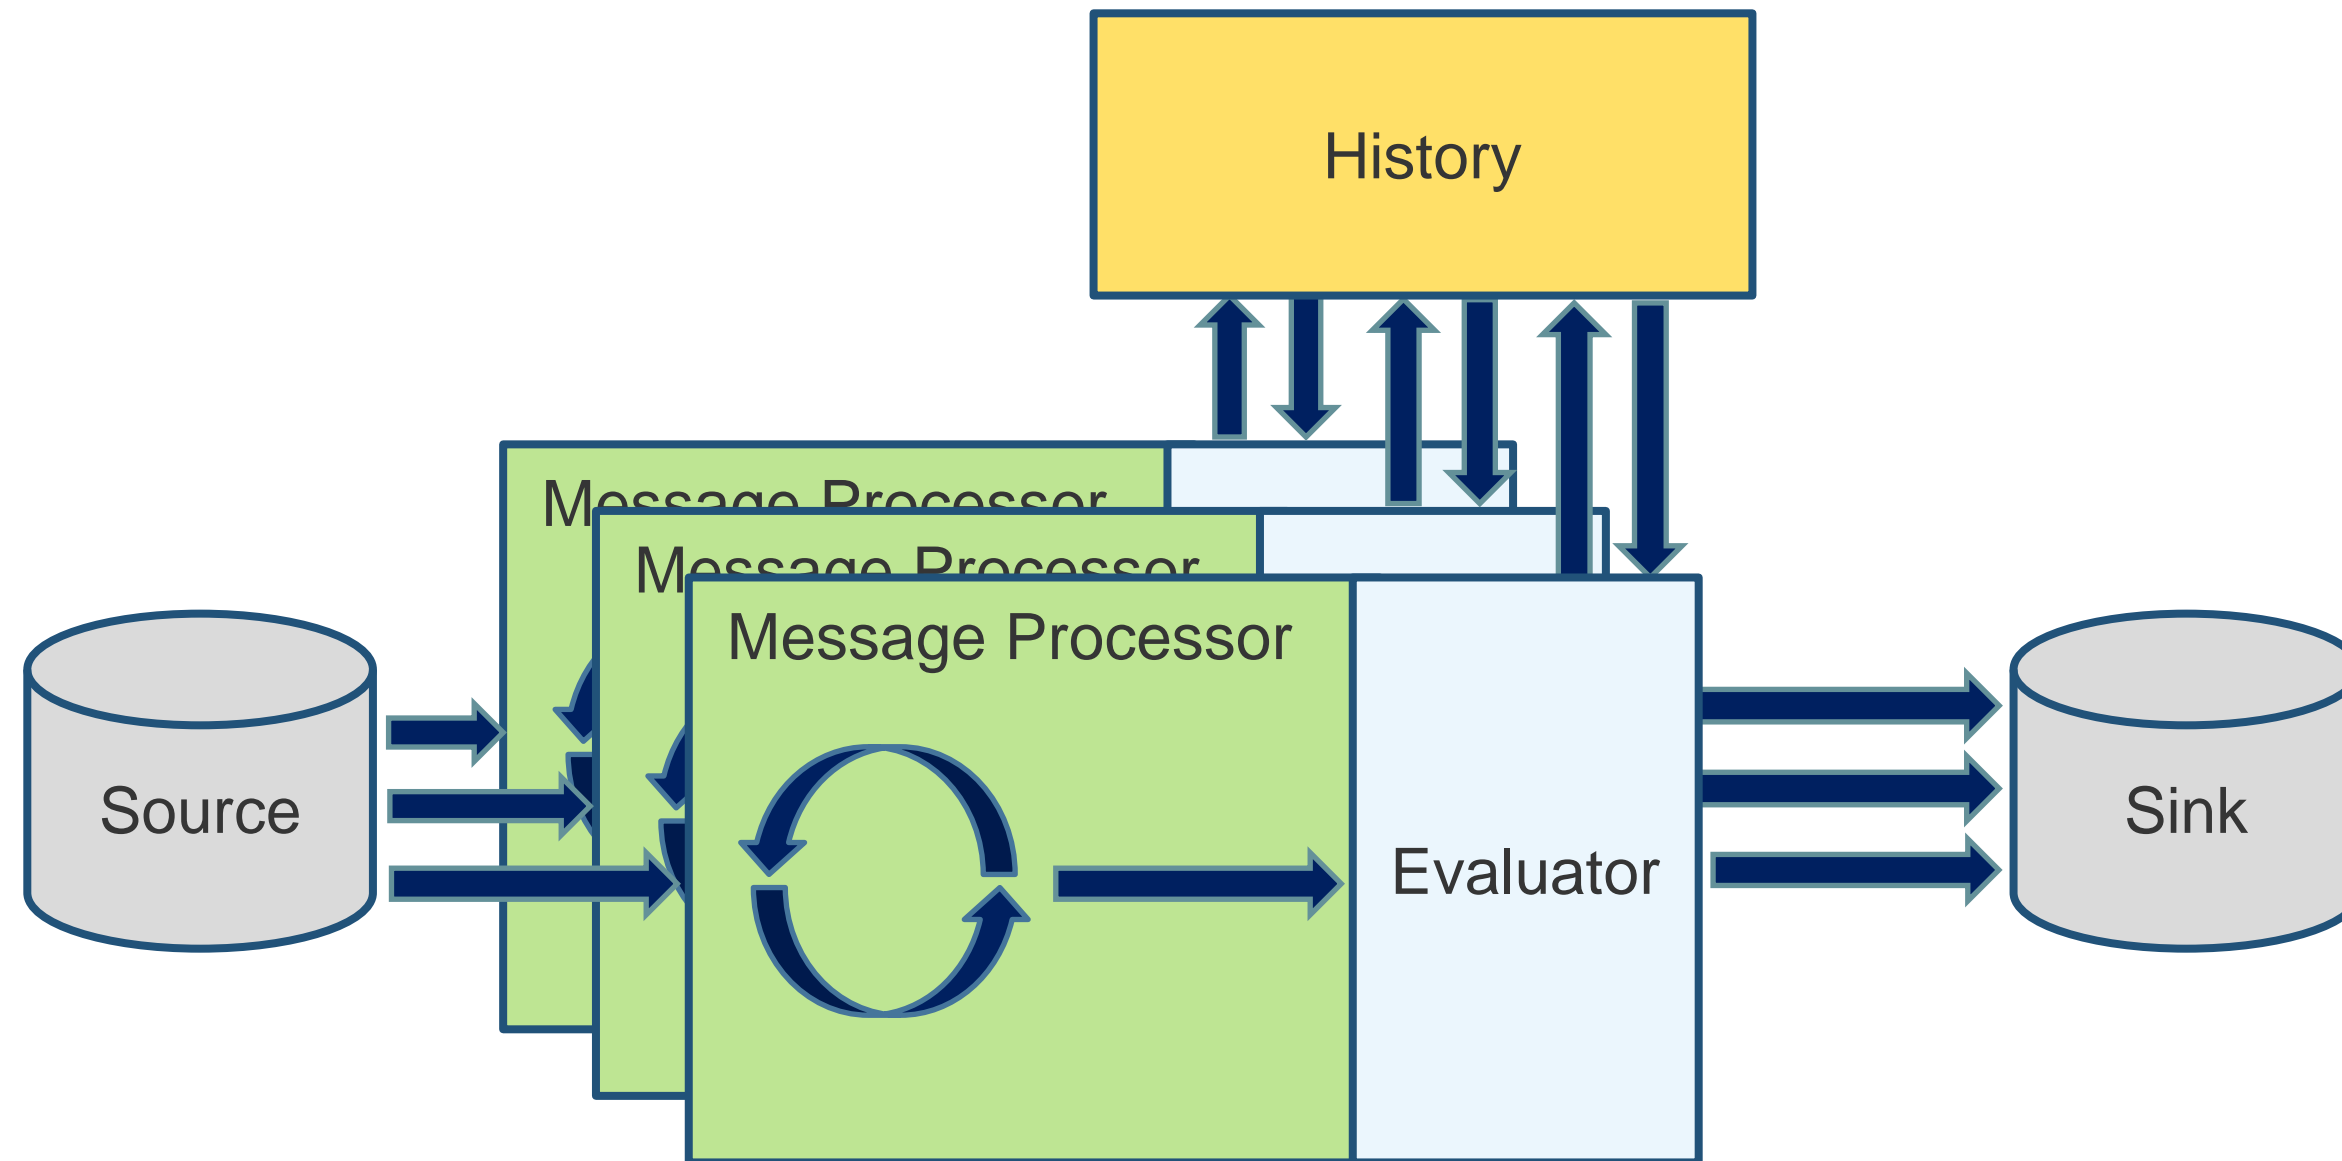
Motivating Problem – Reactive Message Processor



Motivating Problem – Reactive Message Processor

- Receives a continuous stream of input messages
- Generates output only when something special happens
 - History must be maintained to detect “special” changes
- Reading the state of the history is never required
- Every message input requires a write operation to a container
 - Which uses one or more containers
- **It must scale to multiple threads!**

Motivating Problem – Reactive Message Processor



- **Atomic:** Each modification of an update group is treated as a single transaction, which either succeeds completely, or fails completely
- **Consistent:** Each transaction can only bring the update group (and the enclosing container) from one valid state to another, maintaining all invariants
- **Isolated:** Each transaction must ensure that concurrent execution of other transactions leaves its update group (and the container) in the same state that would have been obtained as if the transactions were executed in some valid sequential order

- Instantiate a per-container mutex and perform updates in a single critical section guarded by the mutex
- Upside
 - It works
 - Easy to understand and think about
- Downside
 - **Does not scale**

- Divide the set of elements into individual *shards* such that the members of each element's related group are also in the shard
- Updates to each shard are performed by only one thread dedicated to servicing that shard
- Upside
 - Good performance
- Downside
 - Increased complexity
 - It works IFF the data is amenable to sharding

- Instantiate a per-element mutex, acquire the mutexes for all the elements in the update group, then update and release them all
- Upside
 - Seems like it should work (at least for the case where U_E has exactly one element)
 - Slightly more difficult to understand and think about
- Downside
 - Some mutex implementations are not small
 - What if E_0 's related group contains E_1 , and E_1 's related group contains E_0 ?
 - **Trap!**

Some Solutions – Per-Element Mutex

- Instantiate a per-element mutex, acquire the mutexes for all the elements in the update group, then update and release them all
- Upside
 - Seems like it should work
 - Slightly more difficult to understand about
- Downside
 - Some mutex implementations are not small
 - What if E_0 's related group contains E_1 , and E_1 's related group contains E_0 ?
 - **Trap!**



Another Solution – Strict Timestamp Ordering (STO)

- STO is one of *many* database concurrency control algorithms
- It is a transactional, *timestamp-based* algorithm
- **Timestamp**
 - A monotonically increasing value that indicates the age of a transaction
 - A lower timestamp value (TSV) indicates an older transaction
 - A higher timestamp value indicates a newer transaction
 - A transaction's TSV is its "birthday"
- STO uses timestamps to serialize concurrent transactions

- When each transaction begins it is assigned a unique timestamp from an authoritative, universal timestamp source
- If two transactions are attempting to write the same update group at the same time, the transaction with the lower timestamp goes first
 - Younger transactions always wait for older transactions
 - Older transactions never wait for younger transactions, they give up (roll back) and try again
- STO operation schedules are serializable and deadlock-free
 - Price for deadlock-freedom is the potential restart of a transaction

Strict Timestamp Ordering – General Case

- An element E with
 - A read timestamp rd_tsv
 - A write timestamp wr_tsv
- A transaction TX_1 with timestamp tsv
 - Functions $begin()$, $commit()$, and $rollback()$
 - Facility for TX_1 to "acquire" E
- Function $update()$ to update E
- Function $read()$ to read from E

Transaction TX_1 calls $read(E)$

- If $[TX_1.tsv < E.wr_tsv]$ then
 - A younger transaction (TX_2) has already written E , so call $rollback(TX_1)$
- If $[TX_1.tsv > E_n.wr_tsv]$ then
 - An older transaction (TX_0) is in progress involving E
 - TX_1 waits until all older transactions (i.e., TX_0) are complete
 - TX_1 acquires E , assigns $TX_1.tsv$ to $E.rd_tsv$, and reads E 's contents

Transaction TX_1 calls `update(E)`

- If $[TX_1.tsv < E.rd_tsv]$ then
 - A younger transaction (TX_2) has already read from E , so call `rollback(TX_1)`
- If $[TX_1.tsv < E.wr_tsv]$ then
 - A younger transaction (TX_2) has already written to E , so call `rollback(TX_1)`
- If $[TX_1.tsv > E_n.wr_tsv]$ then
 - An older transaction (TX_0) is in progress involving E
 - TX_1 waits until all older transactions (i.e., TX_0) are complete
 - TX_1 acquires E , assigns $TX_1.tsv$ to $E.wr_tsv$, and updates E 's contents

- For our problem, we assume all operations are updates
 - Then we need only one timestamp per element, `E.tsv`

Transaction `TX1` calls `update(E)`

- If `[TX1.tsv < E.tsv]` then
 - A younger transaction (`TX2`) has already written to `E`, so call `rollback(TX1)`
- If `[TX1.tsv > E.tsv]` then
 - An older transaction (`TX0`) is in progress involving `E`
 - `TX1` waits until all older transactions (i.e., `TX0`) are complete
 - `TX1` acquires `E`, assigns `TX1.tsv` to `E.tsv`, and updates `E`'s contents

- Need some basic synchronization components
 - Mutex
 - Condition variable
 - Atomic pointer
 - Atomic compare and exchange
- Need a class that represents a lockable item (element)
- Need a class that manages transactions on behalf of each thread
 - Transactions are represented by a timestamp value (TSV)

- Threads share containers holding lockable items
- Threads will initiate transactions, and each thread owns its transactions
- A transaction's key operations (begin, acquire, commit, rollback) can only be performed by the owning thread
- A transaction manager acquires lockable items on behalf of its owning thread
 - Other logic in the owning thread performs actual update operations
- In a transaction, the entire update group must be acquired before the thread applies any write operation to any member of that group

Implementing STO in C++ – Design Choices

thread_function(*shared_collection*)

[Create a transaction manager]

while [Work remains to be done]

[Begin a new transaction]

[Find target element *E* in *shared_collection*]

[Determine the membership of the update group to the extent possible]

while [Update group members remain unacquired AND acquisitions have all succeeded]
[Starting with *E*, attempt to acquire the next unacquired member of the update group]
[Revise the membership of the update group, if necessary]

endwhile

if [All members of the update group were acquired]
[Apply write operations to the members of the update group]
[Commit the transaction, loop back to the top, get new work]

else

[Roll back the transaction, loop back to the top, and try again]

endif

endwhile

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
  [Begin a new transaction]
```

```
  [Find target element E in shared_collection]
```

```
  [Determine the membership of the update group to the extent possible]
```

```
  while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
    [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
    [Revise the membership of the update group, if necessary]
```

```
  endwhile
```

```
  if [All members of the update group were acquired]
```

```
    [Apply write operations to the members of the update group]
```

```
    [Commit the transaction, loop back to the top, get new work]
```

```
  else
```

```
    [Roll back the transaction, loop back to the top, and try again]
```

```
  endif
```

```
endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
[Begin a new transaction]
```

```
[Find target element E in shared_collection]
```

```
[Determine the membership of the update group to the extent possible]
```

```
while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
[Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
[Revise the membership of the update group, if necessary]
```

```
endwhile
```

```
if [All members of the update group were acquired]
```

```
[Apply write operations to the members of the update group]
```

```
[Commit the transaction, loop back to the top, get new work]
```

```
else
```

```
[Roll back the transaction, loop back to the top, and try again]
```

```
endif
```

```
endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
[Begin a new transaction]
```

```
[Find target element E in shared_collection]
```

```
[Determine the membership of the update group to the extent possible]
```

```
while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
[Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
[Revise the membership of the update group, if necessary]
```

```
endwhile
```

```
if [All members of the update group were acquired]
```

```
[Apply write operations to the members of the update group]
```

```
[Commit the transaction, loop back to the top, get new work]
```

```
else
```

```
[Roll back the transaction, loop back to the top, and try again]
```

```
endif
```

```
endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```


Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
[Begin a new transaction]
```

```
[Find target element E in shared_collection]
```

```
[Determine the membership of the update group to the extent possible]
```

```
while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
[Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
[Revise the membership of the update group, if necessary]
```

```
endwhile
```

```
if [All members of the update group were acquired]
```

```
[Apply write operations to the members of the update group]
```

```
[Commit the transaction, loop back to the top, get new work]
```

```
else
```

```
[Roll back the transaction, loop back to the top, and try again]
```

```
endif
```

```
endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
    [Create a transaction manager]
```

```
    while [Work remains to be done]
```

```
        [Begin a new transaction]
```

```
        [Find target element E in shared_collection]
```

```
        [Determine the membership of the update group to the extent possible]
```

```
        while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
            [Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
            [Revise the membership of the update group, if necessary]
```

```
        endwhile
```

```
        if [All members of the update group were acquired]
```

```
            [Apply write operations to the members of the update group]
```

```
            [Commit the transaction, loop back to the top, get new work]
```

```
        else
```

```
            [Roll back the transaction, loop back to the top, and try again]
```

```
        endif
```

```
    endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
[Begin a new transaction]
```

```
[Find target element E in shared_collection]
```

```
[Determine the membership of the update group to the extent possible]
```

```
while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
[Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
[Revise the membership of the update group, if necessary]
```

```
endwhile
```

```
if [All members of the update group were acquired]
```

```
[Apply write operations to the members of the update group]
```

```
[Commit the transaction, loop back to the top, get new work]
```

```
else
```

```
[Roll back the transaction, loop back to the top, and try again]
```

```
endif
```

```
endwhile
```

Implementing STO in C++ – Design Choices

```
thread_function(shared_collection)
```

```
[Create a transaction manager]
```

```
while [Work remains to be done]
```

```
[Begin a new transaction]
```

```
[Find target element E in shared_collection]
```

```
[Determine the membership of the update group to the extent possible]
```

```
while [Update group members remain unacquired AND acquisitions have all succeeded]
```

```
[Starting with E, attempt to acquire the next unacquired member of the update group]
```

```
[Revise the membership of the update group, if necessary]
```

```
endwhile
```

```
if [All members of the update group were acquired]
```

```
[Apply write operations to the members of the update group]
```

```
[Commit the transaction, loop back to the top, get new work]
```

```
else
```

```
[Roll back the transaction, loop back to the top, and try again]
```

```
endif
```

```
endwhile
```


Prerequisites

Inclusions, Type Aliases, Forward Declarations

```
//- Stuff we need.  
//  
#include <cstring>  
#include <atomic>  
#include <chrono>  
#include <condition_variable>  
#include <functional>  
#include <future>  
#include <mutex>  
#include <random>  
#include <thread>  
#include <type_traits>  
#include <vector>  
  
using tsv_type      = uint64_t;      //- Timestamp value  
using tx_id_type    = uint64_t;      //- Transaction ID  
using item_id_type  = uint64_t;      //- Item ID  
  
class transaction_manager;  
class lockable_item;  
class stopwatch;
```

Class `lockable_item`

Class Overview – lockable_item

```
class lockable_item
{
public:
    lockable_item();

    item_id_type    id() const noexcept;
    tsv_type        last_tsv() const noexcept;

private:
    friend class transaction_manager;

    using atomic_txm_pointer = std::atomic<transaction_manager*>;
    using atomic_item_id     = std::atomic<item_id_type>;

    atomic_txm_pointer mp_owning_tx;    //- Pointer to transaction manager that owns this item
    tsv_type           m_last_tsv;     //- Timestamp of last owner
    item_id_type        m_item_id;     //- For debugging/tracking/logging

    static atomic_item_id sm_item_id_generator;
};
```

Member Functions – lockable_item

```
inline  
lockable_item::lockable_item()  
:    mp_owning_tx(nullptr), m_last_tsv(0), m_item_id(++sm_item_id_generator)  
{}
```

```
inline item_id_type  
lockable_item::id() const noexcept  
{  
    return m_item_id;  
}
```

```
inline tsv_type  
lockable_item::last_tsv() const noexcept  
{  
    return m_last_tsv;  
}
```

```
lockable_item::atomic_item_id    lockable_item::sm_item_id_generator = 0;
```

Class transaction_manager

Class Overview – transaction_manager

```
class transaction_manager
{
public:
    ~transaction_manager();
    transaction_manager(int log_level, FILE* fp=nullptr);

    tx_id_type id() const noexcept;
    tsv_type tsv() const noexcept;

    void begin();
    void commit();
    void rollback();

    bool acquire(lockable_item& item);

private:
    ...
};
```

Class Overview – transaction_manager

...

private:

```
using item_ptr_list = std::vector<lockable_item*>;  
using mutex         = std::mutex;  
using tx_lock       = std::unique_lock<std::mutex>;  
using cond_var      = std::condition_variable;  
using atomic_tsv     = std::atomic<tsv_type>;  
using atomic_tx_id   = std::atomic<tx_id_type>;
```

```
tx_id_type    m_tx_id;  
tsv_type      m_tx_tsv;  
item_ptr_list m_item_ptrs;  
mutex         m_mutex;  
cond_var      m_cond;  
FILE*         m_fp;  
int           m_log_level;
```

```
static atomic_tsv    sm_tsv_generator;  
static atomic_tx_id  sm_tx_id_generator;
```

...

Class Overview – transaction_manager

...

private:

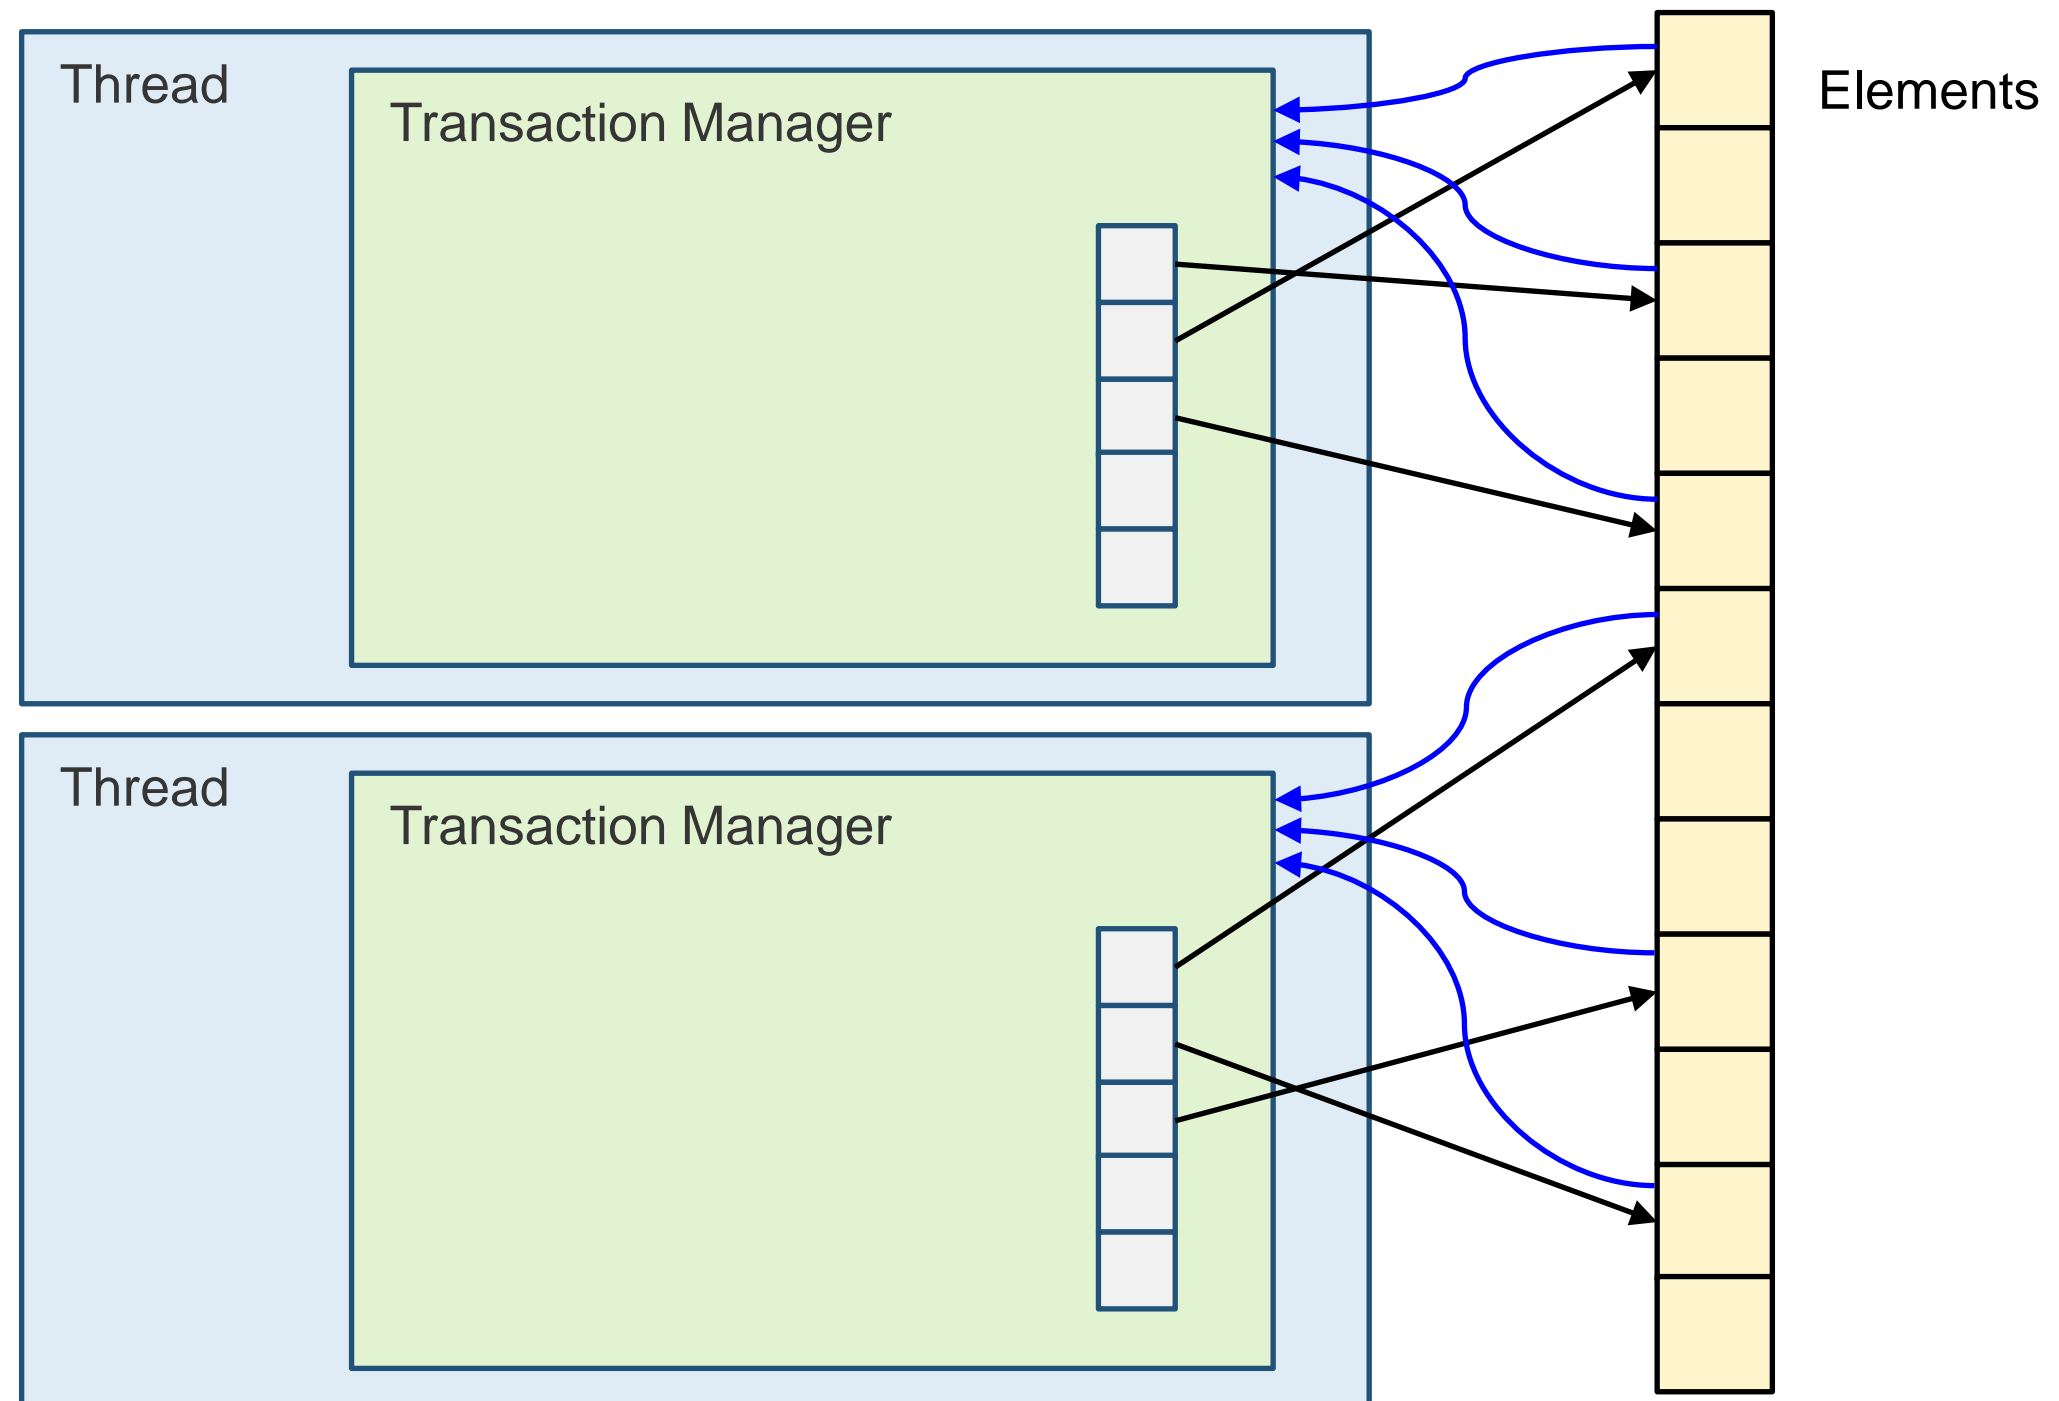
```
using item_ptr_list = std::vector<lockable_item*>;  
using mutex         = std::mutex;  
using tx_lock       = std::unique_lock<std::mutex>;  
using cond_var      = std::condition_variable;  
using atomic_tsv     = std::atomic<tsv_type>;  
using atomic_tx_id   = std::atomic<tx_id_type>;
```

```
tx_id_type    m_tx_id;  
tsv_type      m_tx_tsv;  
item_ptr_list m_item_ptrs;  
mutex         m_mutex;  
cond_var      m_cond;  
FILE*         m_fp;  
int           m_log_level;
```

```
static atomic_tsv    sm_tsv_generator;  
static atomic_tx_id  sm_tx_id_generator;
```

...

TransactionManager/Element Relationships



Class Overview – transaction_manager

...

```
private:
    void    log_begin() const;
    void    log_commit() const;
    void    log_rollback() const;
    void    log_acquisition_success(lockable_item const& item) const;
    void    log_acquisition_failure(lockable_item const& item) const;
    void    log_acquisition_same(lockable_item const& item) const;
    void    log_acquisition_waiting(lockable_item const& item, transaction_manager* p_curr_tx)
const;
};
```

```
transaction_manager::atomic_tsv    transaction_manager::sm_tsv_generator    = 0;
transaction_manager::atomic_tx_id   transaction_manager::sm_tx_id_generator = 0;
```

Member Functions – transaction_manager

```
transaction_manager::transaction_manager(int log_level, FILE* fp)
:   m_tx_id(++sm_tx_id_generator)
,   m_tx_tsv(0)
,   m_item_ptrs()
,   m_mutex()
,   m_cond()
,   m_fp(fp)
,   m_log_level(log_level)
{
    m_item_ptrs.reserve(100u);
}
```

Member Functions – transaction_manager

```
inline tx_id_type  
transaction_manager::id() const noexcept  
{  
    return m_tx_id;  
}
```

```
inline tsv_type  
transaction_manager::tsv() const noexcept  
{  
    return m_tx_tsv;  
}
```

Member Functions – transaction_manager::begin()

```
void  
transaction_manager::begin()  
{  
    m_mutex.lock();  
    m_tx_tsv = ++sm_tsv_generator;  
    m_mutex.unlock();  
}
```

Member Functions – transaction_manager::begin()

```
void  
transaction_manager::begin()  
{  
    m_mutex.lock();  
    m_tx_tsv = ++sm_tsv_generator;  
    m_mutex.unlock();  
}
```

Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock    lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();
}
```


Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock    lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();
}
```

Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock      lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();
}
```

Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock      lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();
}
```

Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock      lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();
}
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```


Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – transaction_manager::acquire()

```
bool
transaction_manager::acquire(lockable_item& item)
{
    while (true)
    {
        transaction_manager*    p_curr_tx = nullptr;

        if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
        {
            m_item_ptrs.push_back(&item);

            if (m_tx_tsv > item.m_last_tsv)
            {
                item.m_last_tsv = m_tx_tsv;
                return true;
            }

            ...
        }
    }
}
```


Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – transaction_manager::acquire()

```
while (true)
{
    transaction_manager*    p_curr_tx = nullptr;

    if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
    {
        m_item_ptrs.push_back(&item);

        if (m_tx_tsv > item.m_last_tsv)
        {
            item.m_last_tsv = m_tx_tsv;
            return true;
        }
        else
        {
            return false;
        }
    }
    else
    ...
}
```

Member Functions – transaction_manager::acquire()

```
while (true)
{
    transaction_manager*    p_curr_tx = nullptr;

    if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
    {
        m_item_ptrs.push_back(&item);

        if (m_tx_tsv > item.m_last_tsv)
        {
            item.m_last_tsv = m_tx_tsv;
            return true;
        }
        else
        {
            return false;
        }
    }
    else
    ...
}
```

Atomic Compare and Exchange

```
template<class T>
inline bool
atomic<T>::compare_exchange_strong(T& expected, T desired,
                                   std::memory_order mo = std::memory_order_seq_cst ) noexcept;
```

- Paraphrasing cppreference.com:
 - Atomically compares the representation of **this* with that of *expected*, and if those are bitwise-equal, replaces the representation of **this* with *desired* and returns *true*;
 - Otherwise, loads the actual value stored in **this* into *expected* and returns *false*.

Member Functions – transaction_manager::acquire()

```
while (true)
{
    transaction_manager*    p_curr_tx = nullptr;

    if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
    {
        m_item_ptrs.push_back(&item);

        if (m_tx_tsv > item.m_last_tsv)
        {
            item.m_last_tsv = m_tx_tsv;
            return true;
        }
        else
        {
            return false;
        }
    }
    else
    ...
}
```

Member Functions – transaction_manager::acquire()

```
while (true)
{
    transaction_manager*    p_curr_tx = nullptr;

    if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
    {
        m_item_ptrs.push_back(&item);

        if (m_tx_tsv > item.m_last_tsv)
        {
            item.m_last_tsv = m_tx_tsv;
            return true;
        }
        else
        {
            return false;
        }
    }
    else
    ...
}
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – transaction_manager::acquire()

```
...  
  
if (m_tx_tsv > item.m_last_tsv)  
{  
    item.m_last_tsv = m_tx_tsv;  
    return true;  
}  
else  
{  
    return false;  
}  
}  
else  
{  
    if (p_curr_tx == this)  
    {  
        return true;  
    }  
    ...  
}
```


Member Functions – transaction_manager::acquire()

```
...  
  
if (m_tx_tsv > item.m_last_tsv)  
{  
    item.m_last_tsv = m_tx_tsv;  
    return true;  
}  
else  
{  
    return false;  
}  
}  
else  
{  
    if (p_curr_tx == this)  
    {  
        return true;  
    }  
    ...  
}
```

Member Functions – transaction_manager::acquire()

```
...  
  
if (m_tx_tsv > item.m_last_tsv)  
{  
    item.m_last_tsv = m_tx_tsv;  
    return true;  
}  
else  
{  
    return false;  
}  
}  
else  
{  
    if (p_curr_tx == this)  
    {  
        return true;  
    }  
    ...  
}
```

Member Functions – transaction_manager::acquire()

```
...  
  
if (m_tx_tsv > item.m_last_tsv)  
{  
    item.m_last_tsv = m_tx_tsv;  
    return true;  
}  
else  
{  
    return false;  
}  
}  
else  
{  
    if (p_curr_tx == this)  
    {  
        return true;  
    }  
    ...  
}
```

Member Functions – `transaction_manager::acquire()`

```
acquire(lockable_item item)
    while (true)
        if [I acquire item]
            [Add item to the list of items that I own]

            if [My timestamp is newer than item's last timestamp]
                [Set item's timestamp to my timestamp]
                [Success – return and keep going]
            else
                [Failure – return and roll back]
            endif
        else
            if [I already own item]
                [Success – return and keep going]
            else
                [I'll wait until the current owner releases item and then re-try]
            endif
        endif
    endwhile
```

Member Functions – transaction_manager::acquire()

```
...  
  
if (m_tx_tsv > item.m_last_tsv)  
{  
    item.m_last_tsv = m_tx_tsv;  
    return true;  
}  
else  
{  
    return false;  
}  
}  
else  
{  
    if (p_curr_tx == this)  
    {  
        return true;  
    }  
    ...  
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```


Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock      lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::commit()

```
void
transaction_manager::commit()
{
    tx_lock      lock(m_mutex);

    while (m_item_ptrs.size() != 0)
    {
        m_item_ptrs.back()->mp_owning_tx.store(nullptr);
        m_item_ptrs.pop_back();
    }
    m_cond.notify_all();    //- This is being called by a TM object in another thread
}
```

Member Functions – transaction_manager::acquire()

```
else
{
    if (p_curr_tx == this)
    {
        return true;
    }
    else
    {
        tx_lock    lock(p_curr_tx->m_mutex);

        while (item.mp_owning_tx.load() == p_curr_tx)
        {
            if (p_curr_tx->m_tx_tsv > m_tx_tsv)
            {
                return false;
            }
            p_curr_tx->m_cond.wait(lock);
        }
    }
}
}
```

Member Functions – transaction_manager::acquire()

```
while (true)
{
    transaction_manager*    p_curr_tx = nullptr;

    if (item.mp_owning_tx.compare_exchange_strong(p_curr_tx, this))
    {
        m_item_ptrs.push_back(&item);

        if (m_tx_tsv > item.m_last_tsv)
        {
            item.m_last_tsv = m_tx_tsv;
            return true;
        }
        else
        {
            return false;
        }
    }
    else
    ...
}
```


Testing

- Create functions to update a collection of shared items
 - Make sure that data races are possible and can be detected
- Measure single-threaded updates – baseline
- Measure multi-threaded updates with induced data races
- Measure multi-threaded updates guarded by a single critical section
- Measure multi-threaded transactional updates
- 8-Core (16 hypercores), Ubuntu 18.04, GCC 10.2

Testing – Includes, etc.

```
#include "transaction.hpp"

using namespace tx;
using namespace std;
using namespace std::chrono_literals;

//- Forward declarations and common type aliases common to the test functions below().
//
struct test_item;

using item_list  = std::vector<test_item>;
using index_list = std::vector<size_t>;

using entropy    = random_device;
using prn_gen    = mt19937_64;
using int_dist   = uniform_int_distribution<>;
using hasher     = hash<string_view>;
```

Testing – test_item

```
// - An updatable type susceptible to data races.
//
struct test_item : public lockable_item
{
    static constexpr size_t    buf_size = 32;

    char    ma_chars[buf_size];

    void    st_update(FILE* fp, prn_gen& gen, int_dist& char_dist);
    void    tx_update(transaction const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist);
};
```

Testing – test_item::st_update()

```
// - Updates on-board data for the non-STO tests. Checks to see if a race has occurred.
//
void
test_item::st_update(FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, item %zd\n", this->id());
    }
}
```

Testing – test_item::st_update()

```
// - Updates on-board data for the non-STO tests. Checks to see if a race has occurred.
//
void
test_item::st_update(FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, item %zd\n", this->id());
    }
}
```

Testing – test_item::st_update()

```
// - Updates on-board data for the non-STO tests. Checks to see if a race has occurred.
//
void
test_item::st_update(FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, item %zd\n", this->id());
    }
}
```

Testing – test_item::st_update()

```
// - Updates on-board data for the non-STO tests. Checks to see if a race has occurred.
//
void
test_item::st_update(FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, item %zd\n", this->id());
    }
}
```


Testing – test_item::tx_update()

// - Updates on-board data for the ST0 tests. Checks to see if a race has occurred.

//

void

test_item::tx_update(transaction_manager const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist)

{

char local_chars[buf_size];

string_view local_view(local_chars, buf_size);

string_view shared_view(ma_chars, buf_size);

hasher hash;

for (size_t i = 0; i < buf_size; ++i)

{

local_chars[i] = ma_chars[i] = (char) char_dist(gen);

}

if (hash(shared_view) != hash(local_view))

{

fprintf(fp, "RACE FOUND!, TX %zd item %zd\n", tx.id(), this->id());

}

}

Testing – test_item::tx_update()

```
// - Updates on-board data for the ST0 tests. Checks to see if a race has occurred.
//
void
test_item::tx_update(transaction_manager const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, TX %zd item %zd\n", tx.id(), this->id());
    }
}
```

Testing – test_item::tx_update()

```
//- Updates on-board data for the ST0 tests.  Checks to see if a race has occurred.
//
void
test_item::tx_update(transaction_manager const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!,  TX %zd  item %zd\n", tx.id(), this->id());
    }
}
```

Testing – test_item::tx_update()

```
//- Updates on-board data for the ST0 tests.  Checks to see if a race has occurred.
//
void
test_item::tx_update(transaction_manager const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, TX %zd item %zd\n", tx.id(), this->id());
    }
}
```

Testing – test_item::tx_update()

```
// - Updates on-board data for the ST0 tests. Checks to see if a race has occurred.
//
void
test_item::tx_update(transaction_manager const& tx, FILE* fp, prn_gen& gen, int_dist& char_dist)
{
    char          local_chars[buf_size];
    string_view    local_view(local_chars, buf_size);
    string_view    shared_view(ma_chars, buf_size);
    hasher         hash;

    for (size_t i = 0; i < buf_size; ++i)
    {
        local_chars[i] = ma_chars[i] = (char) char_dist(gen);
    }

    if (hash(shared_view) != hash(local_view))
    {
        fprintf(fp, "RACE FOUND!, TX %zd item %zd\n", tx.id(), this->id());
    }
}
```

Transactional Element Access Test

Transactional Element Access Test

```
void
tx_access_test(item_list& items, FILE* fp, size_t tx_count, size_t refs_count)
{
    entropy        rd;
    prn_gen        gen(rd());
    int_dist       refs_index_dist(0, (int)(items.size()-1));
    int_dist       refs_count_dist(1, (int) refs_count);
    int_dist       char_dist(0, 127);

    stopwatch      sw;
    index_list     indices;
    size_t         index;

    transaction_manager tx(1, fp);
    bool           acquired;

    sw.start();

    for (size_t i = 0; i < tx_count; ++i)
    {
        ...
    }
}
```

Transactional Element Access Test

```
void
tx_access_test(item_list& items, FILE* fp, size_t tx_count, size_t refs_count)
{
    entropy      rd;
    prn_gen      gen(rd());
    int_dist     refs_index_dist(0, (int)(items.size()-1));
    int_dist     refs_count_dist(1, (int) refs_count);
    int_dist     char_dist(0, 127);

    stopwatch    sw;
    index_list    indices;
    size_t        index;

    transaction_manager tx(1, fp);
    bool           acquired;

    sw.start();

    for (size_t i = 0; i < tx_count; ++i)
    {
        ...
    }
}
```


Transactional Element Access Test

```
void
tx_access_test(item_list& items, FILE* fp, size_t tx_count, size_t refs_count)
{
    entropy      rd;
    prn_gen      gen(rd());
    int_dist     refs_index_dist(0, (int)(items.size()-1));
    int_dist     refs_count_dist(1, (int) refs_count);
    int_dist     char_dist(0, 127);

    stopwatch    sw;
    index_list   indices;
    size_t       index;

    transaction_manager tx(1, fp);
    bool         acquired;

    sw.start();

    for (size_t i = 0; i < tx_count; ++i)
    {
        ...
    }
}
```

Transactional Element Access Test

...

```
for (size_t i = 0; i < tx_count; ++i)
{
    //- Compute the size of the update group
    //
    indices.clear();
    refs_count = refs_count_dist(gen);

    //- Compute the membership of the update group
    //
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = refs_index_dist(gen);
        indices.push_back(index);
    }
}
```

...

Transactional Element Access Test

...

```
for (size_t i = 0; i < tx_count; ++i)
{
    //- Compute the size of the update group
    //
    indices.clear();
    refs_count = refs_count_dist(gen);

    //- Compute the membership of the update group
    //
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = refs_index_dist(gen);
        indices.push_back(index);
    }
}
```

...

Transactional Element Access Test

...

```
//- Compute the membership of the update group
//
for (size_t j = 0; j < refs_count; ++j)
{
    index = refs_index_dist(gen);
    indices.push_back(index);
}
```

```
tx.begin();
acquired = true;
```

```
//- Acquire the members of the update group
//
for (size_t j = 0; acquired && j < refs_count; ++j)
{
    index = indices[j];
    acquired = tx.acquire(items[index]);
}
```

...

Transactional Element Access Test

...

```
//- Acquire the members of the update group
//
for (size_t j = 0; acquired && j < refs_count; ++j)
{
    index = indices[j];
    acquired = tx.acquire(items[index]);
}
```

```
//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
```

...

Transactional Element Access Test

...

// - Acquire the members of the update group

//

```
for (size_t j = 0; acquired && j < refs_count; ++j)
{
    index = indices[j];
    acquired = tx.acquire(items[index]);
}
```

// - Modify the members of the update group

//

```
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
```

...

Transactional Element Access Test

...

```
//- Acquire the members of the update group
//
for (size_t j = 0; acquired && j < refs_count; ++j)
{
    index    = indices[j];
    acquired = tx.acquire(items[index]);
}
```

```
//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
```

...

Transactional Element Access Test

...

```
//- Acquire the members of the update group
//
for (size_t j = 0; acquired && j < refs_count; ++j)
{
    index    = indices[j];
    acquired = tx.acquire(items[index]);
}
```

```
//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
```

...

Transactional Element Access Test

...

```
//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
else
{
    tx.rollback();
}
...
```

Transactional Element Access Test

...

```
//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
else
{
    tx.rollback();
}
...
```

Transactional Element Access Test


```
...

//- Modify the members of the update group
//
if (acquired)
{
    for (size_t j = 0; j < refs_count; ++j)
    {
        index = indices[j];
        items[index].tx_update(tx, fp, gen, char_dist);
    }
    tx.commit();
}
else
{
    tx.rollback();
}
}


sw.stop();
fprintf(fp, "TX %zd took %d msec\n", tx.id(), sw.milliseconds_elapsed<int>());
}
```

Test Results


Results – 10M items / 100K transactions / 20 refs / 16 threads




Results – 10M items / 100K transactions / 20 refs / 16 threads




Results – 10M items / 100K transactions / 20 refs / 16 threads




Results – 1M items / 100K transactions / 20 refs / 16 threads




Results – 1M items / 100K transactions / 20 refs / 16 threads




Results – 1M items / 100K transactions / 20 refs / 16 threads




Results – 100K items / 100K transactions / 20 refs / 16 threads




Results – 100K items / 100K transactions / 20 refs / 16 threads




Results – 100K items / 100K transactions / 20 refs / 16 threads




Results – 10K items / 100K transactions / 20 refs / 16 threads




Results – 10K items / 100K transactions / 20 refs / 16 threads



Results – 10K items / 100K transactions / 20 refs / 16 threads



Results – 10K items / 100K transactions / 20 refs / 16 threads



Summary

- These tools operate upon containers and with elements, but don't require changing the containers themselves
- There is an assumption that the container's internal structure is unchanged while transactions are in progress
 - Consider the case of a **vector** resize
 - Consider the case of adding an element to a **map**
 - Could this be handled by a per-container shared mutex?

- So far, only `std::vector` has been used
 - The maximum number of elements is pre-allocated and resizes don't occur
- To obtain a container that is resizable
 - Create a home-grown hash table using `std::vector`
 - Each element of the vector is a hash bucket
 - Hash buckets have member functions for adding, finding, erasing elements
 - Hash buckets are locked during transactions, and their contents updated
 - With a good hash function, lookup time is quite fast (but not as fast as indexing)

- Some threads could starve
 - Transactions might become stale
- Other container types may be amenable... ?
- Lots of room for more work

Thank You for Attending!

Talk: github.com/BobSteagall/CppCon2020

Blog: bobsteagall.com